



Survey of Fault-tolerant LLM Training

JIANG, Zhihan

Ph.D. Student

The Chinese University of Hong Kong

2023-11-16

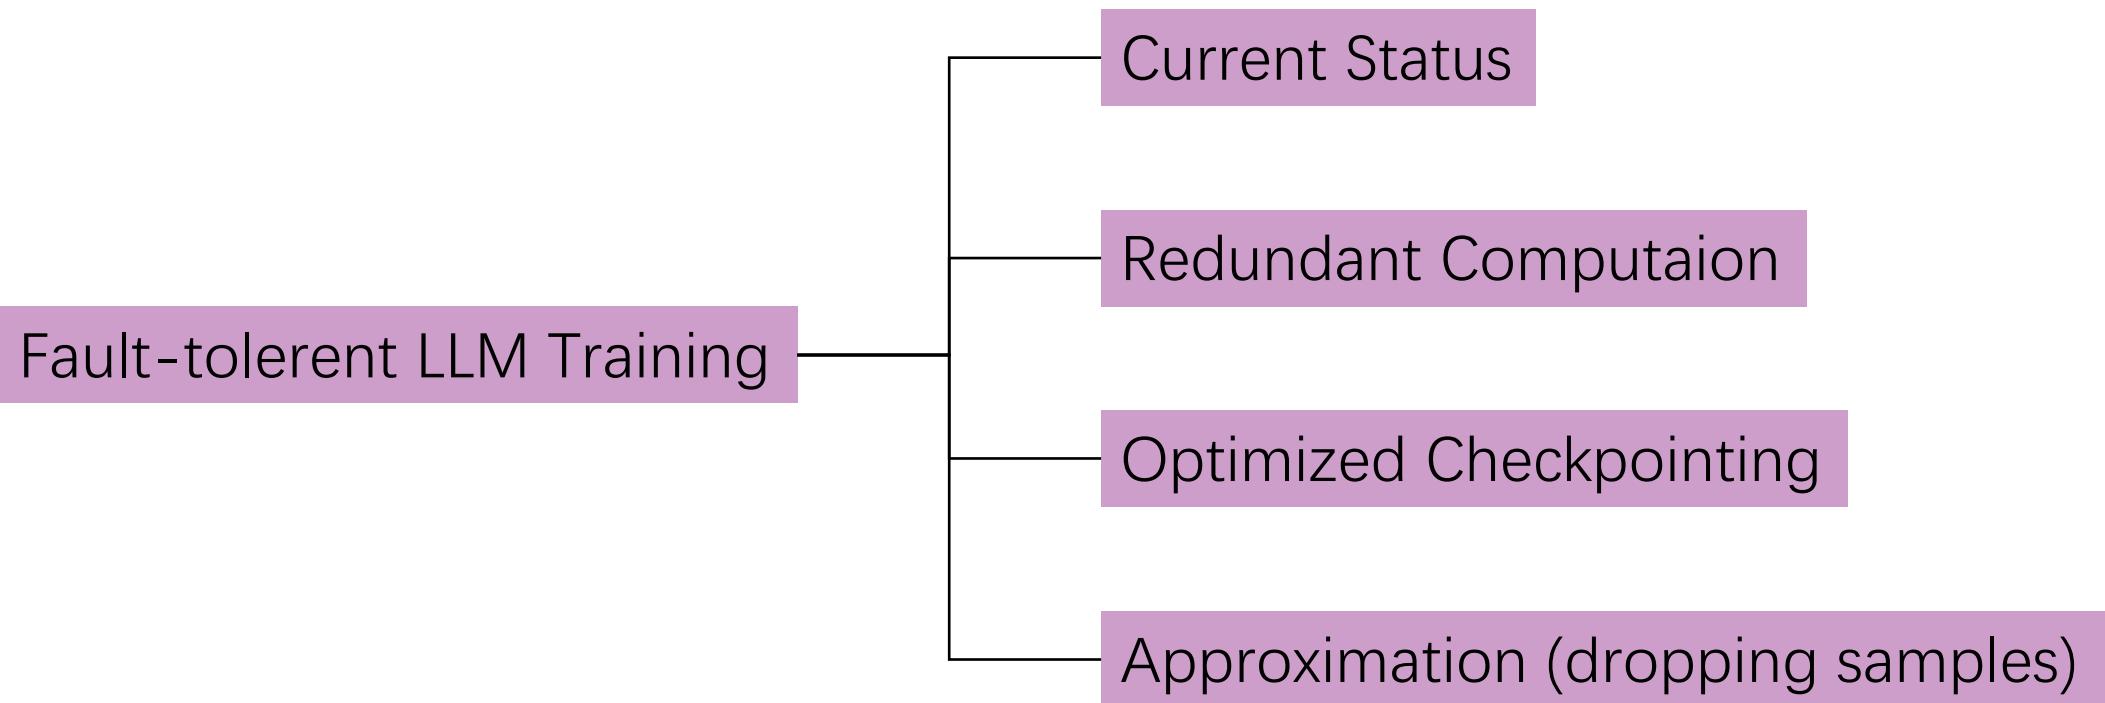


Automated Reliable Intelligent
Software Engineering

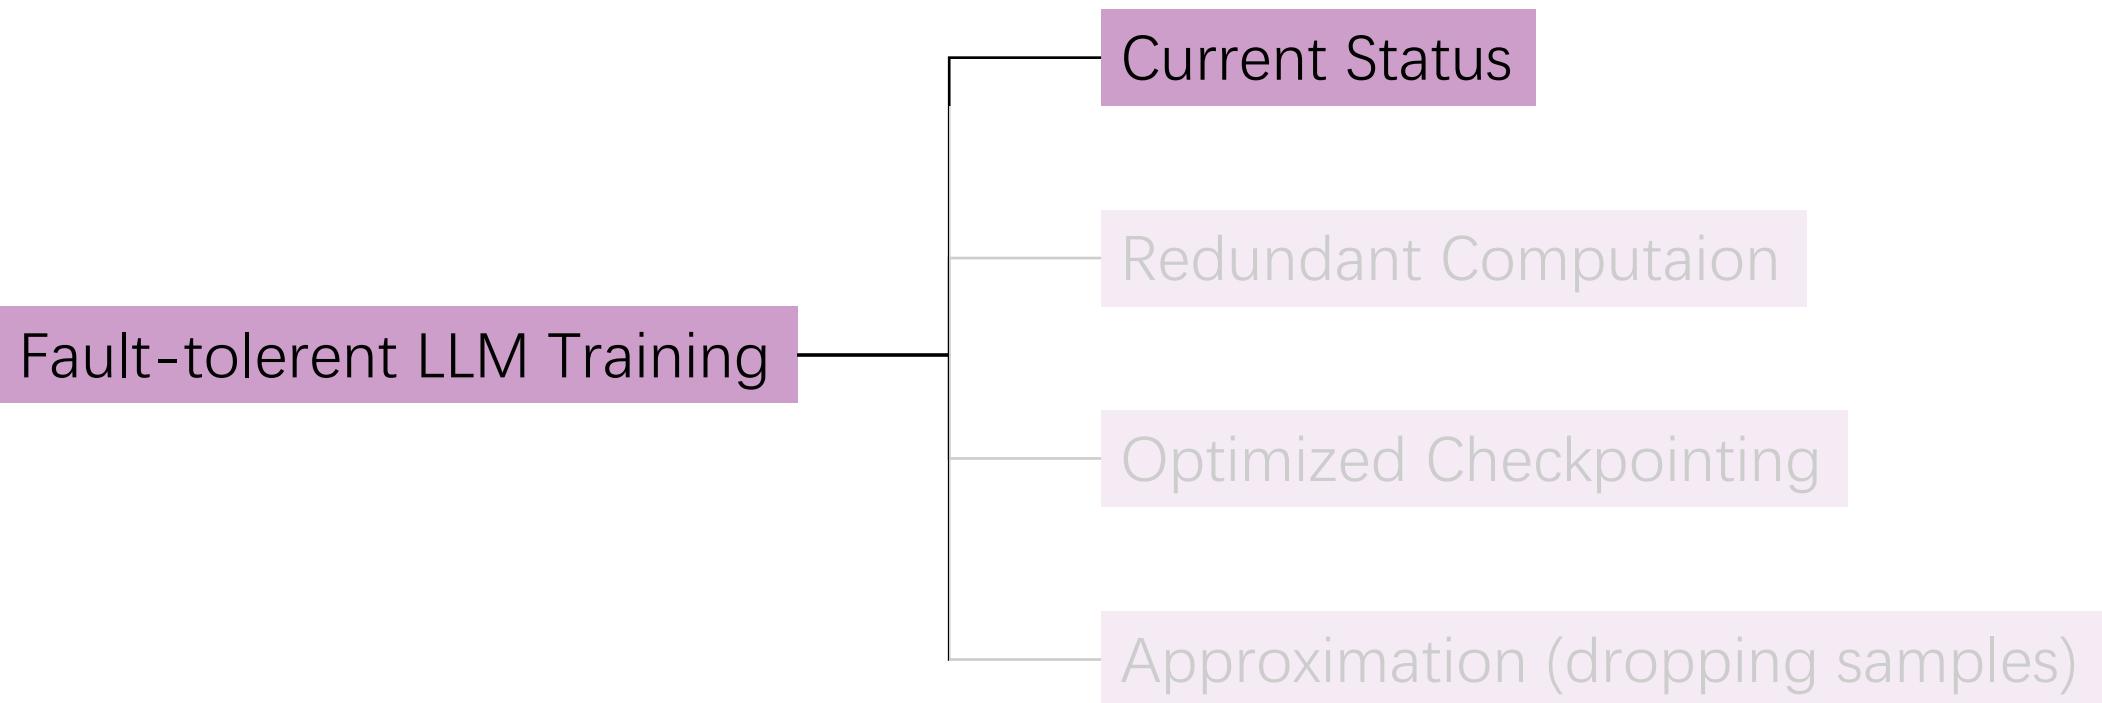


香港中文大學
The Chinese University of Hong Kong

Outline

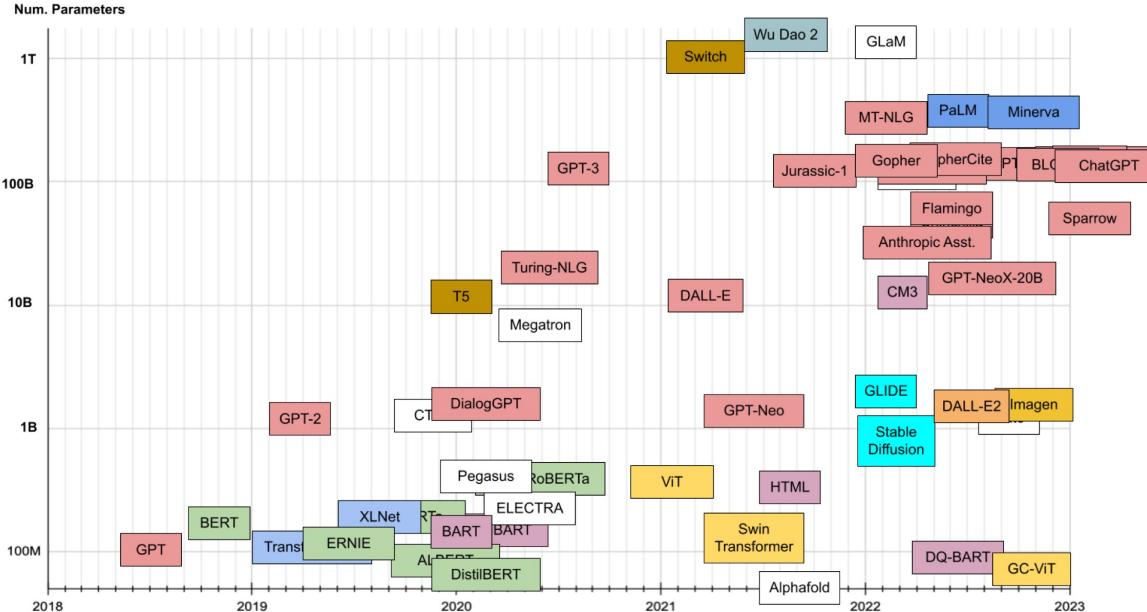


Outline





Why we need fault-tolerant LLM training?



The sizes of LLMs and training data are scaling up

Model	# A100-40GB	Throughput (tokens/sec)	Hours / 1B Tokens	Multi-node speedup	Price / 1B Tokens		
					Oracle [BM.GPU4.8]	AWS [p4d.24xlarge] on-demand	AWS [p4d.24xlarge] spot
GPT3-125M	8	191,693	1.45	1.00x	\$35.36	\$47.49	\$14.24
GPT3-125M	32	739,328	0.38	3.86x	\$36.67	\$49.25	\$14.77
GPT3-125M	128	2,471,936	0.11	12.90x	\$43.87	\$58.92	\$17.67
GPT3-350M	8	69,427	4.00	1.00x	\$97.62	\$131.11	\$39.33
GPT3-350M	32	269,926	1.03	3.89x	\$100.44	\$134.89	\$40.46
GPT3-350M	128	958,464	0.29	13.81x	\$113.14	\$151.96	\$45.58
GPT3-760M	8	57,549	4.83	1.00x	\$117.77	\$158.17	\$47.45
GPT3-760M	32	218,317	1.27	3.79x	\$124.18	\$166.78	\$50.03
GPT3-760M	128	696,320	0.40	12.10x	\$155.74	\$209.16	\$62.74
GPT3-1.3B	8	38,912	7.14	1.00x	\$174.18	\$233.93	\$70.17
GPT3-1.3B	32	152,781	1.82	3.93x	\$177.45	\$238.32	\$71.49
GPT3-1.3B	128	561,152	0.50	14.42x	\$193.25	\$259.55	\$77.86

The training machines and time are also increasing

the computational demand for large-scale AI models doubles approximately every 10 months



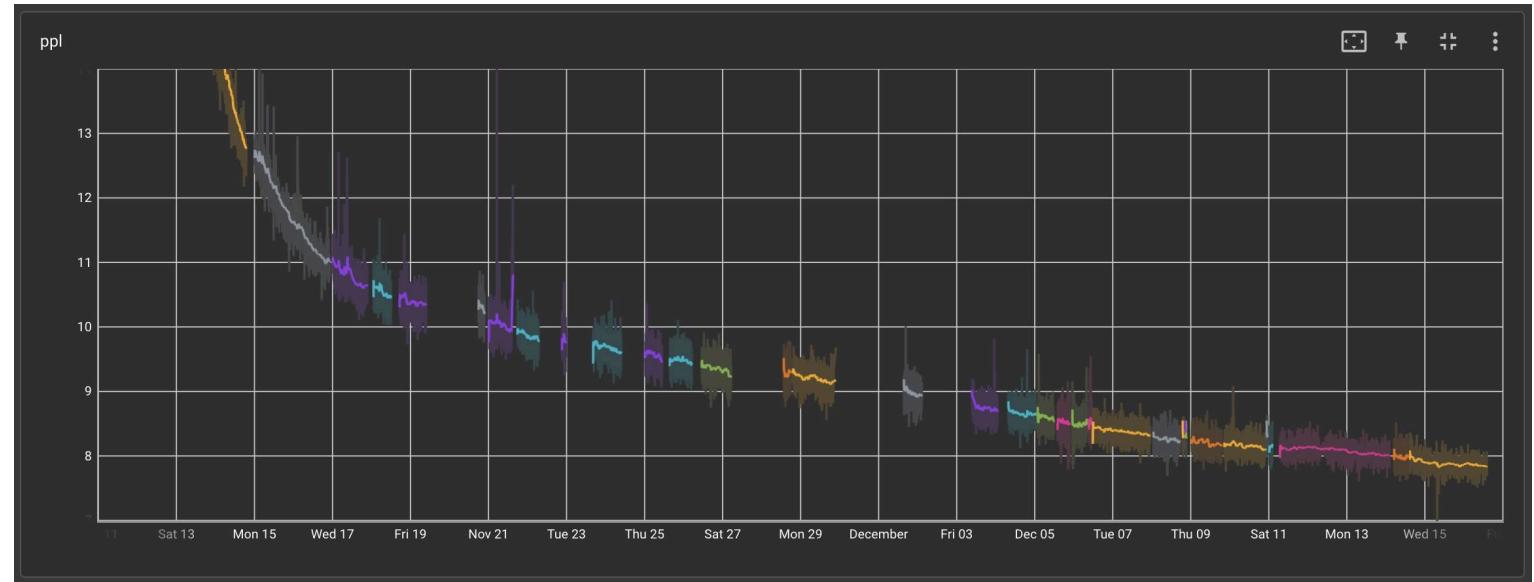
Why we need fault-tolerant LLM training?

- As the size of the model parameters and training machines increases, the probability of failures during training also significantly increases.

35+ manual restarts

70+ automatic restarts

100+ cycling hosts



All in all, working around **infrastructure issues** has dominated the last two weeks of the team's time, given that these hardware issues can take the experiment down for hours at any time of the day.

Since the sleepless night of Thanksgiving break, this past week has been filled with gradient overflow errors / loss scale hitting the minimum threshold (which was put in place to avoid underflowing) which also causes training to grind to a halt. We restarted from previous checkpoints a couple of times.



Why we need fault-tolerant LLM training?

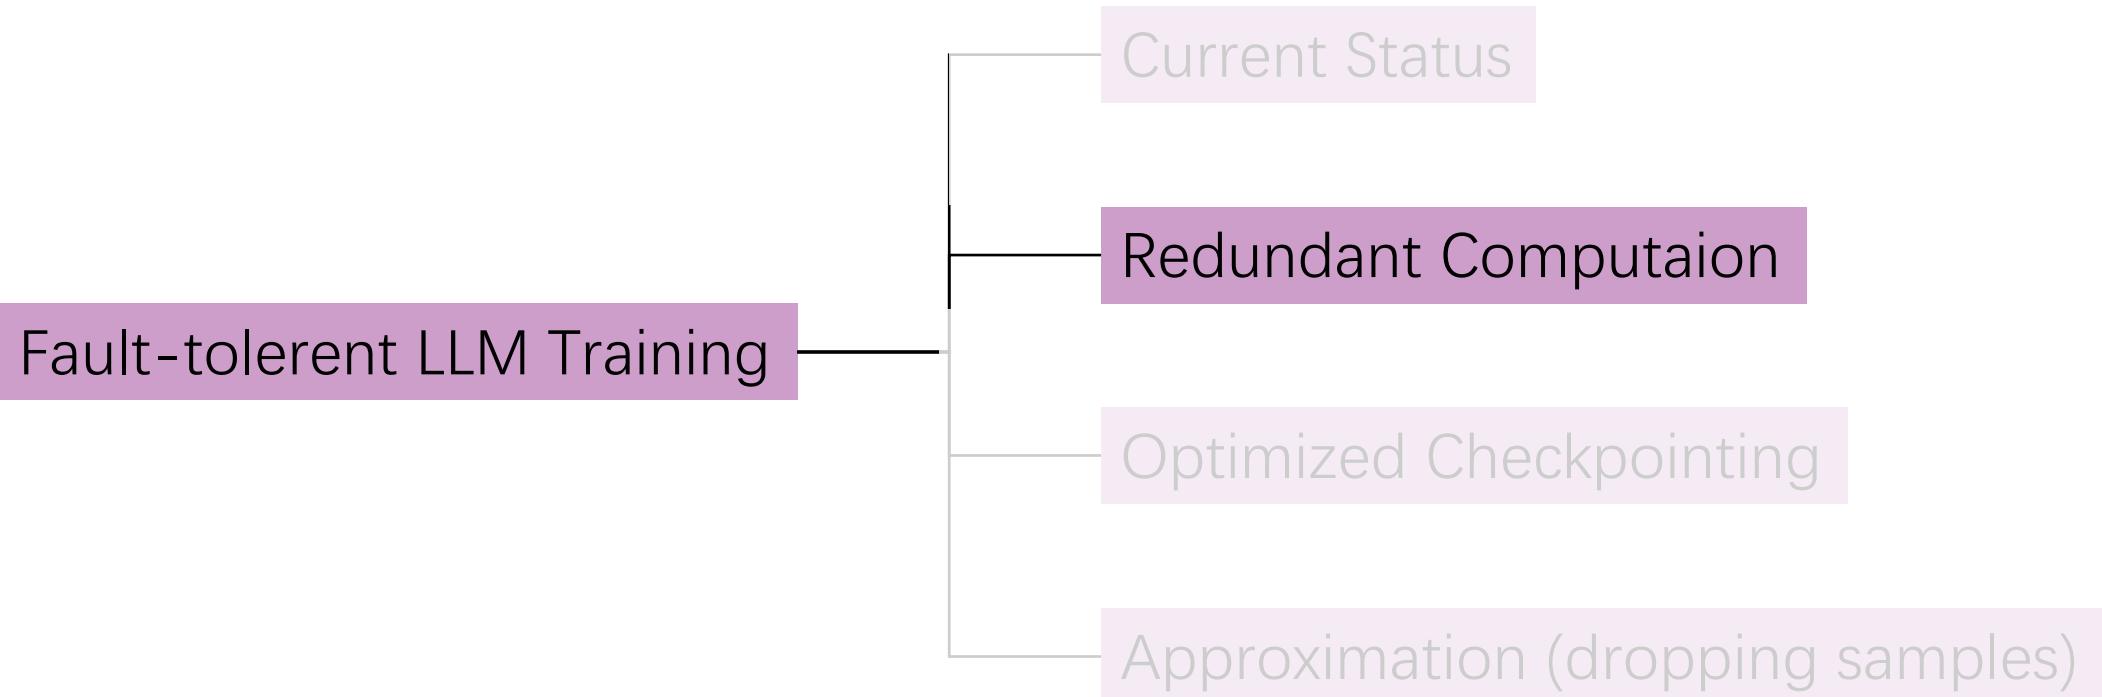
LLMs Training Task Error Statistics (From May 2023 to July 2023, running on SenseCore cluster)

Error Categorization	Number of Tasks	Root Cause
Storage Read/Write Errors	34	Due to synchronization anomalies of storage servers, significant variations in the time overhead for file storage or object storage occur across different nodes. This leads to communication waiting timeout or socket timeout in tasks.
Network Communication Errors	43	Incorrect insertion of IB network card; Uneven load distribution in RDMA traffic; Misconfigured RoCE network switch; Expired ARP cache IP; Ethernet card or link is not connected.
Node Hardware and Software Errors	66	GPU ECC errors; GPU failure; Node not ready; Insufficient shared memory; Pod sends SIGTERM signal to exit; Pod OOMKilled; Node image pull timeout; Local storage exceeding limit error.
User Code and Training Environment Errors	179	Error in creating duplicate files with the same name; Data type conversion error; Python Segmentation fault; CUDA runtime version mismatch with CUDA driver; AttributeError; torch.cuda.OutOfMemoryError; RuntimeError; ModuleNotFoundError; NameError; AssertionError; OSError.
Others	55	System hang without error output; Occasional socket timeout errors without specific issues identified during troubleshooting; Random occurrences of Pod processes with exit code -9 errors; Pod startup failures.

Approach	Fault Tolerant	Anomaly Detection	Checkpoint Optimization
DeepSpeed	✗	✗	✓
PyTorch Elastic	✓	✗	✗
Horovod Elastic	✓	✗	✓
Singularity	✓	✗	✓
DeepFreeze	✗	✗	✓
PAI	✓	✗	✗
ModelArts	✓	✓	✗
Azure	✗	✗	✓
TRANSOM	✓	✓	✓

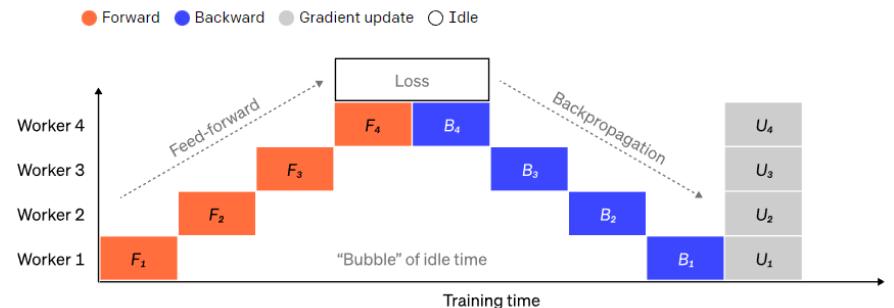
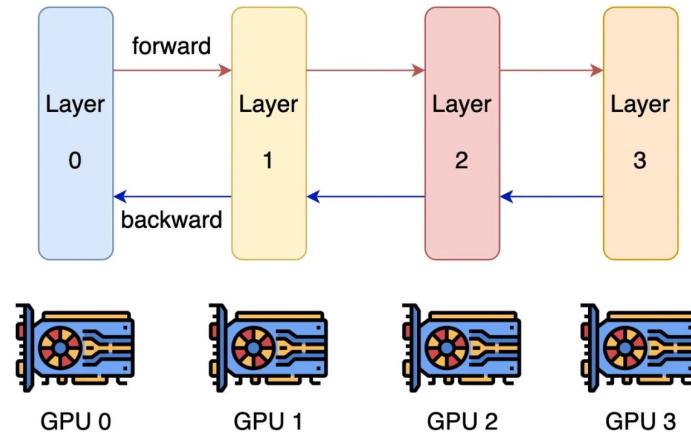
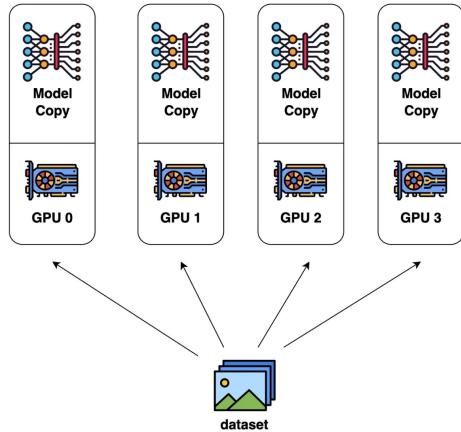
Attention to the fault tolerance of large language models is increasing!

Outline





Parallel Training in large-scale DNNs



Data Parallelism

$$C = A \times B$$

Non-distributed

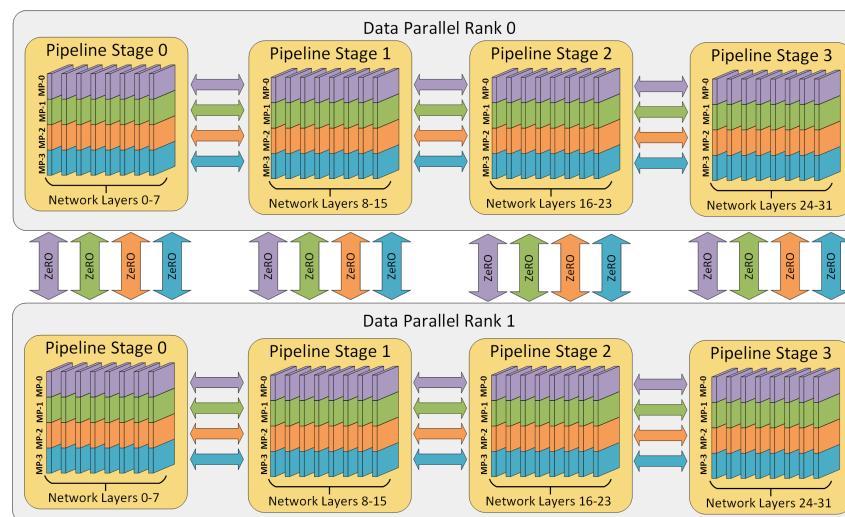
$$C = A \times B$$

all-gather along column

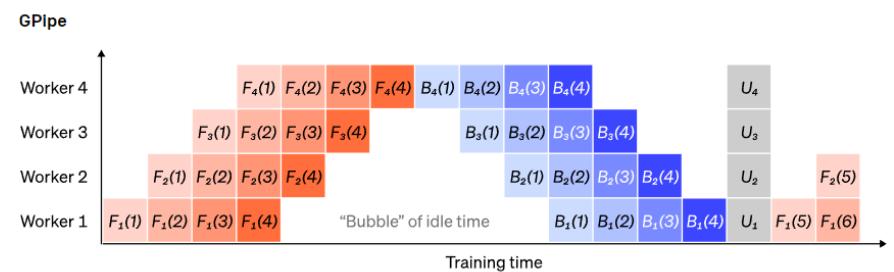
$$C = A \times B$$

Column-Splitting Tensor Parallel

Pipeline Parallelism



Tensor Parallelism



PipeDream

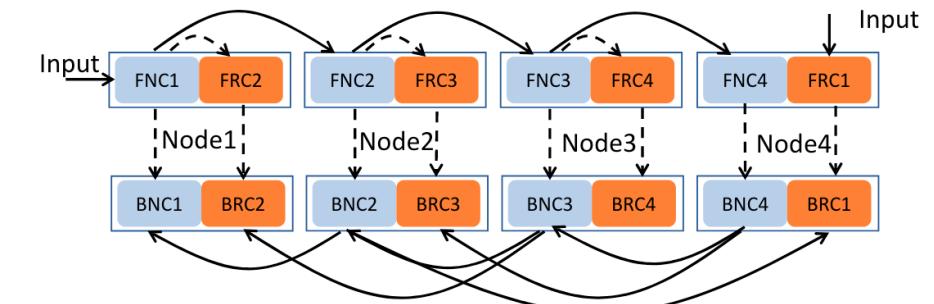
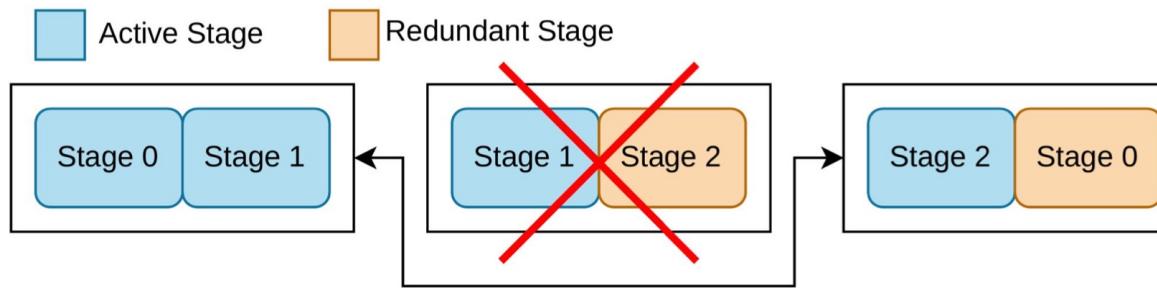
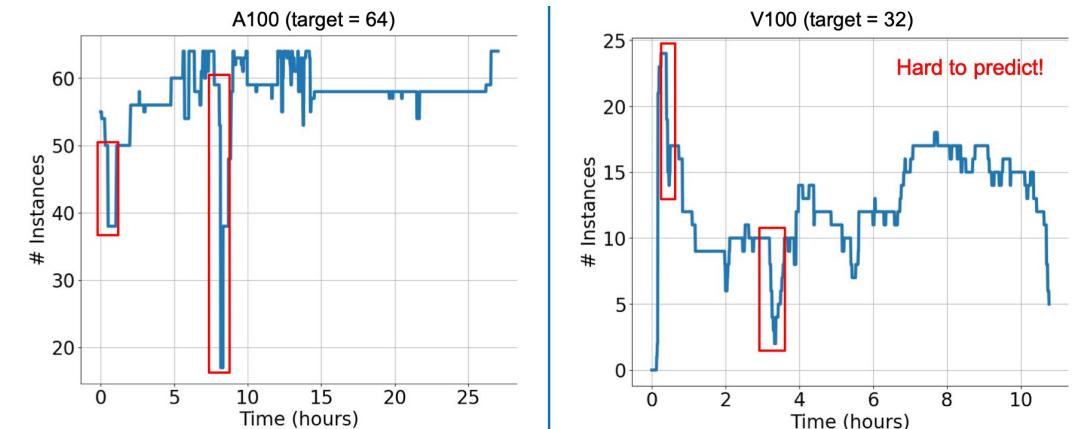


Optimized Pipeline



Bamboo: pipeline redundant computation

- Model sizes are increasing, the cost of training is higher
- Spot instances can lower costs, but have high failure rates
 - up to 70% cheaper
 - Preemptions can be unavoidable
- Redundant Computation (RC)
 - inspired by disk redundancies such as RAID
 - each node carries its own shard of layers and its successor's shard (exploit locality)

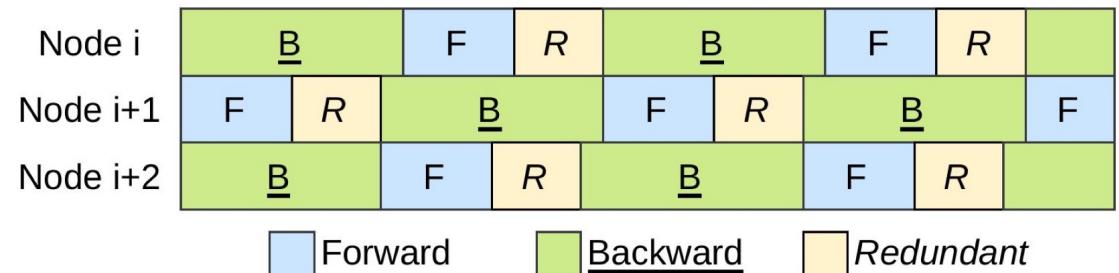




Bamboo: pipeline redundant computation

Challenge 1: high overhead

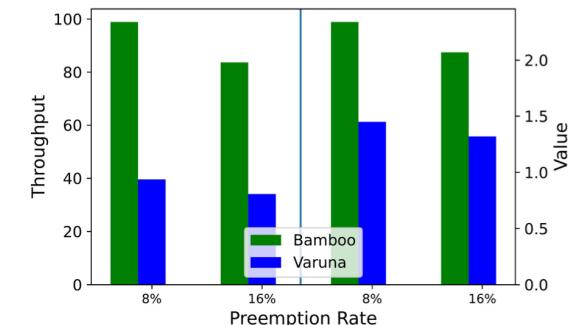
- pipeline parallelism has bubbles
- use this idle time to minimize redundancy overhead



Challenge 2: high GPU memory usage

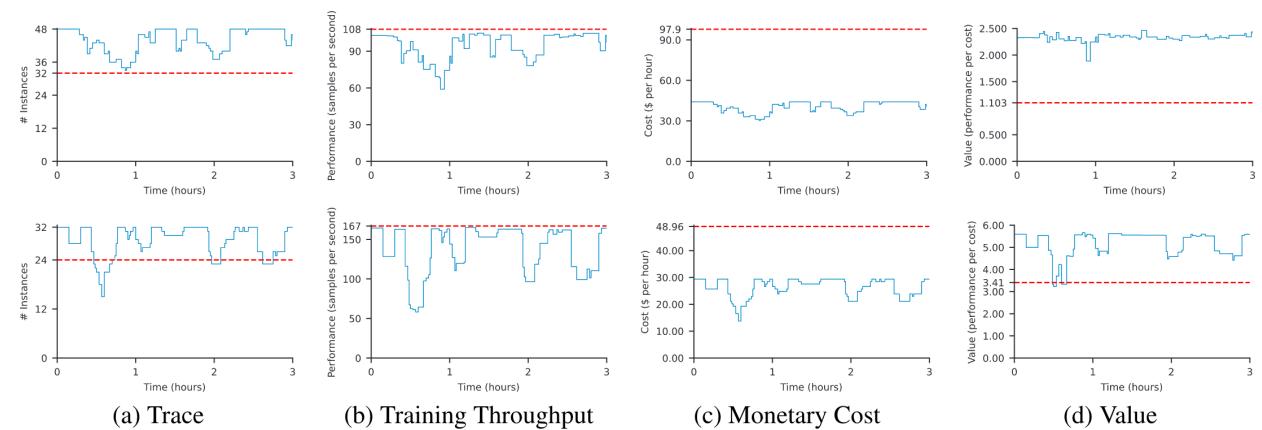
- swap out the intermediate results of each node's FRC into the node's CPU memory (offload less frequently used tensors to CPU memory)

Model	Dataset	Samples	D	P
ResNet-152 [22]	ImageNet [32]	300,000	4	12
VGG-19 [63]	ImageNet [32]	1,000,000	4	6
AlexNet [32]	ImageNet [32]	1,000,000	4	6
GNMT-16 [68]	WMT16 EN-De	200,000	4	6
BERT-Large [15]	Wikicorpus En [15]	2,500,000	4	12
GPT-2 [49]	Wikicorpus En [15]	500,000	4	12



Challenge 3: consecutive failure

- make consecutive nodes in each pipeline come from different zones



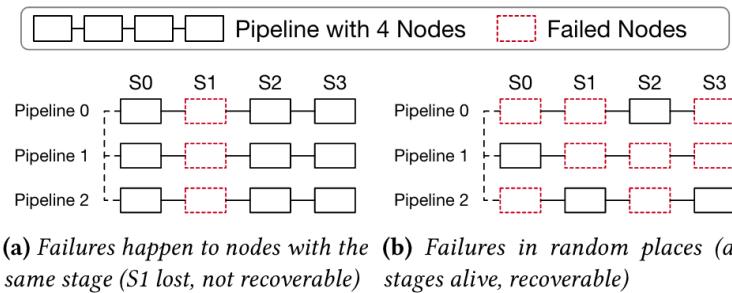
compared to on-demand instances



Oobleck: scheduling pipeline templates

- Scenarios: **hybrid-parallel training (DP + PP)**
- Model state redundancy in DP is free redundancy

Can we leverage this redundancy in DP for reliability?

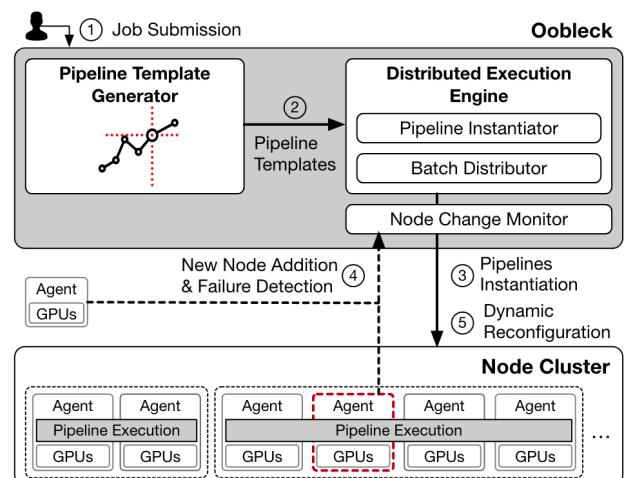
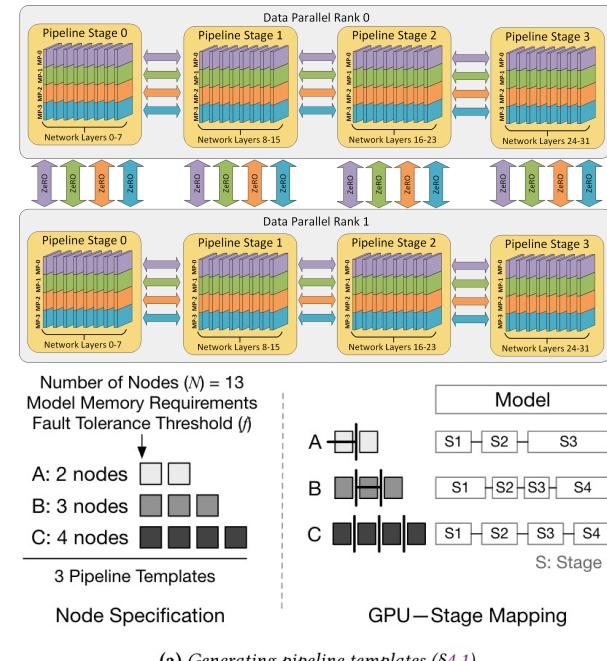


(a) Failures happen to nodes with the same stage (S1 lost, not recoverable)
(b) Failures in random places (all stages alive, recoverable)

The core design: **pipeline templates**

- a pipeline specification that defines how many nodes should be assigned to a pipeline, how many stages to create, and how to map model layers in stages to GPUs.

Decoupling planning (pipeline template generation) from execution (pipeline instantiation) enables fast failure recovery;



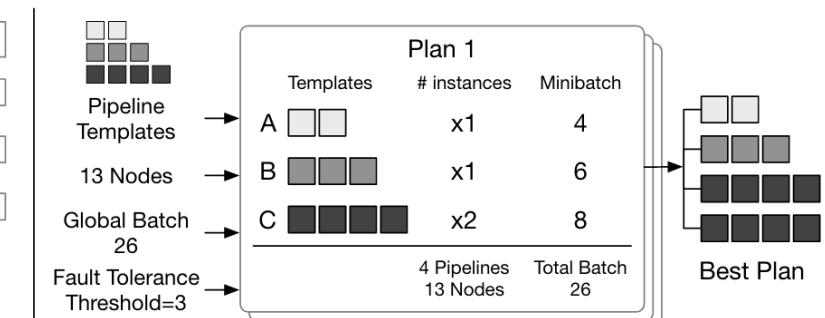
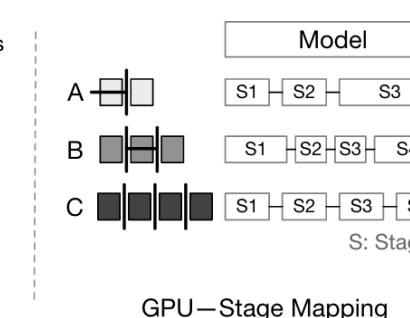
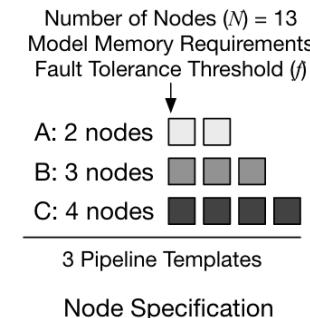


Oobleck: scheduling pipeline templates

Oobleck Planning Algorithm

- Generating pipeline templates

- Node specification
- GPU–Stage mapping



(a) Generating pipeline templates (§4.1)

- Pipeline instantiation

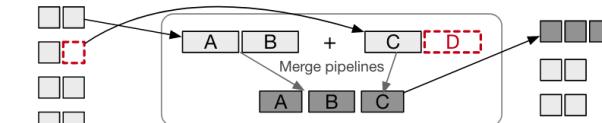
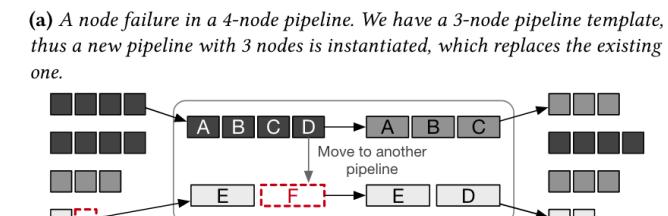
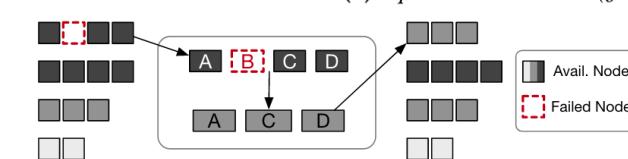
- Enumerating all instantiation options
- Calculating throughput with batch distribution

- Dynamic Reconfiguration

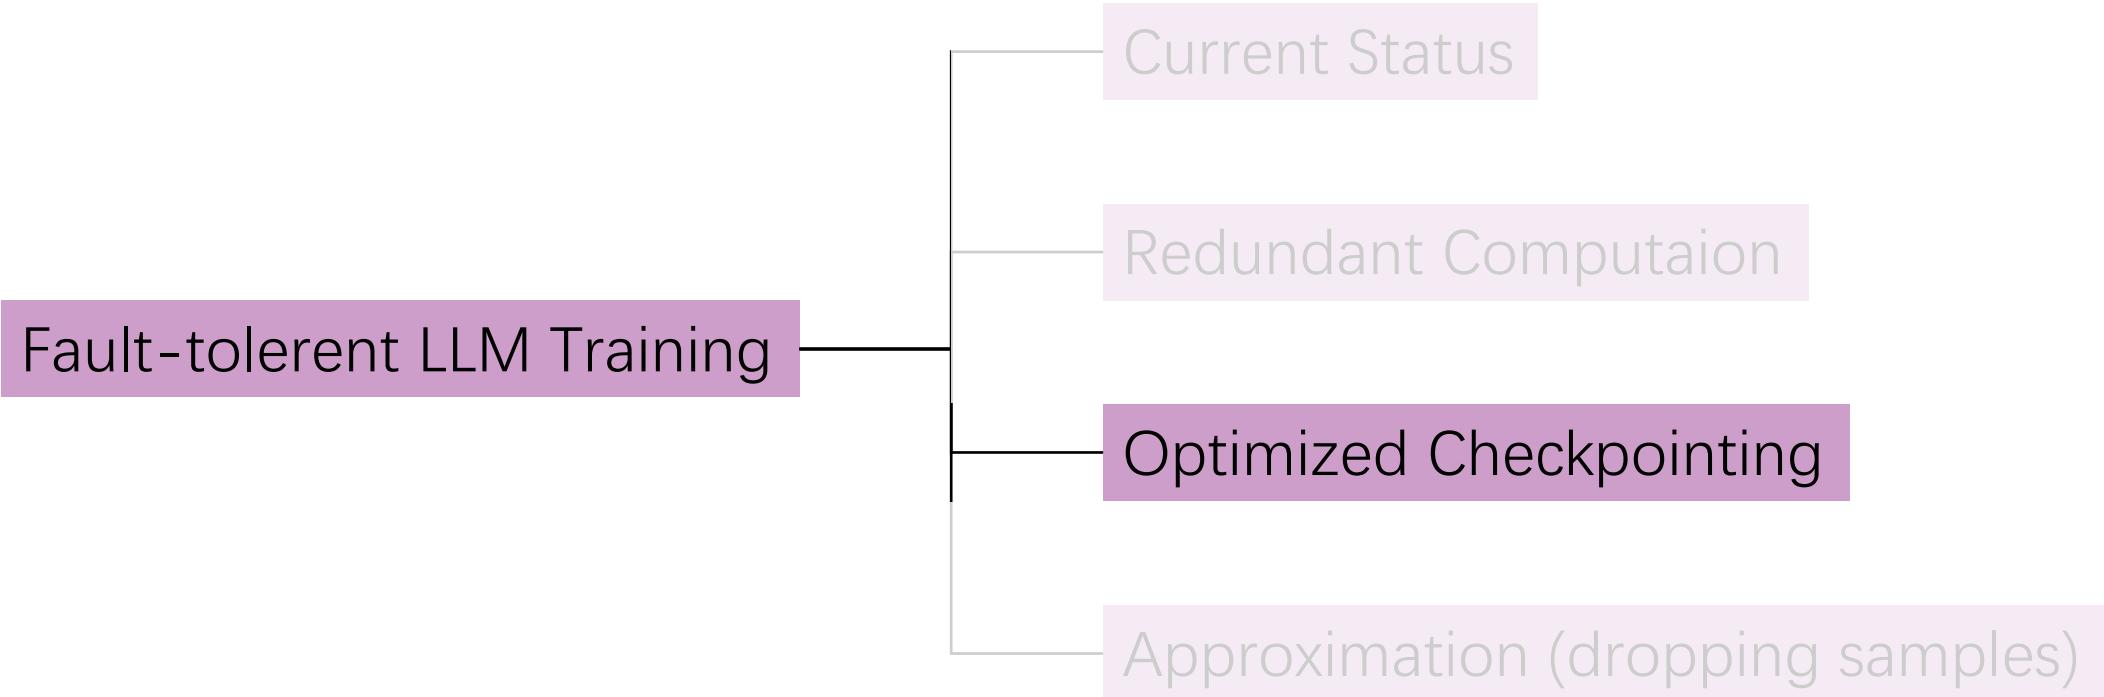
- Pipeline re-instantiation
- Batch redistribution

- Other designs

- Model synchronization between heterogeneous pipelines in a layer granularity



Outline



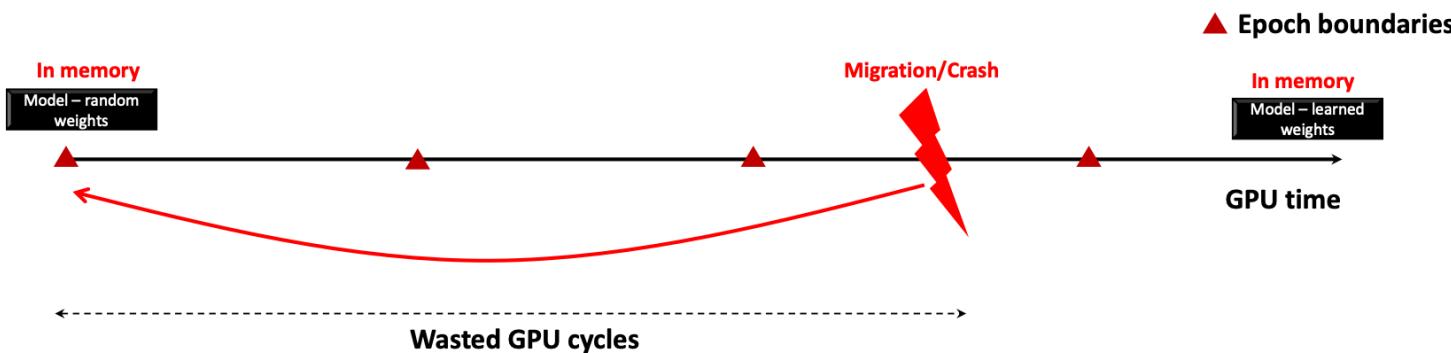


CheckFreq: fine-grained and pipelined checkpointing

Three main challenges in large-scale DNNs checkpointing:

- **Checkpoint stalls:** how to minimize cost of checkpoint?
- **Checkpointing frequency:** how often to checkpoint?
- **Data invariant:** how to resume correctly?

Any interruption can wipe out the model parameters learned so far in memory,
restarting this expensive process!



Checkpoint of a 128B LLM:
~ 1.536 TB

Trade-off between low-overhead and high frequency of checkpointing

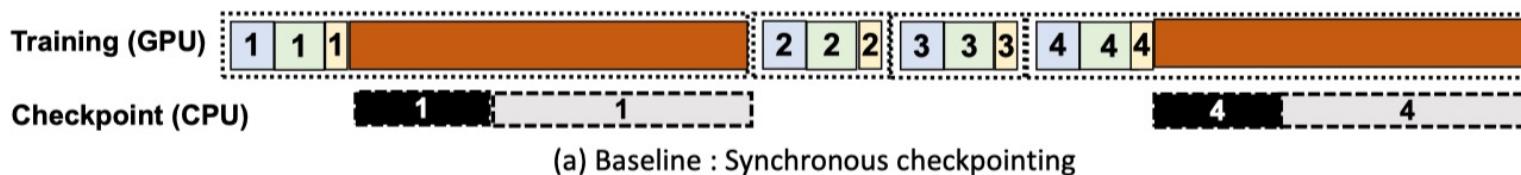


CheckFreq: fine-grained and pipelined checkpointing

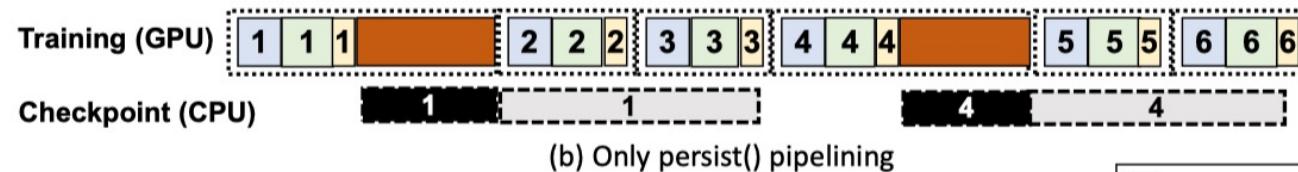
Checkpoint stalls: how to minimize cost of checkpoint?

2-phase DNN-aware checkpointing
Low checkpoint stalls

- Synchronous checkpointing introduces checkpoint stalls => Runtime overhead
- Low-cost checkpointing mechanism that is split into two pipelined operations:
 - Snapshot() : Serialize and copy into a user-space buffer
 - Persist() : Write out the serialized contents to disk

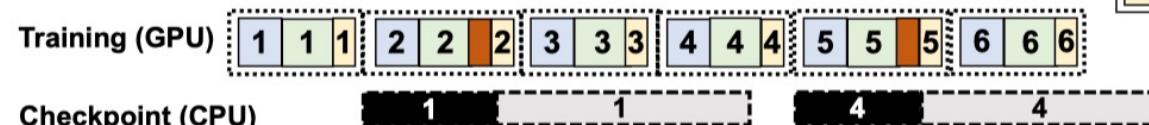


(a) Baseline : Synchronous checkpointing



(b) Only persist() pipelining

Forward pass	Snapshotting
Backward pass	Disk IO
Weight update	Checkpoint stall



(c) Snapshot() and persist() pipelining

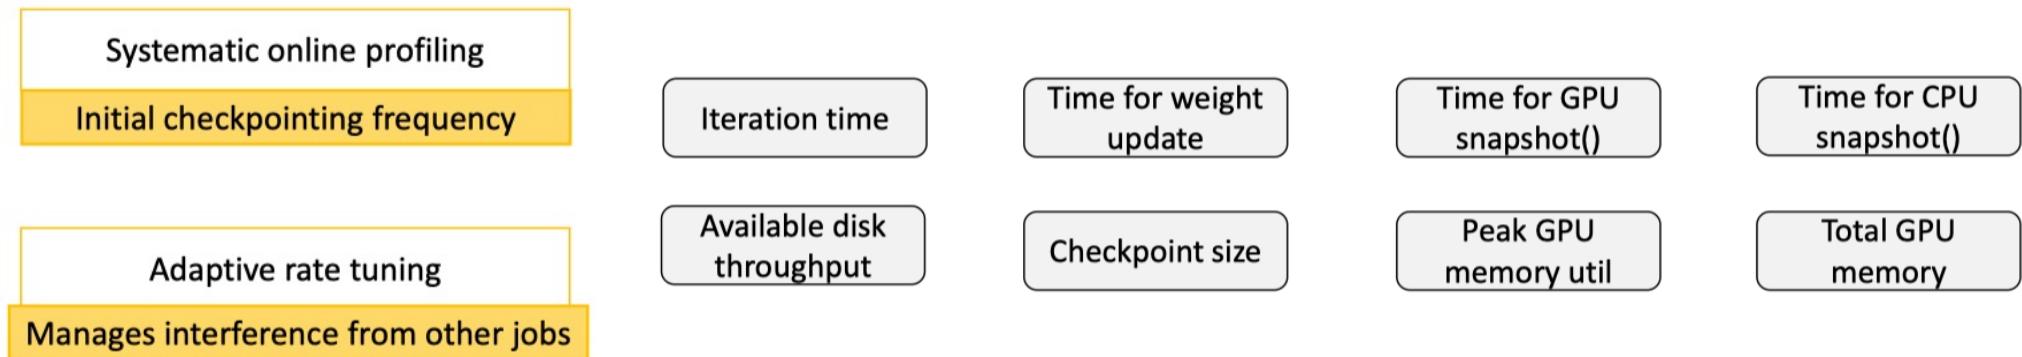


CheckFreq: fine-grained and pipelined checkpointing

Checkpointing frequency: how often to checkpoint?



automatically profiles metrics and determines frequency



Data invariant: how to resume correctly?



epoch seeded pseudo-random transformations

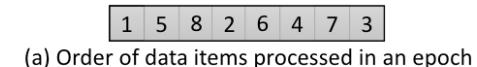
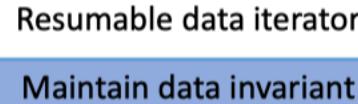


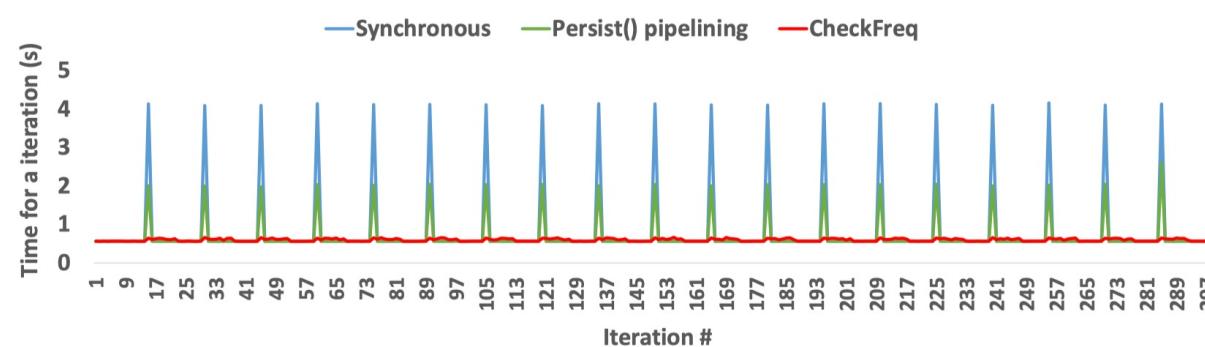
Figure 5: **Resuming iterator state.** When iterator state is not resumable, an epoch might miss data items when job is interrupted (items 3,6,7 are missed in b). CheckFreq (c) ensures that training resumes from exactly where it left off.



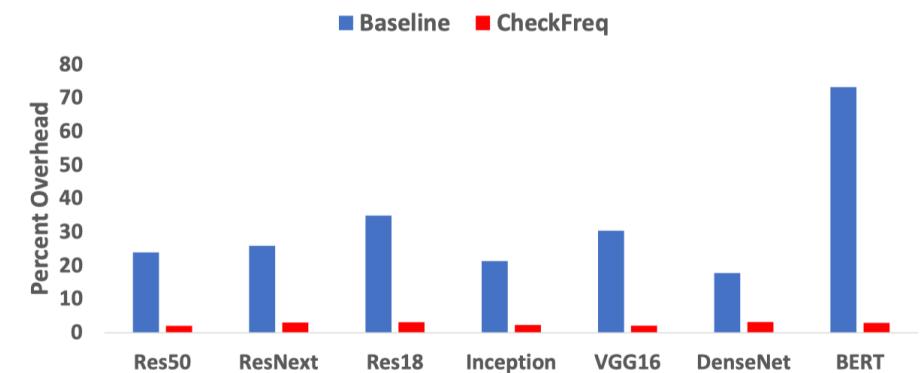
CheckFreq: fine-grained and pipelined checkpointing

Evaluation on 7 different DNNs:

1. CheckFreq reduces checkpoint stalls



2. CheckFreq reduces checkpoint overhead



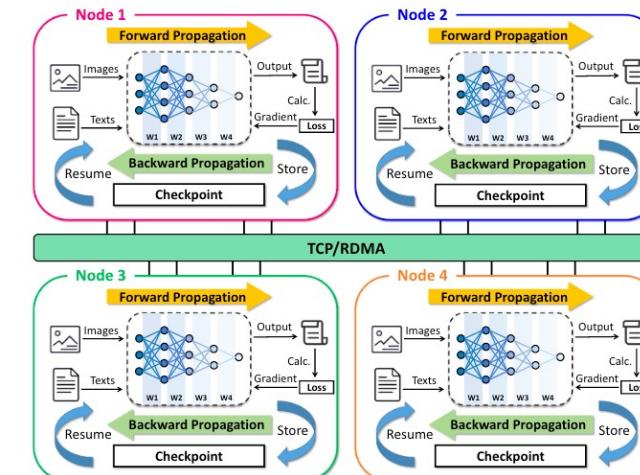
3. CheckFreq lowers recovery time: from hours to seconds

Model	Epoch-based (s)	CheckFreq (s)
Res18	840	5
Res50	2100	24
VGG16	5700	25
ResNext	7080	32
DenseNet	2340	7
Inception	3000	27
BERT	4920	85



LightCheck: pipelined checkpoint to PM

- Existing methods (e.g., CheckFreq) still cannot fully utilize the parallelism among computation, communication, and checkpointing



LightCheck pipelines checkpointing with comp. and comm. in a layer-wise way

- Leverage the data dependency.
- Multiple nodes synchronize the model parameters layer-by-layer during communication.

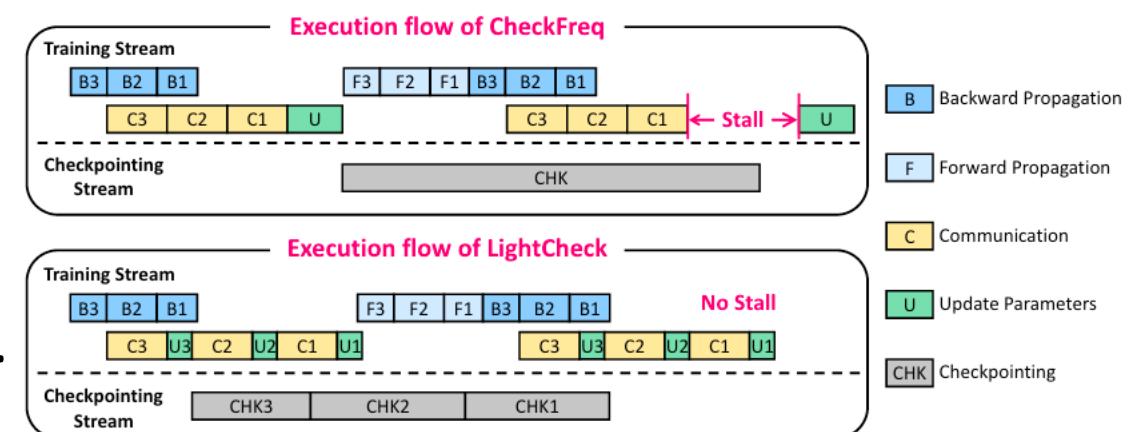


Fig. 3: The execution flows of LightCheck and CheckFreq.



LightCheck: pipelined checkpoint to PM

Persistent Memory (PM) has received extensive attention

- Large capacity with near-DRAM performance
- The unified virtual addressing (UVA) technique enables zero-copy access over PCIe between GPU and PM
- The asynchronous layer-wise checkpointing is done via CUDA streams and events

Evaluation

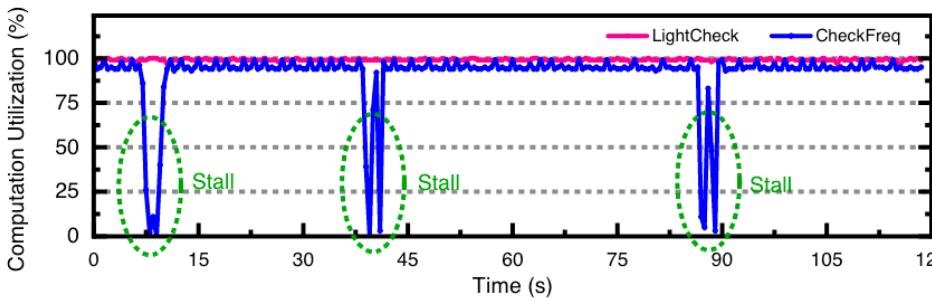
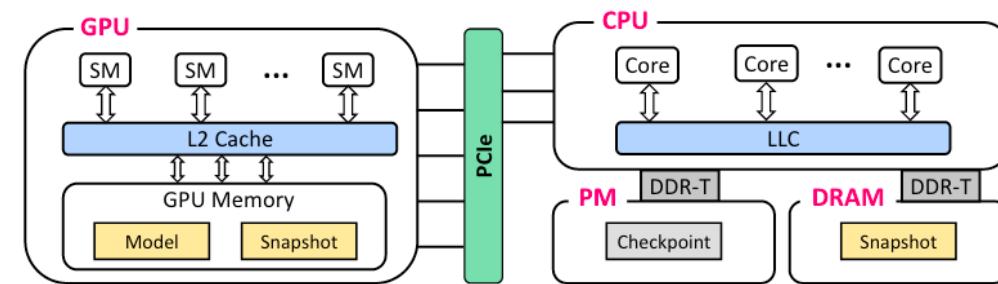
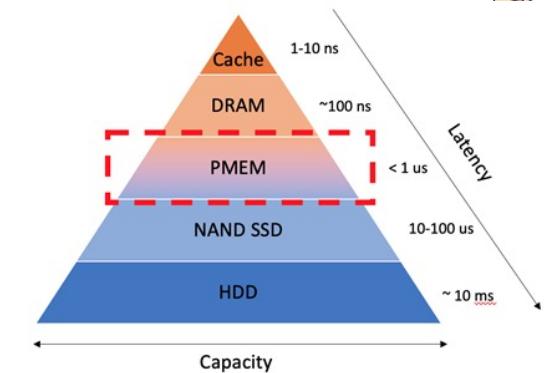


Fig. 8: The GPU computation utilizations of LightCheck and CheckFreq.



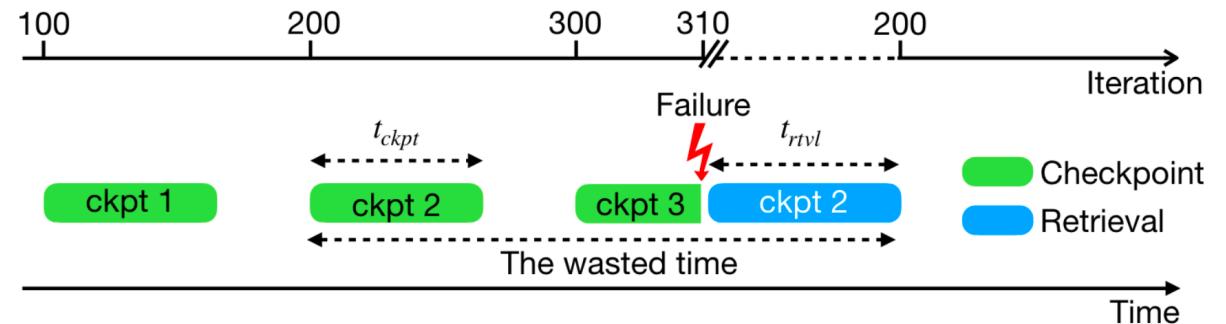
Models	Total Training Time (h)			
	No Failure	LightCheck	CheckFreq	<code>torch.save</code>
ResNet-18	10.7	10.9	11.1	11.2
VGG_16	67.5	69.7	72.3	73.7
Inception-V3	79.7	83.0	84.1	85.4
AlexNet	6.0	6.1	6.3	6.5
GPT-2	161.7	164.6	171.1	179.7
BERT	501.3	514.8	537.5	598.6



GEMINI: hierarchical storage of checkpoints

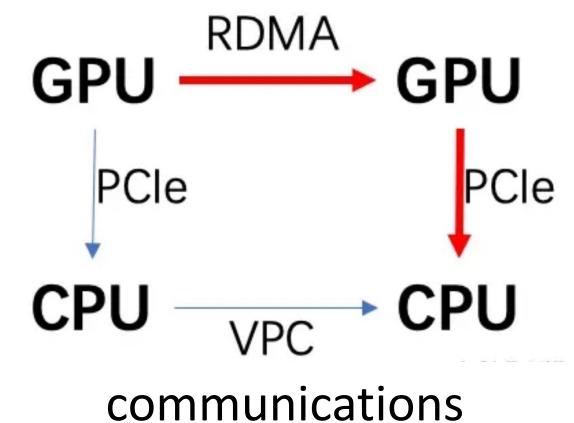
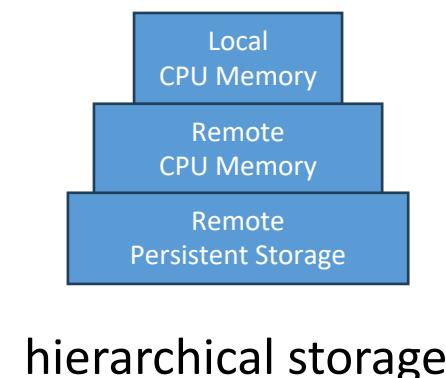
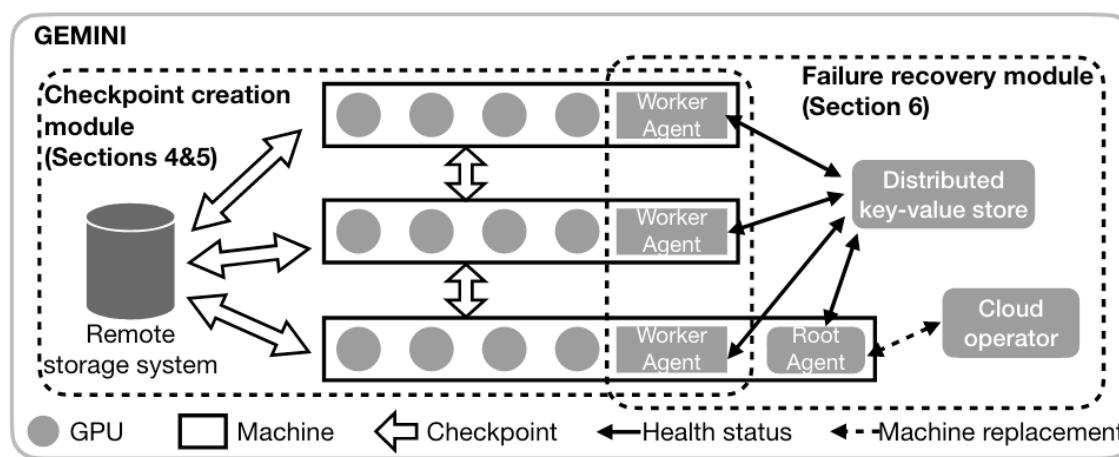
Three main factors in checkpoint:

- checkpoint time
- checkpoint frequency
- retrieval time



Leverage the high bandwidth of CPU memory to achieve fast failure recovery!

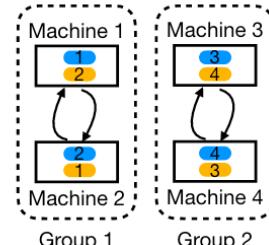
- How to maximize the probability of a successful failure recovery from CPU memory?
- how to minimize the interference of checkpoint traffic with model training?



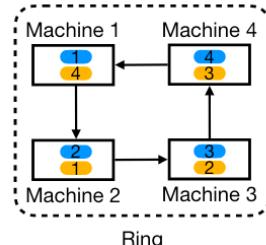


GEMINI: hierarchical storage of checkpoints

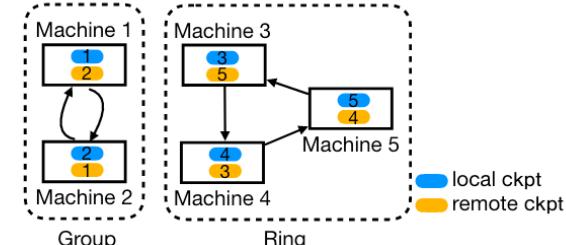
- Maximize the probability of a successful failure recovery from CPU memory
 - Store redundant checkpoints and proposes a placement strategy that maximizes the probability**



(a) Group placement strategy.



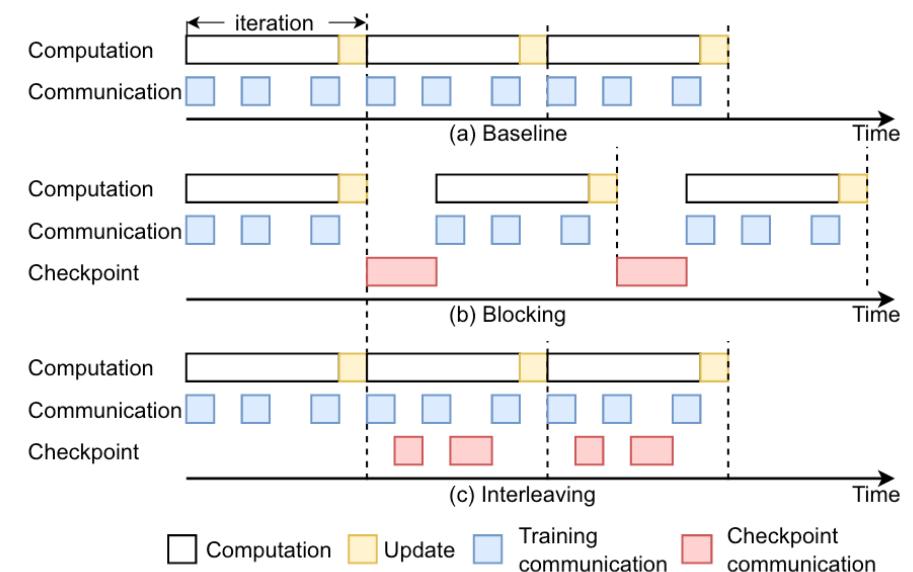
(b) Ring placement strategy.



(c) Mixed placement strategy.

- Minimize the interference of checkpoint traffic with model training?

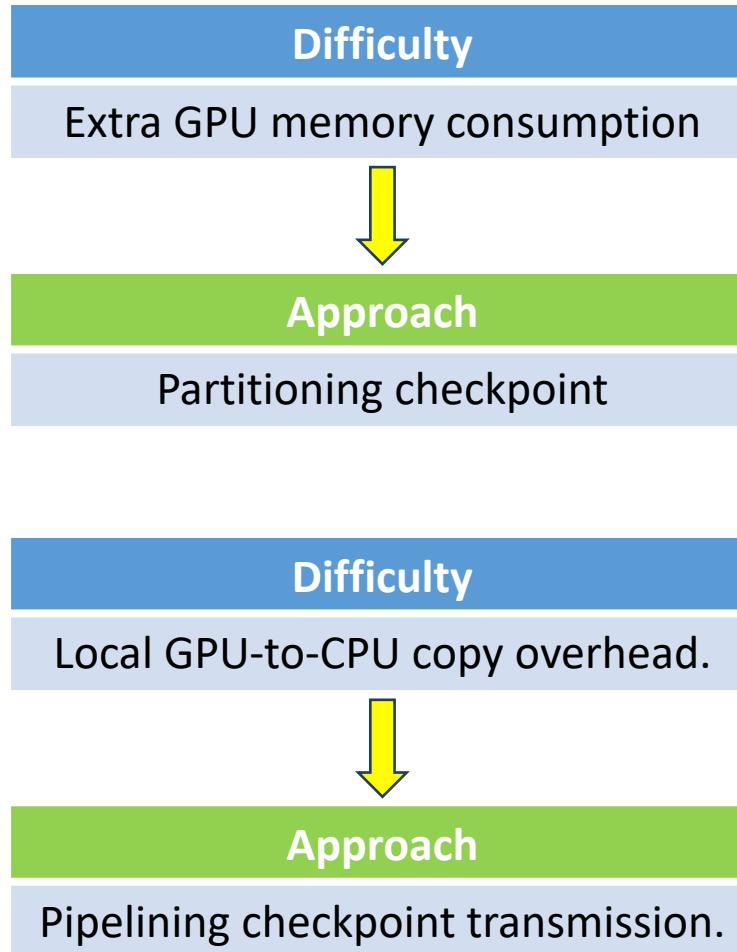
- Design a deliberate communication scheduling algorithm for interleaving these two types of traffic to minimize the interference on training throughput.**
- Method: online profiling for several iterations of training (e.g., 20) without checkpointing.



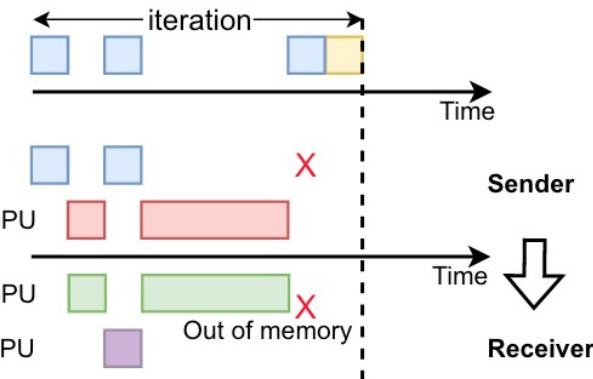


GEMINI: hierarchical storage of checkpoints

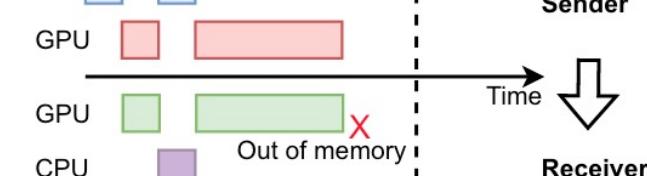
- Other difficulties and approaches



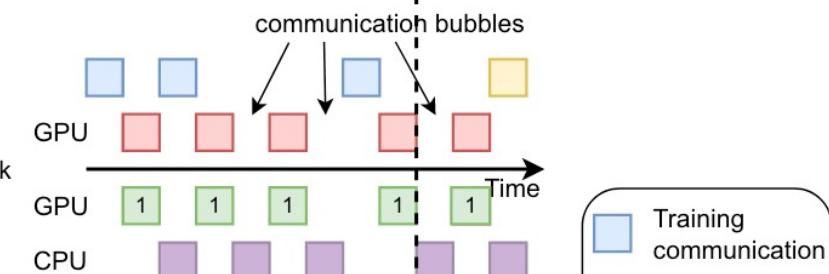
(a) Baseline without checkpointing. Only training communications and Update are displayed.



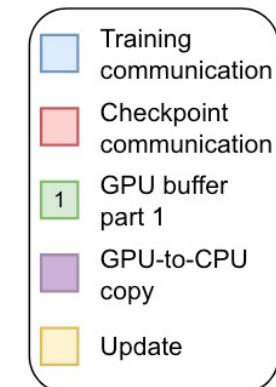
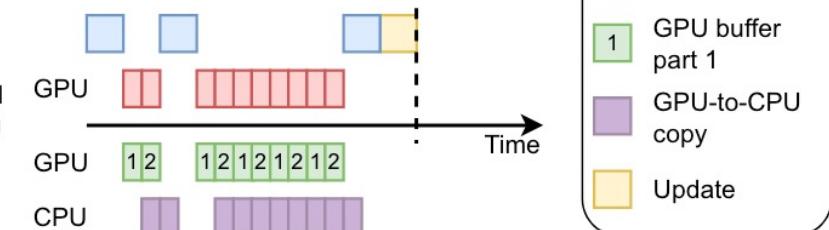
(b) The whole checkpoint is stored in GPU, causing OOM errors.



(c) Checkpoint is partitioned, but GPU-to-CPU copies block checkpoint communications.



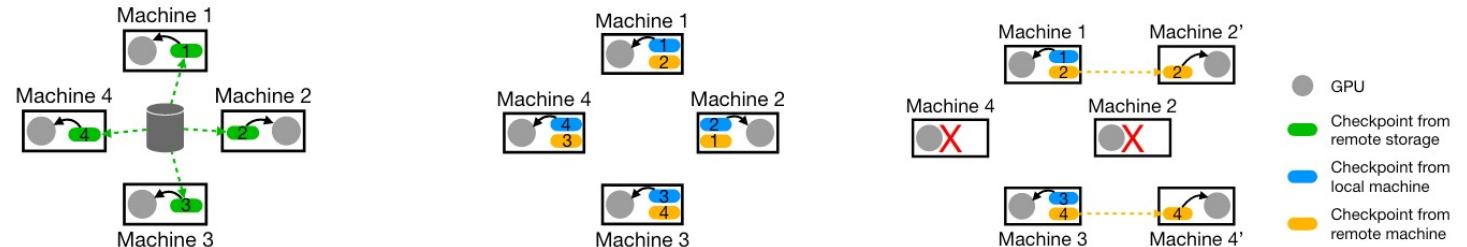
(d) The GPU buffer is splitted into two parts. Checkpointing to CPU memory is pipelined.





GEMINI: hierarchical storage of checkpoints

- Recovery training from failures
 - software failures
 - hardware failures
 - replicas exist
 - replicas not exist



(a) Existing solutions for any types of failures. All checkpoints are always retrieved from the remote persistent storage.

(b) GEMINI for software failures. Checkpoints are at local and the retrieval time is negligible.

(c) GEMINI for failures with two machines replaced. The newly added machines retrieve checkpoints from alive machines.

Evaluation

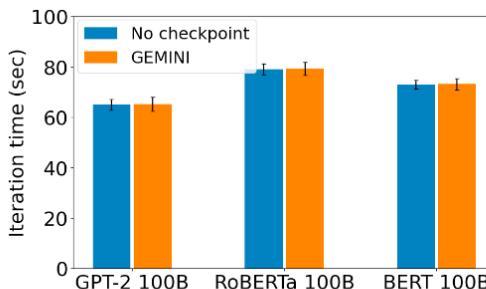


Figure 7. The iteration time of three large models without checkpoints and with GEMINI.

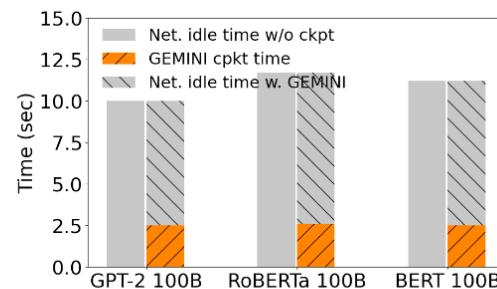
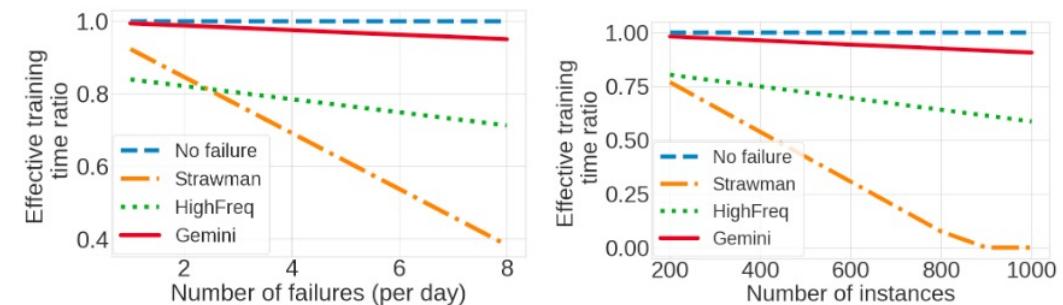
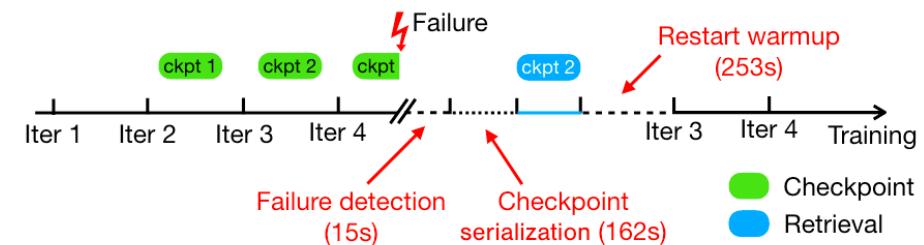


Figure 8. The network idle time of three large models without checkpoints and with GEMINI.



(a) Different failure rates. (b) Different instance numbers.





➤ SWIFT: update-undo and logging-based recovery

- The overhead of checkpointing is high
- **Crash-consistency problem:** parameters from some workers are updated while the others are not.

Swift Design:

- **Update-undo** (replica available):
 - survivors undo the update for the updated parameters
- **Logging-based recovery** (replica unavailable):
 - is done asynchronously
 - records intermediate communication (i.e., intermediate activations in the forward pass and the gradients in the backward pass)
 - replacement workers download logs for recovery

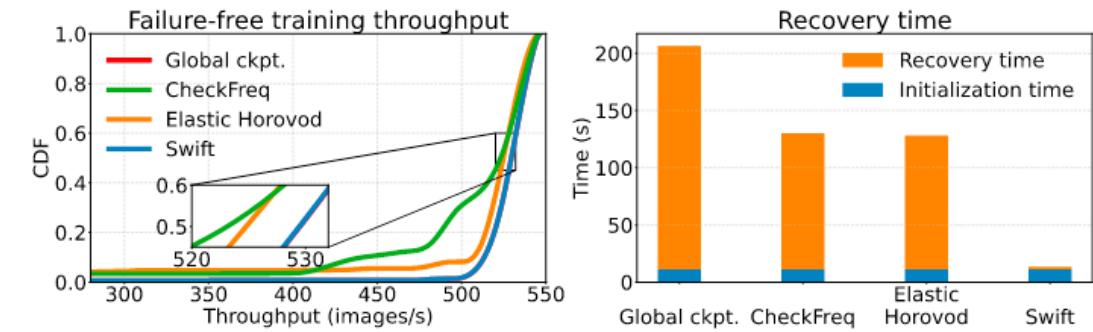


Figure 1. Replication-based recovery for Wide-ResNet-50.

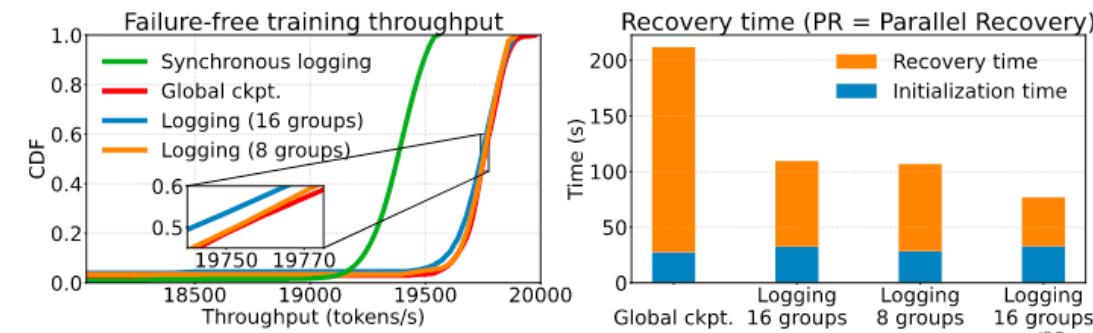
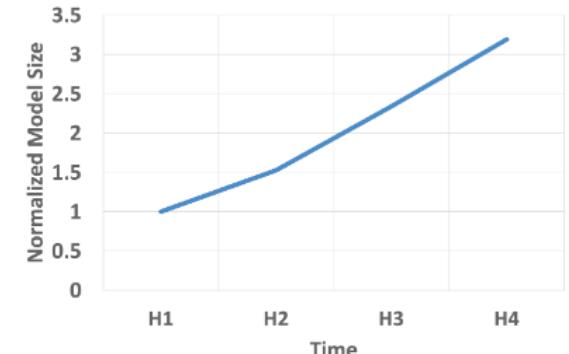
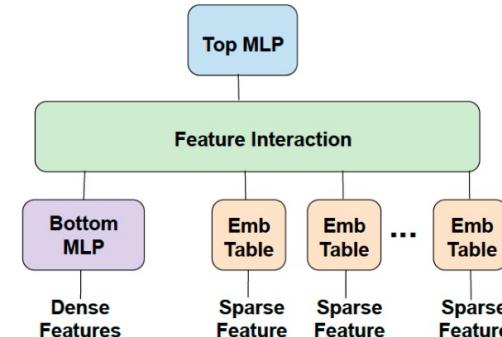


Figure 2. Logging-based recovery for BERT-128.

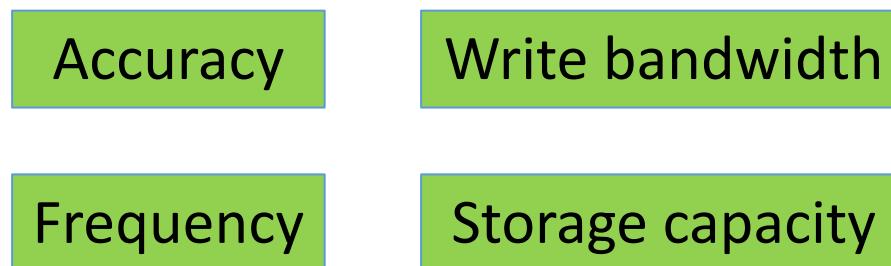


Check-N-Run: reduce checkpoint size

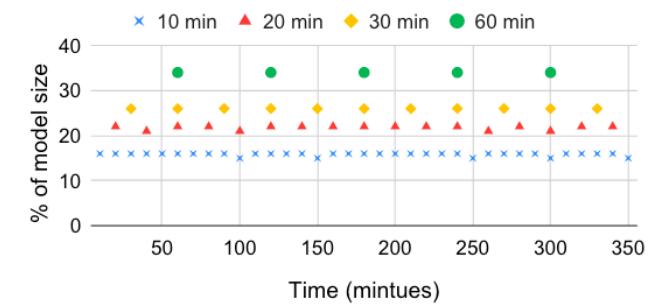
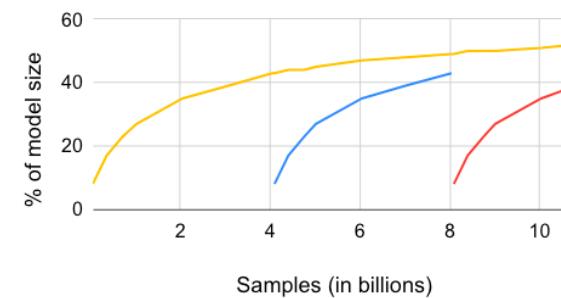
- Checkpoints of large-scale recommendation system (RS) is essential



- Challenges of checkpointing



Recommendation Model Architecture

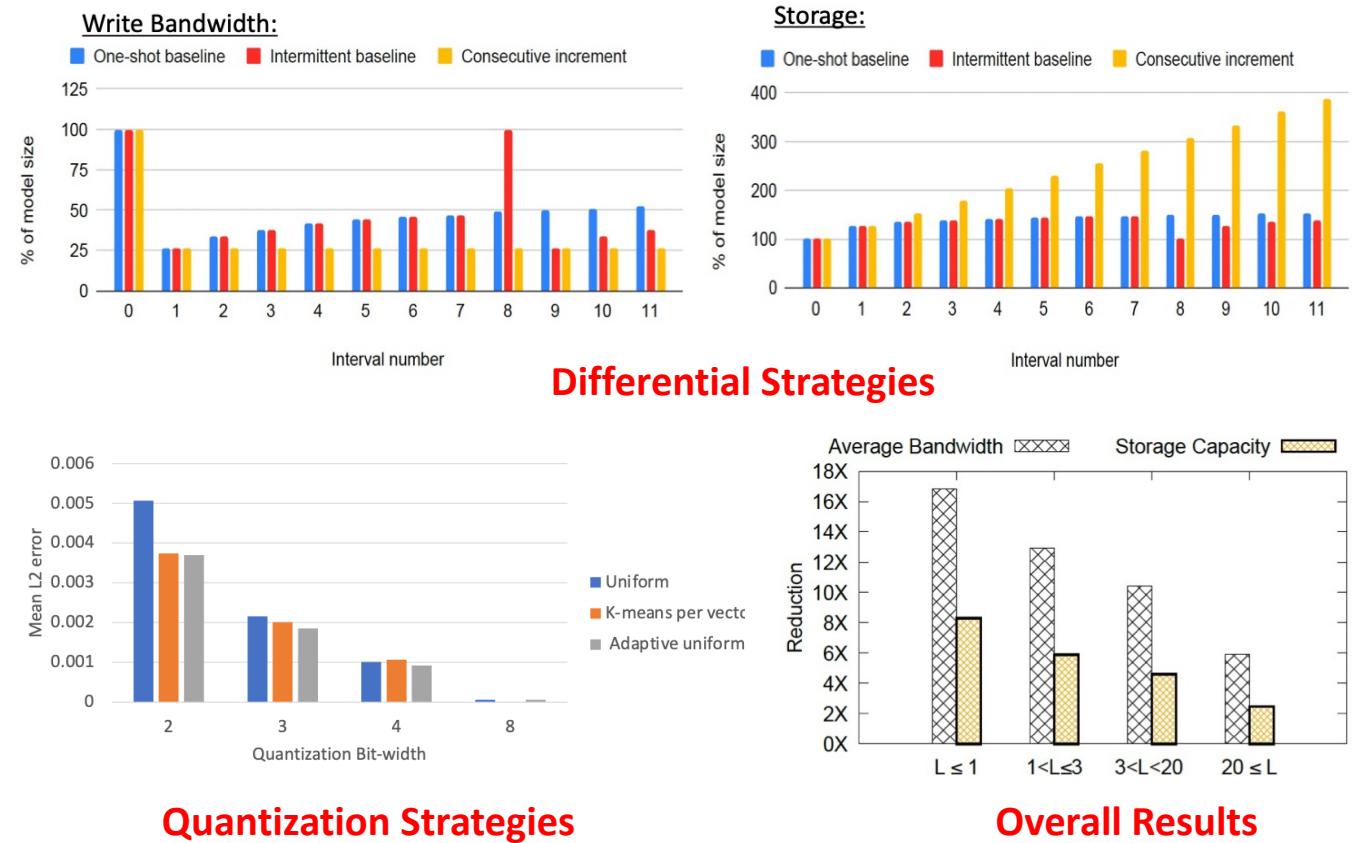
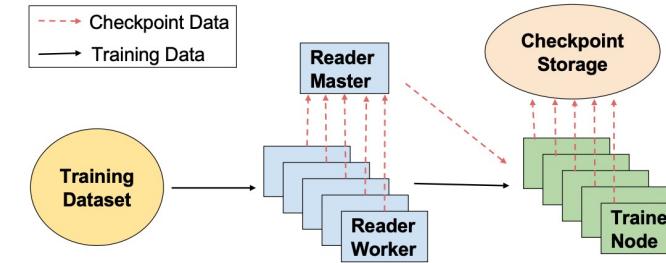


at each iteration only a tiny fraction of the model is updated

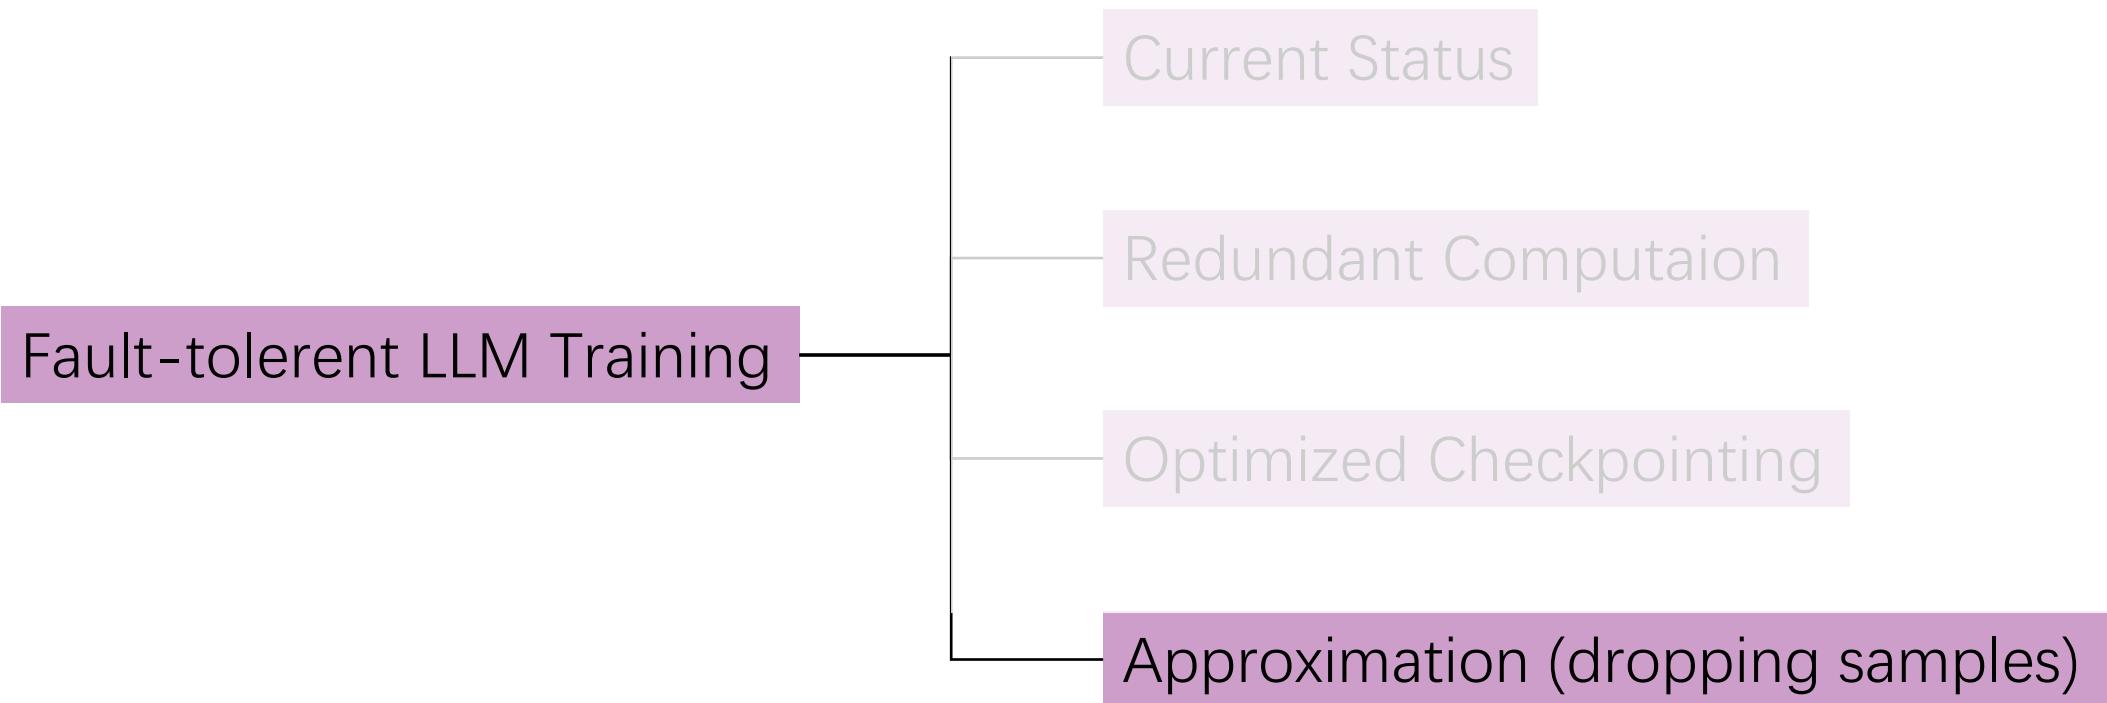


Check-N-Run: reduce checkpoint size

- Strategy 1: differential checkpointing
 - Motivation: model accesses are sparse
 - One-Shot Differential Checkpoint
 - Consecutive Incremental Checkpoint
 - Intermittent Differential Checkpoint
- Strategy 2: checkpoint quantization
 - Compress checkpoint without degrading training accuracy
 - Uniform quantization
 - Non-uniform quantization using k-means
 - Adaptive uniform quantization
- Strategy 3: Decoupling
 - Separate snapshot and persist operations.



Outline

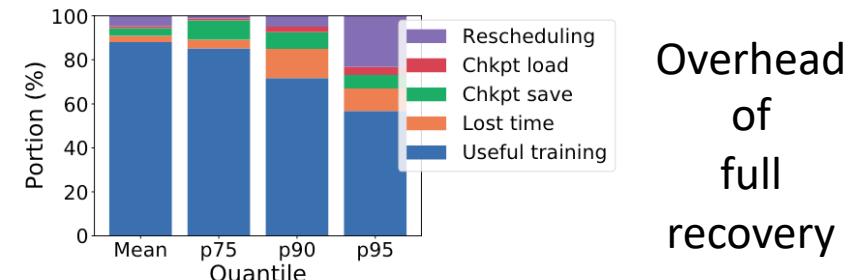




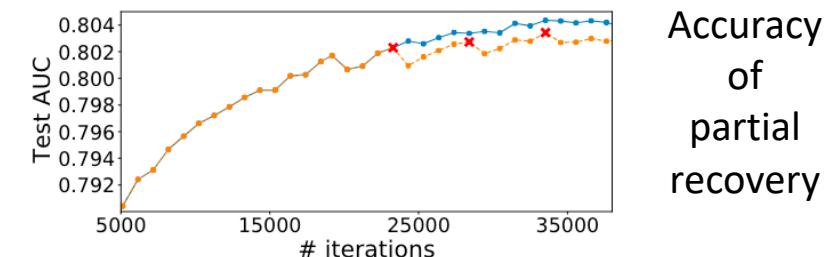
CPR: trade-off between overhead and accuracy

Failures are frequent in large-scale model training

- **Full recovery** will lead to the loss of computation
- **Partial recovery** can harm the model accuracy



Overhead
of
full
recovery



Accuracy
of
partial
recovery

Can we balance between the overhead and the accuracy?

Portion of Lost Samples (PLS)

- the portion of the training data samples whose effect on the model was lost due to a failure
- is a function of checkpoint saving interval, the failure rate and the number of parameter server nodes
- can be used as a metrics to trade-off the **performance overhead** and **the model accuracy**

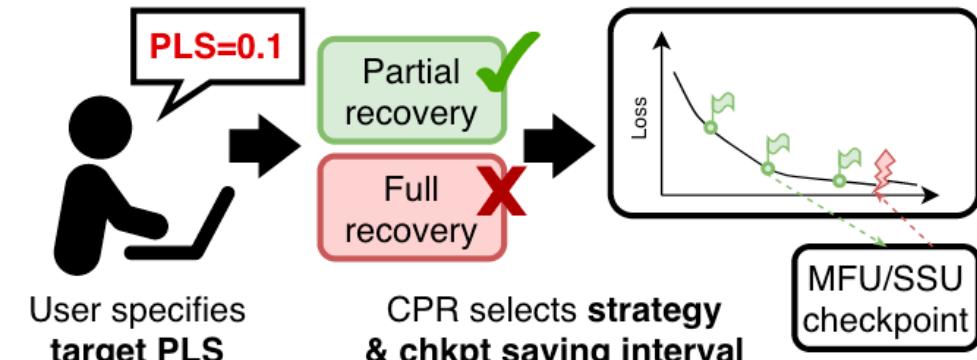
$$\mathbb{E}[PLS] = \frac{0.5T_{save}}{T_{fail}N_{emb}}$$



CPR: trade-off between overhead and accuracy

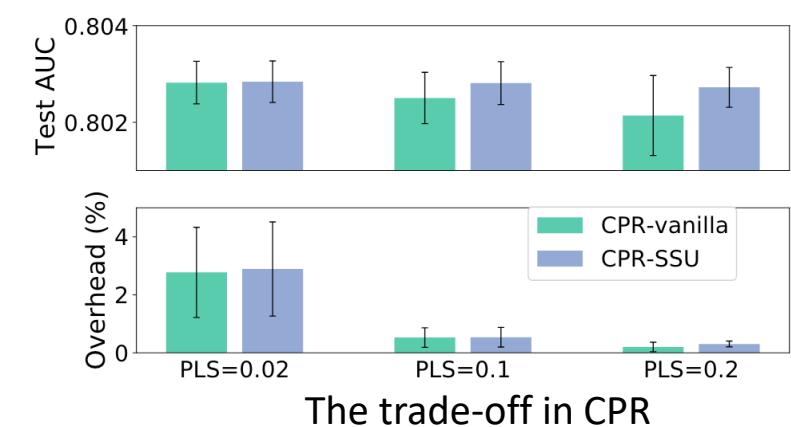
CPR selects between full recovery and partial recovery based on the benefit analysis

- PLS-based checkpointing
 - first choose the saving frequency based on specified PLS
 - trade-off between the overheads of two strategies
 - if the expected benefit is small, choose full recovery
 - otherwise, choose partial recovery
- Frequency-based prioritization
 - With the limited I/O bandwidth, prioritizing to save important updates can make the final model quality to improve.
 - SCAR: prioritize saving parameters with larger changes
 - CPR-MFU: prioritize saving the most-frequently-used parameters
 - CPR-SSU: sub-sample and prioritize used embedding vectors



$$O_{total} \approx O_{save} \frac{T_{total}}{T_{save}} + (O_{load} + \frac{T_{save}}{2} + O_{res}) \frac{T_{total}}{T_{fail}} \quad (1)$$

$$O_{total_par} \approx O_{save} \frac{T_{total}}{T_{save}} + (O_{load} + O_{res}) \frac{T_{total}}{T_{fail}} \quad (2)$$





Conclusion

Redundant Computation

- Pipeline RC in the bubble time to avoid high overhead
- Leverage the inherent redundancy in DP
- Decoupling the planning from the execution

Optimized Checkpointing

- Pipeline checkpointing with computation and communication.
- Leverage hierarchical storage to improve checkpointing frequency
- Avoid the lost of computation
- Leverage partial features and quantization to reduce size

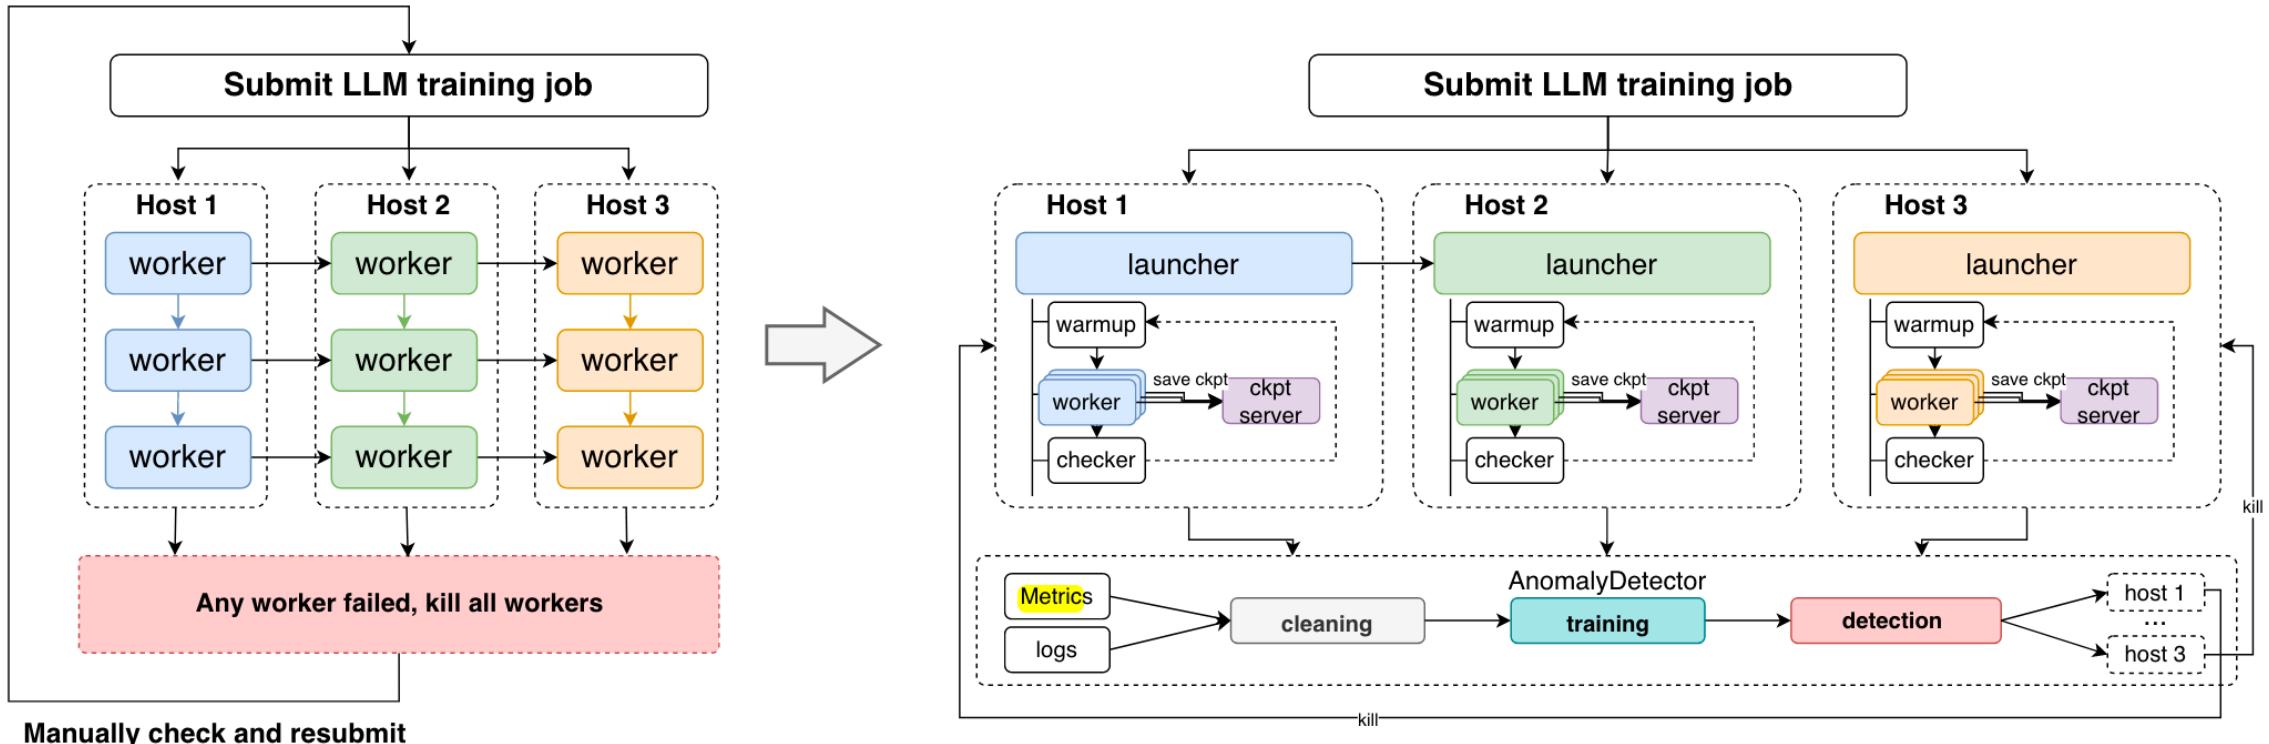
Approximation

- Trade-off between the overhead of full recovery and accuracy of partial recovery



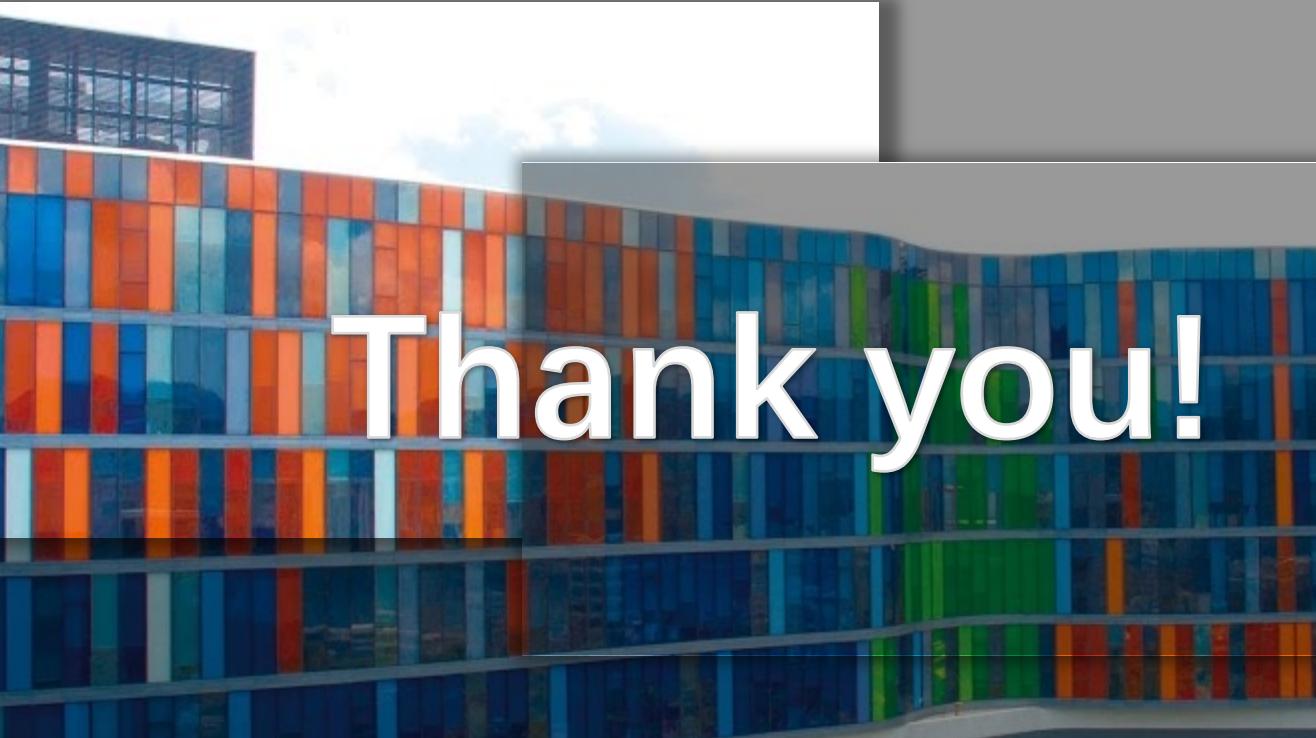
Conclusion

Future direction



Automated anomaly detection

Efficient failure recovery



Thank you!



香港中文大學
The Chinese University of Hong Kong