



# TraceMesh: Scalable and Streaming Sampling for Distributed Traces

---

**Zhuangbin Chen<sup>1</sup>, Zhihan Jiang<sup>2</sup>, Yuxin Su<sup>1</sup>,  
Michael R. Lyu<sup>2</sup>, Zibin Zheng<sup>1</sup>**

<sup>1</sup>Sun Yat-sen University

<sup>2</sup>The Chinese University of Hong Kong

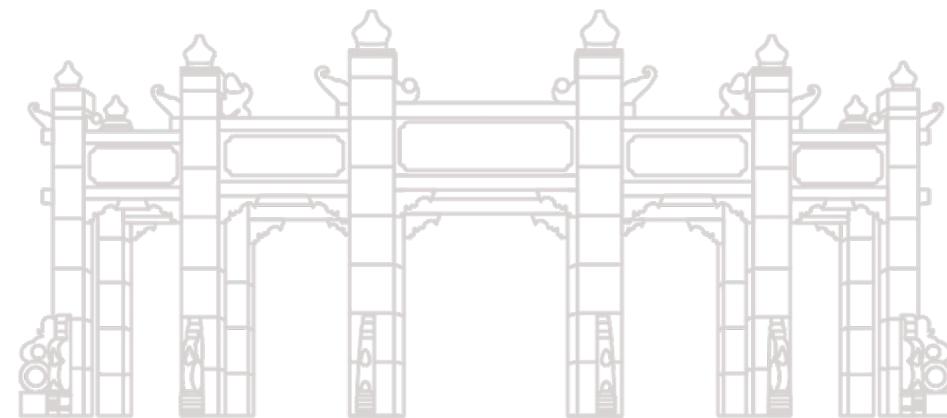
[chenzhb36@mail.sysu.edu.cn](mailto:chenzhb36@mail.sysu.edu.cn)

# Outline

---

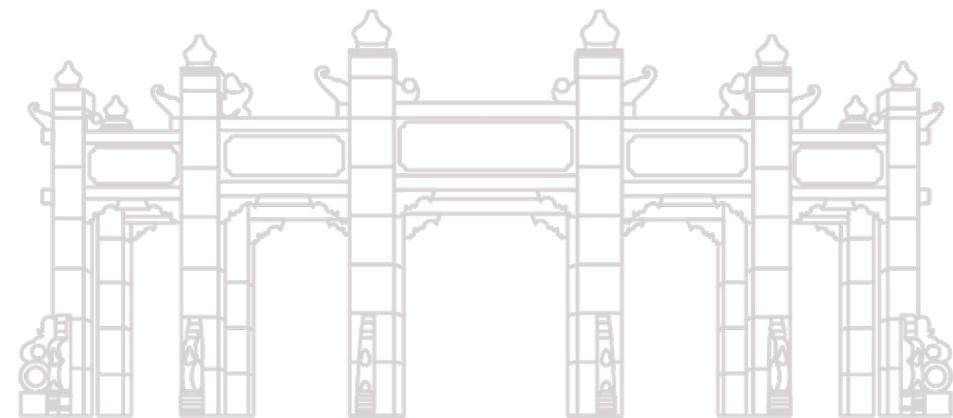


- **Introduction & Motivation**
- **TraceMesh Overview & Design**
- **Evaluation**
- **Conclusion**

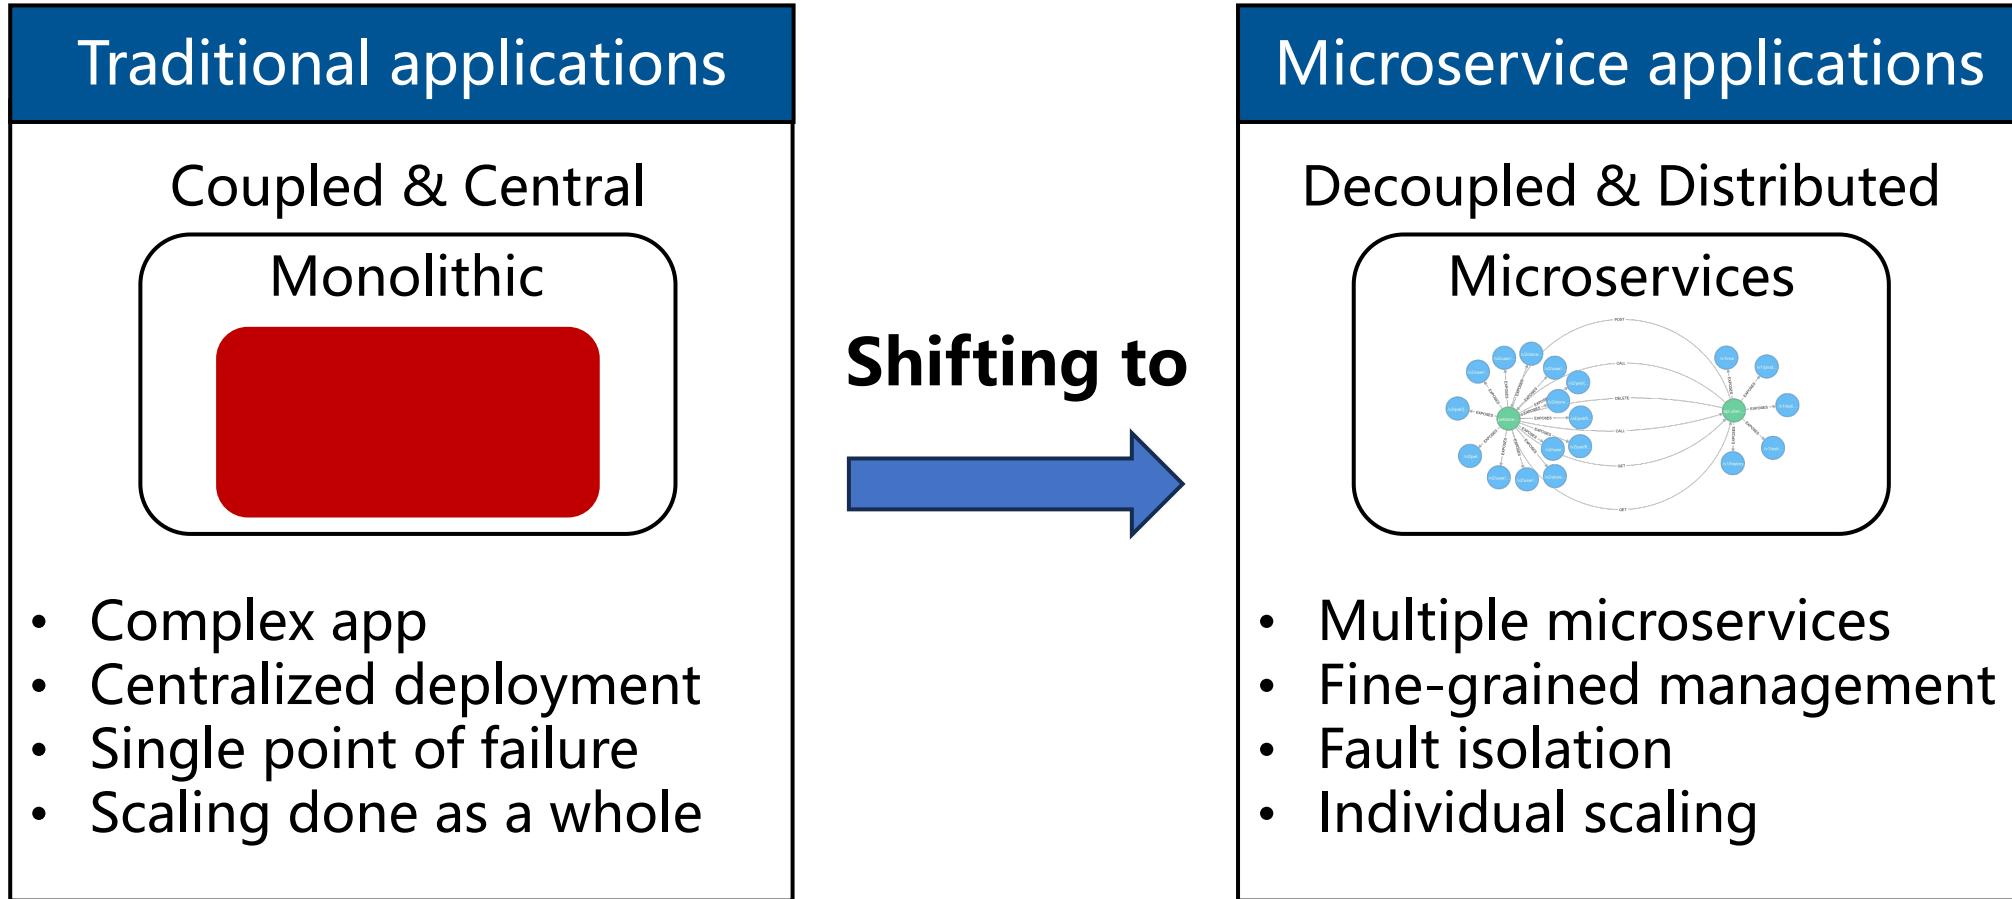




# Introduction & Motivation

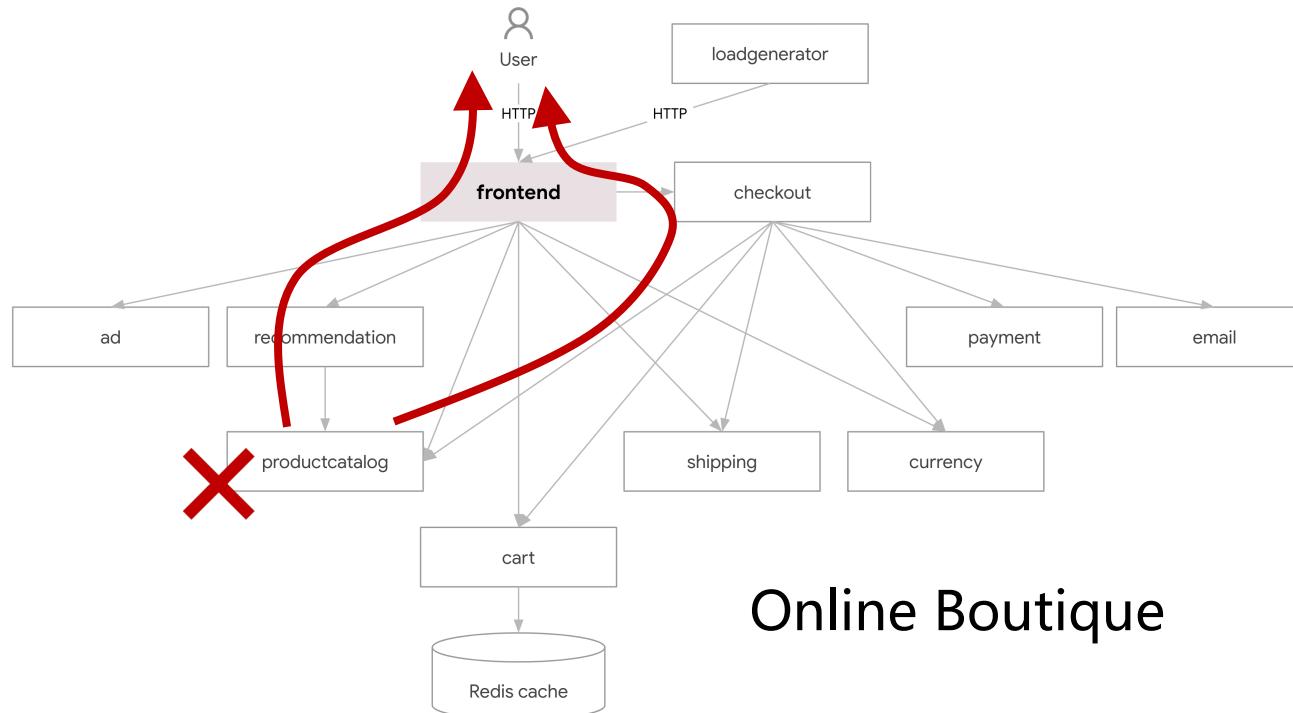


# Shifting to Microservices



# Complex Reliability Management

- E2E behavior can be affected by any microservice
- Failures cascading effect
- Symptoms and root causes can be far apart



# Distributed Tracing

- A directed tree recording executions across microservices
- Trace events: time, operations, messages, attributes
- Critical for performance monitoring, root cause analysis, etc.

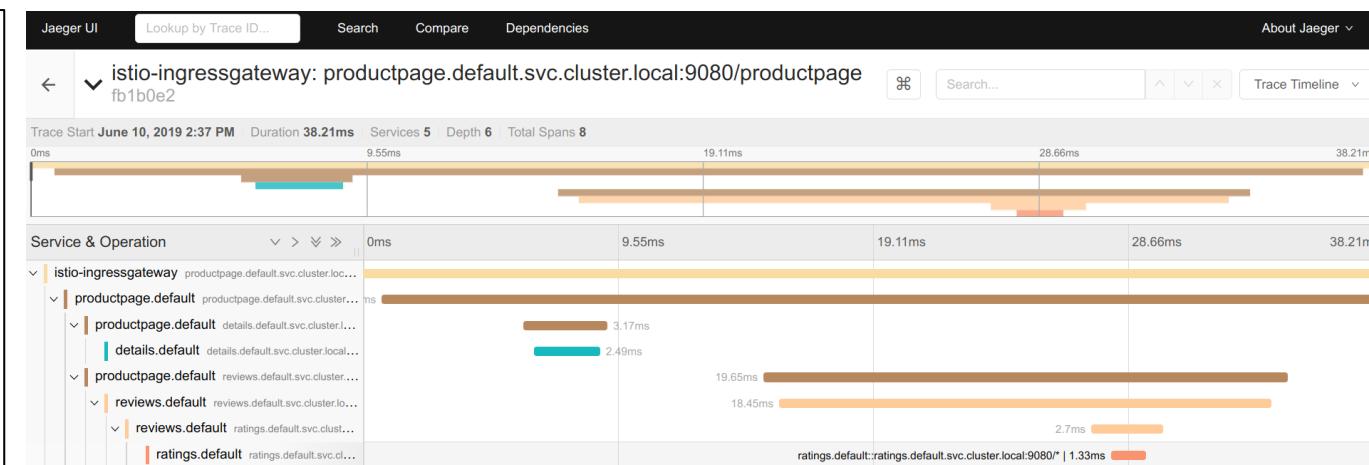
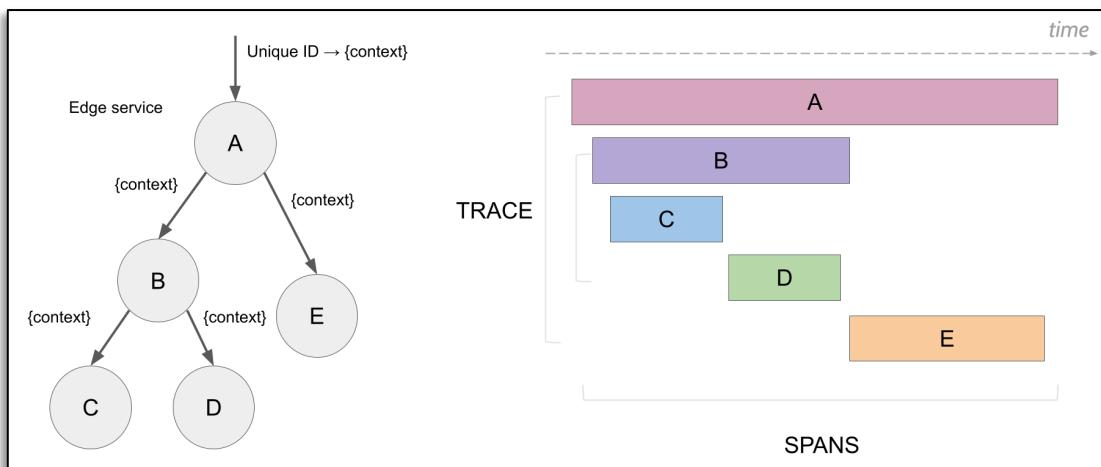
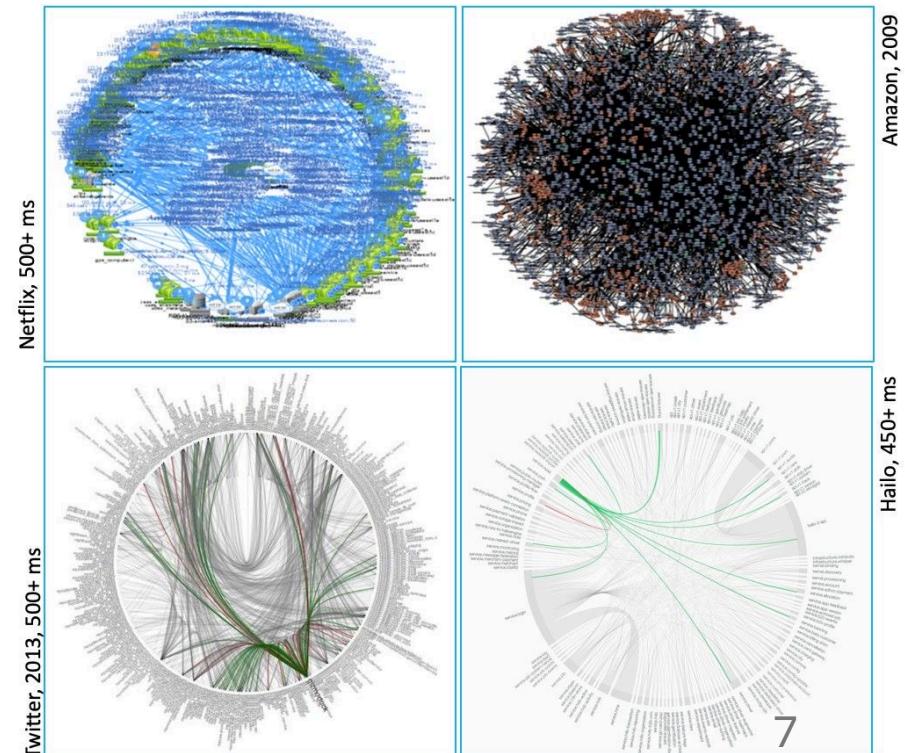


Figure 1 source: <https://medium.com/nikeengineering/hit-the-ground-running-with-distributed-tracing-core-concepts-ff5ad47c7058>  
Figure 2 source: <https://istio.io/latest/docs/tasks/observability/distributed-tracing/jaeger/>

# Trace Data are Large in Size

- Complex microservices systems
  - Uber is composed of thousands of microservices
  - WeChat system hosts more than 3,000 services
- Huge volume of trace data
  - E.g., Google generates about 1,000 TB of raw traces on a daily basis
  - **Significant overhead to trace generation, collection, and ingestion**



Amazon, 2009

Twitter, 2013, 5000+ ms

# Most Trace Data are Useless

- Most traces are recurring and useless
  - Cloud services often guarantee more than 99% uptime in their SLAs
  - Only edge-case traces are interesting and **should be sampled**
  - **By definition, these traces are rare**

AMAZON ELASTIC COMPUTE CLOUD (EC2) (INSTANCE LEVEL)	
Monthly Uptime Percentage	Service Credit Percentage
Less than 99.5% but equal to or greater than 99.0%	10%
Less than 99.0% but equal to or greater than 95.0%	30%
Less than 95.0%	100%

[Amazon Compute Full SLA](#)

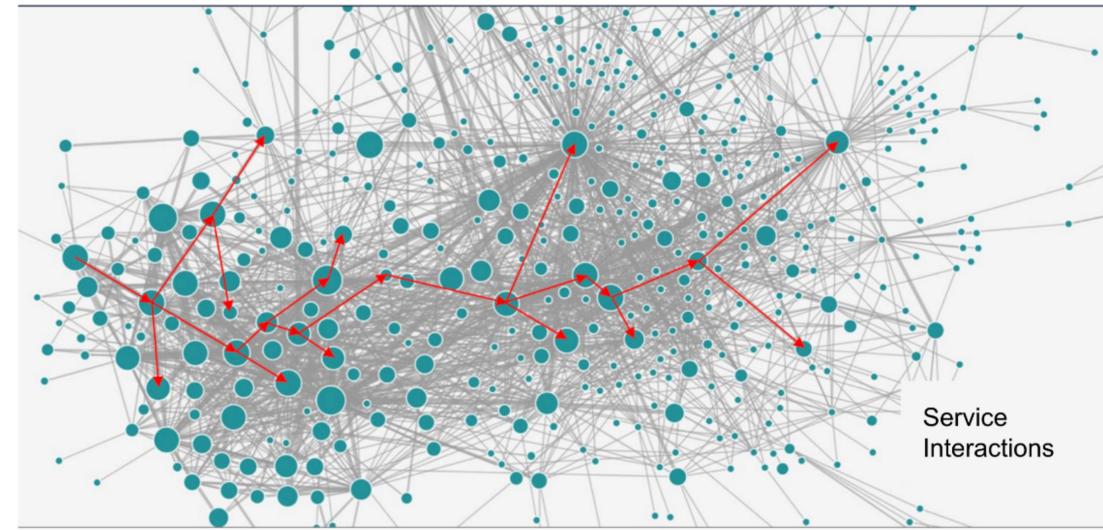
Compute | Containers

# Trace Sampling

- Head-based sampling
  - Decision made **prior** to request execution
  - Uniformly sample traces at random with a sampling rate, e.g., 1%
  - Could result in common-case traces
- Tail-based sampling
  - Decision made **after** request execution
  - Biased sampling based on latency, HTTP status code, etc.
  - Learning-based approaches to identify uncommon traces

# Challenges of Tail-based Trace Sampling

- Large feature dimension
  - Curse of dimensionality
  - Compromised performance and scalability
- Frequent pattern change
  - Microservices undergo continuous updates
  - New trace patterns could emerge

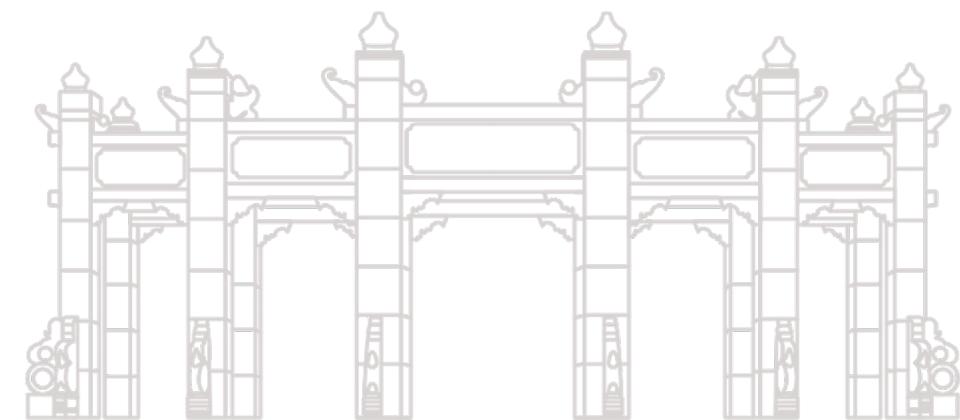


Microservices call graph collected at Uber

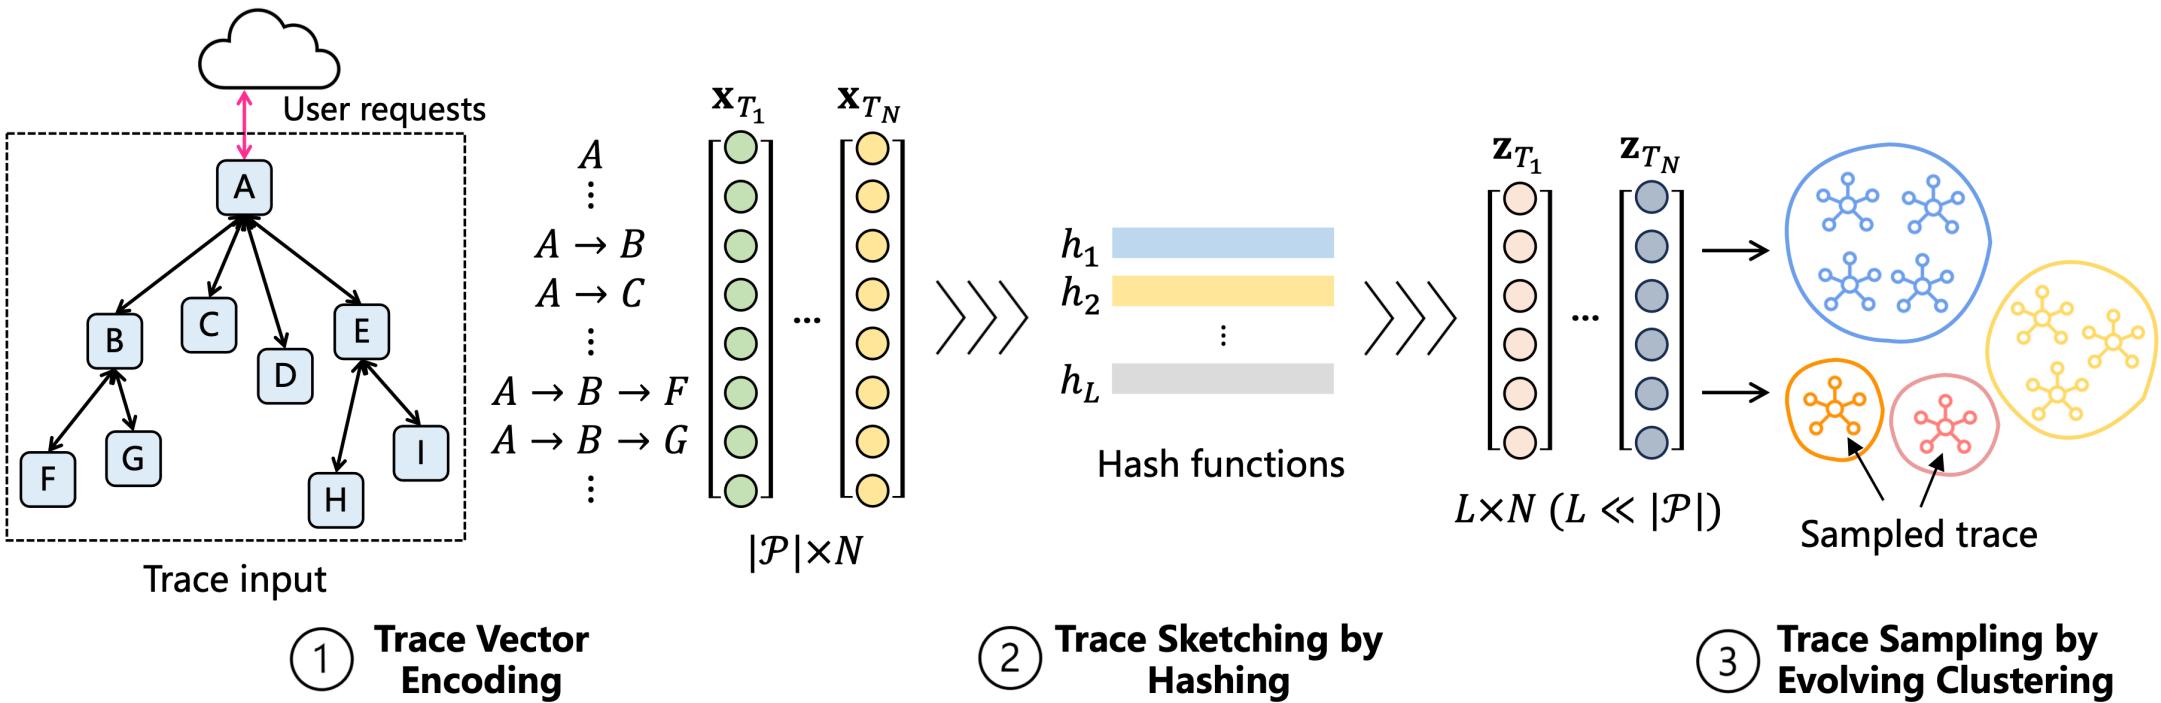
**Not well addressed in existing approaches!**



# TraceMesh Overview & Design



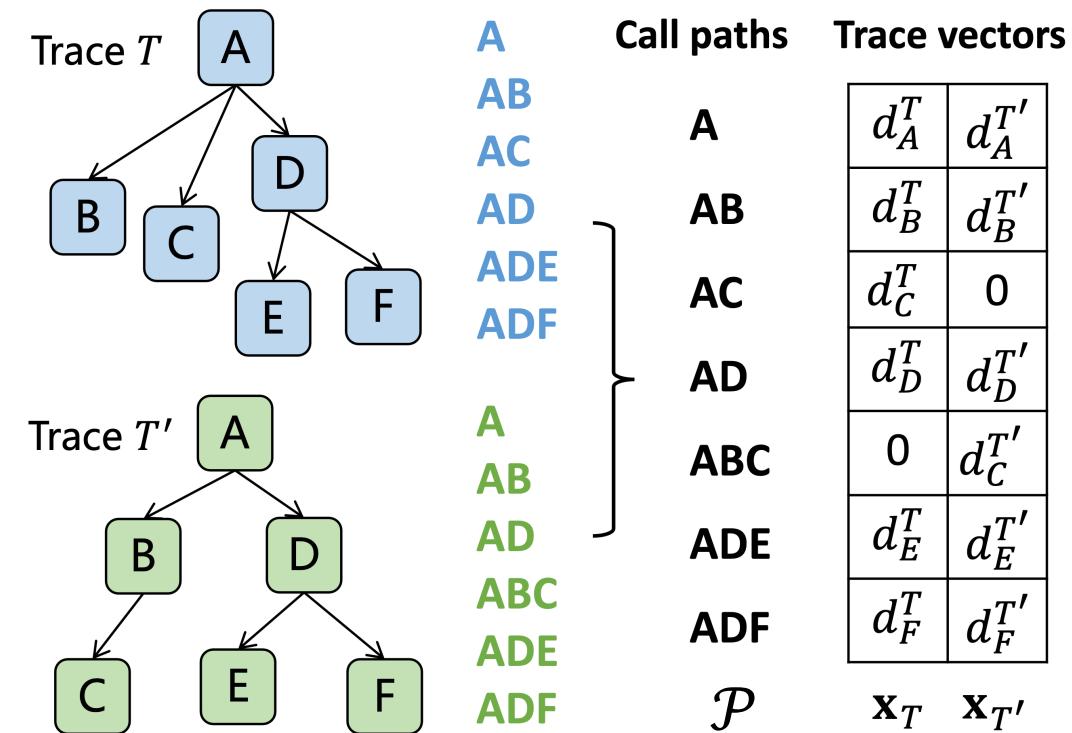
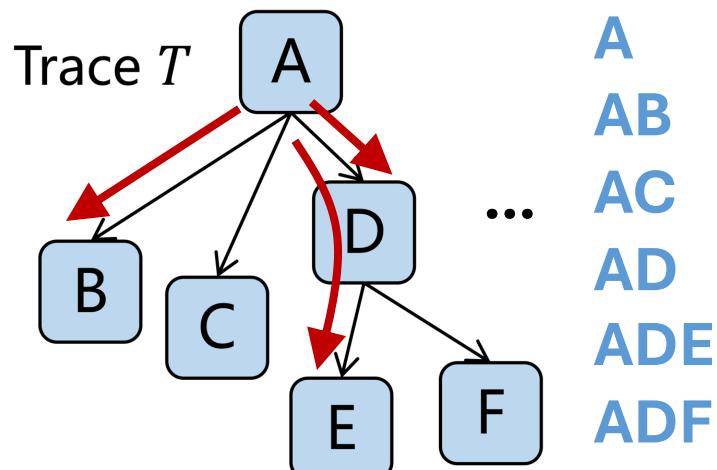
# TraceMesh Overview



1. Transform graph traces into numerical feature vectors
2. Perform streaming trace vector encoding
3. Adaptively sample target traces by clustering

# Step 1: Trace Vector Encoding

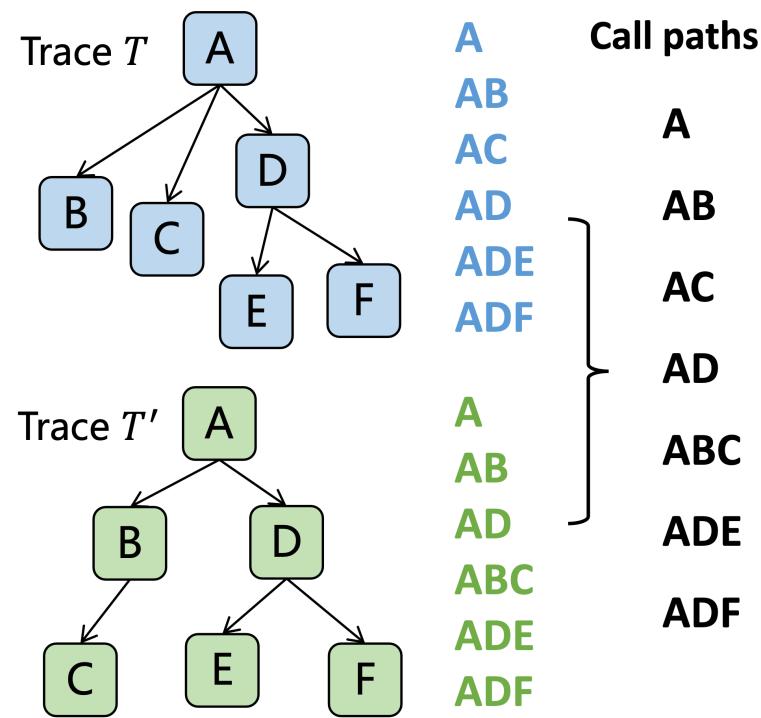
- Two types of trace features are extracted:
  - Structural feature
  - Temporal feature



The complete set of call paths can be large and dynamic!

# Step 2: Trace Sketching by Hashing

- The first challenge:



## Locality-Sensitive Hashing (LSH)

Project high-dimensional vectors into a low-dimensional space while preserving their similarity

$\mathcal{P}$  can be very large

# Step 2: Trace Sketching by Hashing

- The projection process for a trace vector  $\mathbf{x}$  in  $\mathbb{R}^{|\mathcal{P}|}$ :
  1. Instantiate  $L$  projection vectors  $\mathbf{r}_1, \mathbf{r}_2, \dots, \mathbf{r}_L \in \{+1, -1\}^{|\mathcal{P}|}$
  2. Compute trace sketch vector as  $\mathbf{z} = [h_{\mathbf{r}_1}(\mathbf{x}), \dots, h_{\mathbf{r}_L}(\mathbf{x})] \in \{+1, -1\}^L$ :

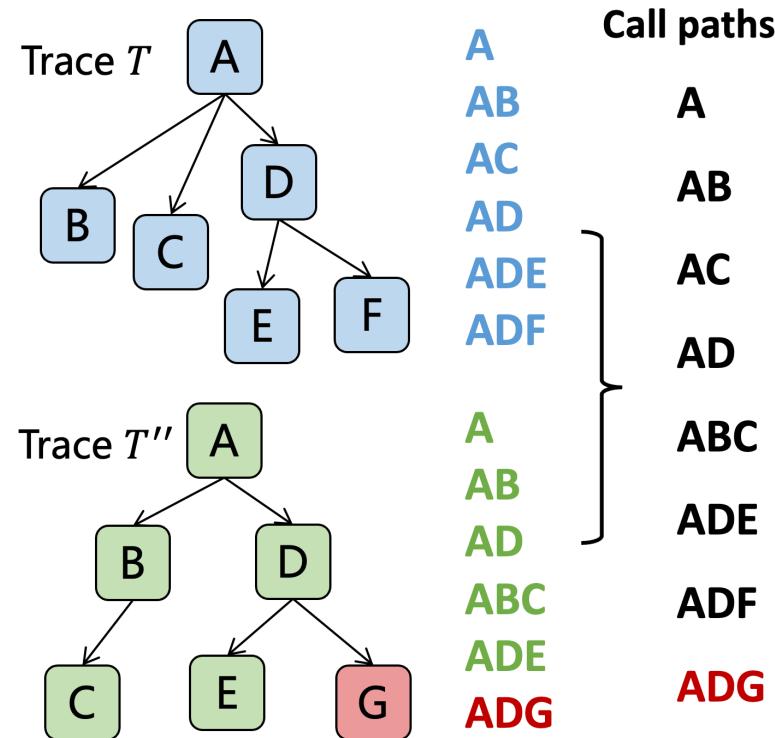
$$h_{\mathbf{r}_l}(\mathbf{x}) = \begin{cases} +1 & \text{if } \mathbf{x} \cdot \mathbf{r}_l \geq 0 \\ -1 & \text{if } \mathbf{x} \cdot \mathbf{r}_l < 0 \end{cases} \quad \text{i.e., } h_{\mathbf{r}_l}(\mathbf{x}) = \text{sign}(\mathbf{x} \cdot \mathbf{r}_l)$$

- Calculate the similarity for two traces by quantifying the agreement level between the sketch vectors:

$$\text{sim}(T, T') \propto \frac{|\{l : \mathbf{z}_T(l) = \mathbf{z}_{T'}(l)\}|}{L}$$

# Step 2: Trace Sketching by Hashing

- The second challenge:



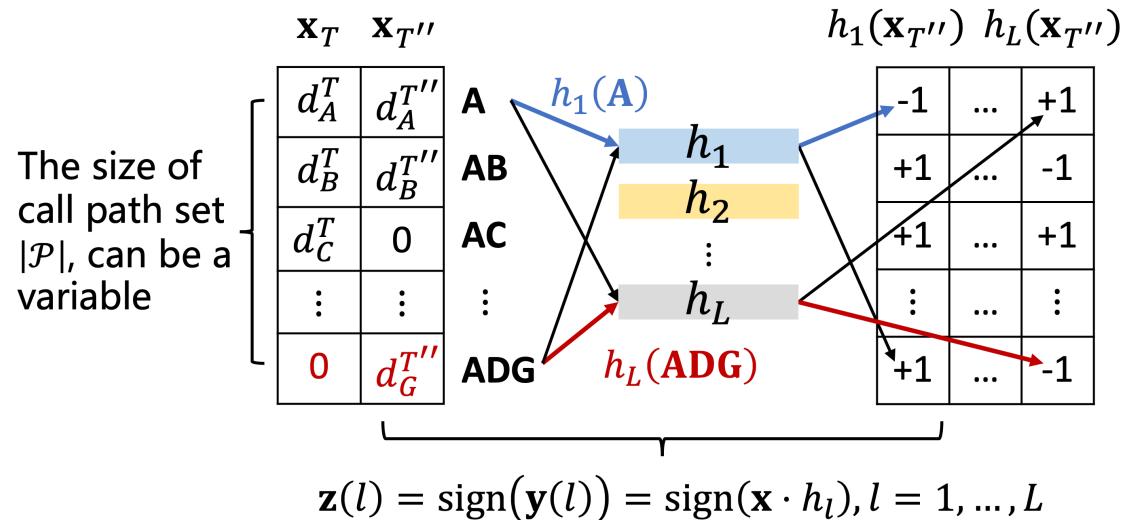
## StreamHash

Replace project vectors with hash functions to deterministically map any call paths to  $\{+1, -1\}$

$\mathcal{P}$  can be dynamic

# Streaming Trace Vector Encoding

- $L$  hash functions  $h_1, h_2, \dots, h_L$  uniformly picked from a family  $\mathcal{H}$  of hash functions



Fixed trace vector size  $L$

$\mathbf{z}_{T''}(1)$	$\text{sign}(-1 * d_A^{T''} + \dots + (+1) * d_G^{T''})$
⋮	⋮
$\mathbf{z}_{T''}(L)$	$\text{sign}(+1 * d_A^{T''} + \dots + (-1) * d_G^{T''})$

The trace  $T''$  with a new path **ADG** can be seamlessly encoded as  $\mathbf{z}_{T''}$

# Step 3: Trace Sampling by Evolving Clustering

- A **density-based clustering** approach for stream data
- Each **micro-cluster (MC)** has three attributes:
  - A weight  $w$ , determined by the traces inside it and time
  - A center  $c$ , the weighted center of the traces inside it
  - A radius  $r$ , the weighted avg distance from points to  $c$
- Two types of MCs with different weight constraints:
  - **Potential MC (PMC)**: contains frequent and usual traces
  - **Outlier MC (OMC)**: contains potentially outlier traces

# Step 3: Trace Sampling by Evolving Clustering

- When a new trace  $T$  arrives, TraceMesh
  - Merges it into the nearest PMC, and samples it with probability:

$$p_T = \mathcal{B} \times \left(1 - \frac{w_p^*}{\sum_{i=0}^{N_{pmc}} w^{(i)}}\right) \quad \mathcal{B}: \text{sampling budget}$$

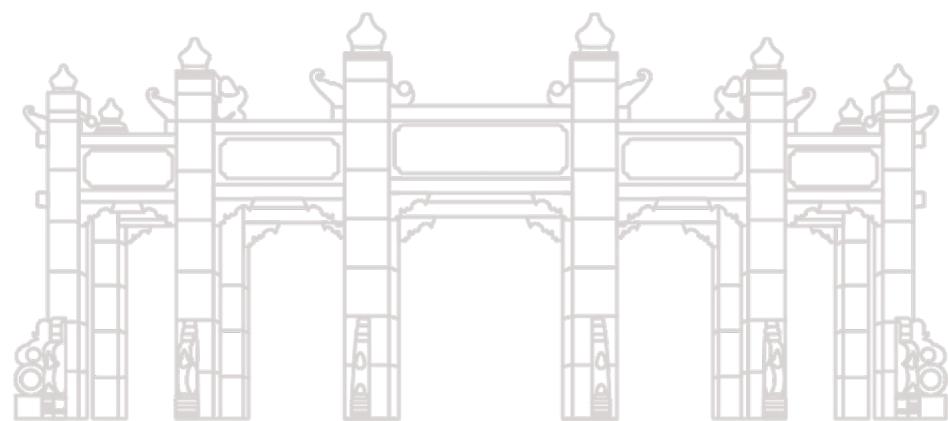
- If failed, merges it into the nearest OMC, and samples it
- If failed, makes  $T$  a new OMC, and samples it

In this process, the **weight and role** of each MC will change.

TraceMesh **dynamically adjusts the sampling strategy** to identify the target traces without over-sampling recurring traces.



# Evaluation



# Datasets

- Traces from open-source microservices
  - Train Ticket
  - Online Boutique
- Traces from production systems

DATASET STATISTICS

<b>Dataset</b>	<b>#Span</b>	<b>#Train</b>	<b>#Test</b>	<b>#Label</b>
Train Ticket	201	653	31,814	624
Online Boutique	43	1,024	78,931	225
Industry	1,308	11,847	497,439	2,453

# Metrics & Baselines

- Metrics

$$\text{Coverage} = \frac{\# \text{ target traces sampled}}{\# \text{ target traces}}$$

$$\text{Sampling rate} = \frac{\# \text{ traces sampled}}{\# \text{ traces}}$$

- Baselines
  - Uniform sampling
  - Perch [SoCC 2018]
  - Sifter [SoCC 2019]
  - Sieve [ICWS 2021]
  - SampleHST [NOMS 2023]

# Experimental Results

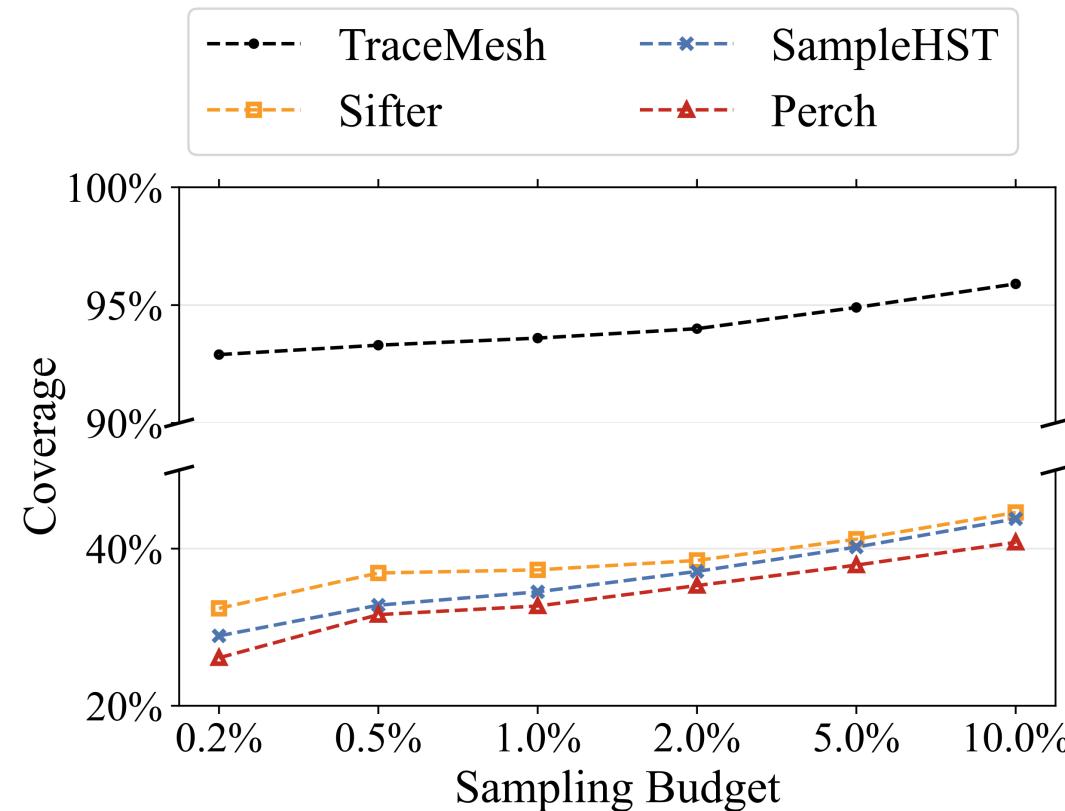
- TraceMesh achieves the best coverage
- TraceMesh has a relatively good sampling rate

EXPERIMENTAL RESULTS OF TRACE SAMPLING

Method	Train Ticket		Online Boutique		Industry	
	Coverage	Sampling Rate	Coverage	Sampling Rate	Coverage	Sampling Rate
Uniform	1.1%	<b>1.00%</b>	1.1%	1.00%	1.0%	1.00%
Sieve	71.8%	6.24%	82.2%	3.91%	61.1%	3.29%
Sifter	50.6%	1.28%	42.2%	1.13%	37.3%	1.16%
SampleHST	46.2%	1.19%	38.2%	1.07%	34.5%	1.10%
Perch	42.9%	<b>1.00%</b>	35.1%	<b>0.98%</b>	32.7%	0.95%
TRACEMESH	<b>98.7%</b>	2.34%	<b>100%</b>	1.35%	<b>93.6%</b>	<b>0.83%</b>

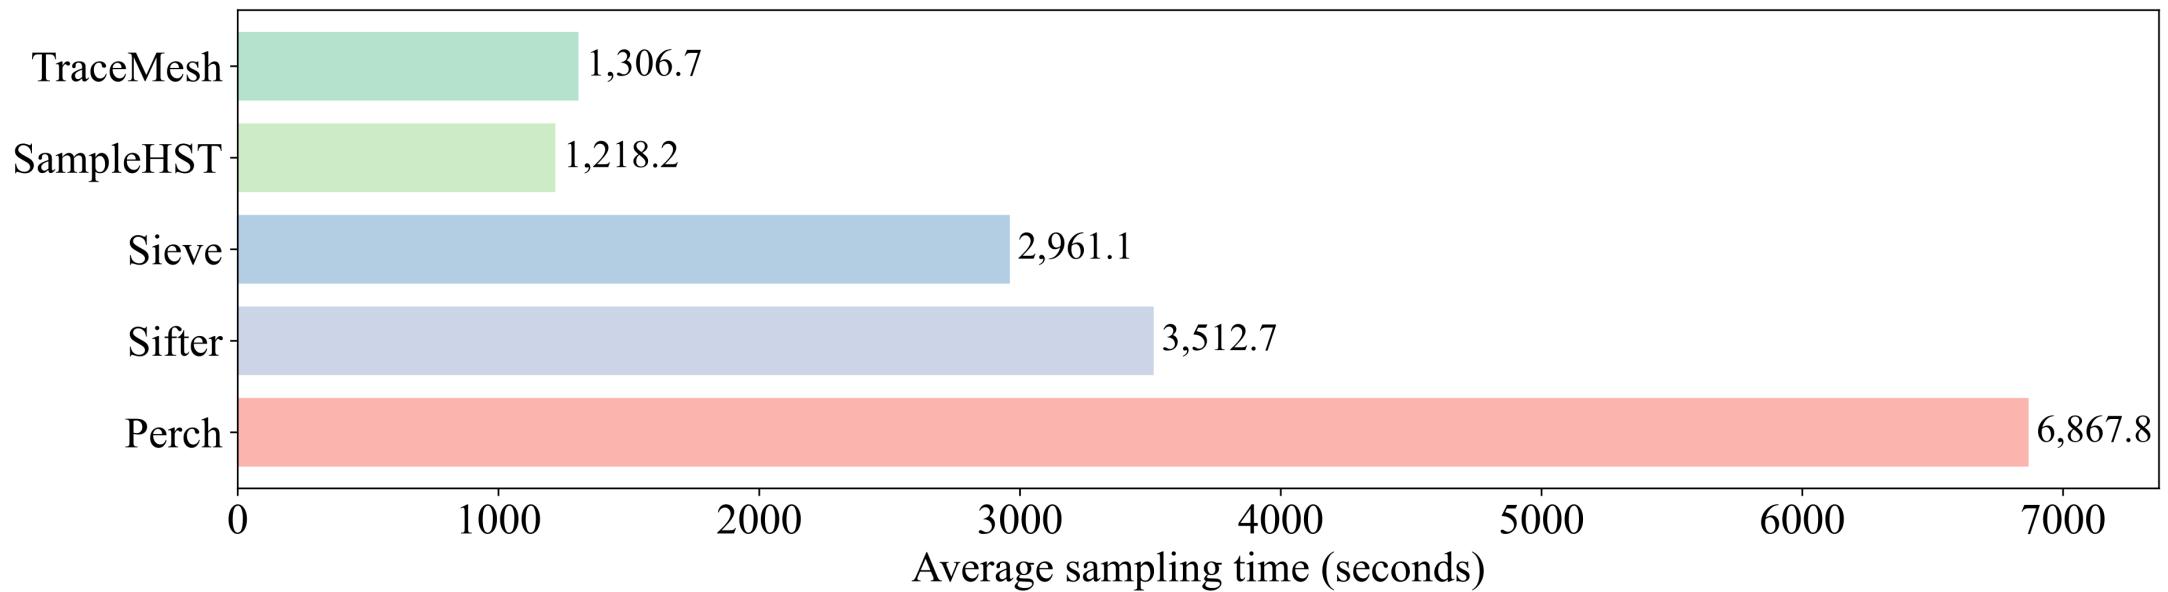
# Coverage vs. Sampling Budget

- Coverage increases with larger sampling budget
- TraceMesh performs trace sampling more accurately



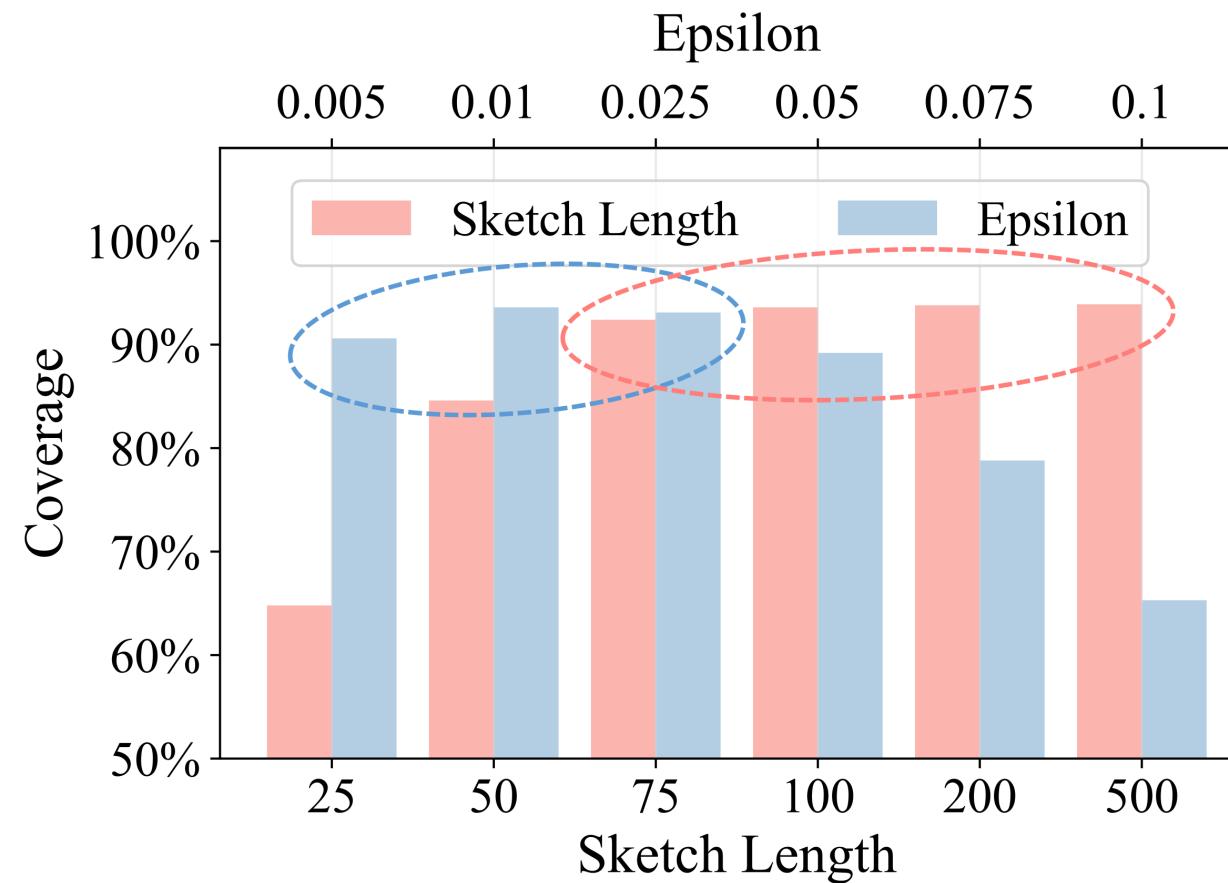
# Efficiency

- TraceMesh performs trace sampling efficiently



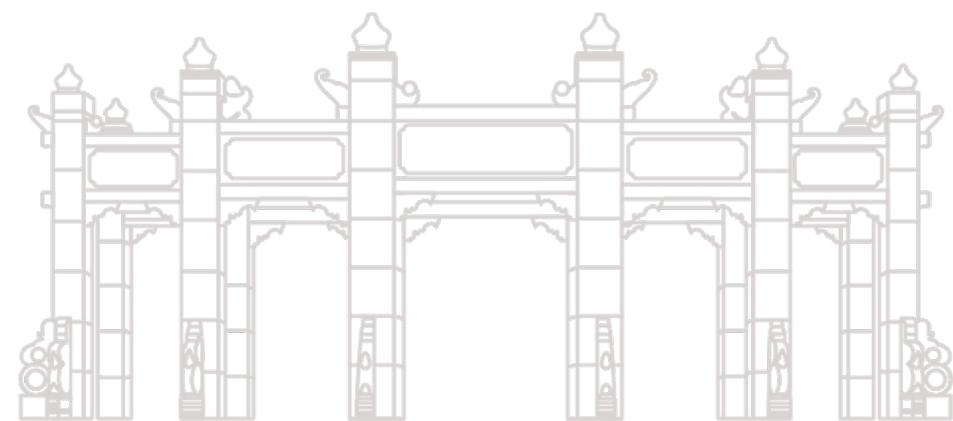
# Parameter Sensitivity

- TraceMesh can achieve stable performance





# Conclusion



# Conclusion

- TraceMesh, a **scalable** and **streaming** trace sampler
  - Large feature dimension: LSH for dimension reduction
  - Frequent pattern changes: StreamHash and evolving clustering
- TraceMesh outperforms baselines in both sampling accuracy and efficiency



# TraceMesh: Scalable and Streaming Sampling for Distributed Traces

**Thanks!**  
**Q & A**

