



LILAC: Log Parsing using LLMs with Adaptive Parsing Cache

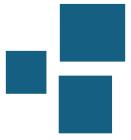
Zhihan Jiang¹, Jinyang Liu¹, Zhuangbin Chen², Yichen Li¹, Junjie Huang¹,
Yintong Huo¹, Pinjia He³, Jiazhen Gu¹ and Michael R. Lyu¹

¹The Chinese University of Hong Kong,

²Sun Yat-sen University,

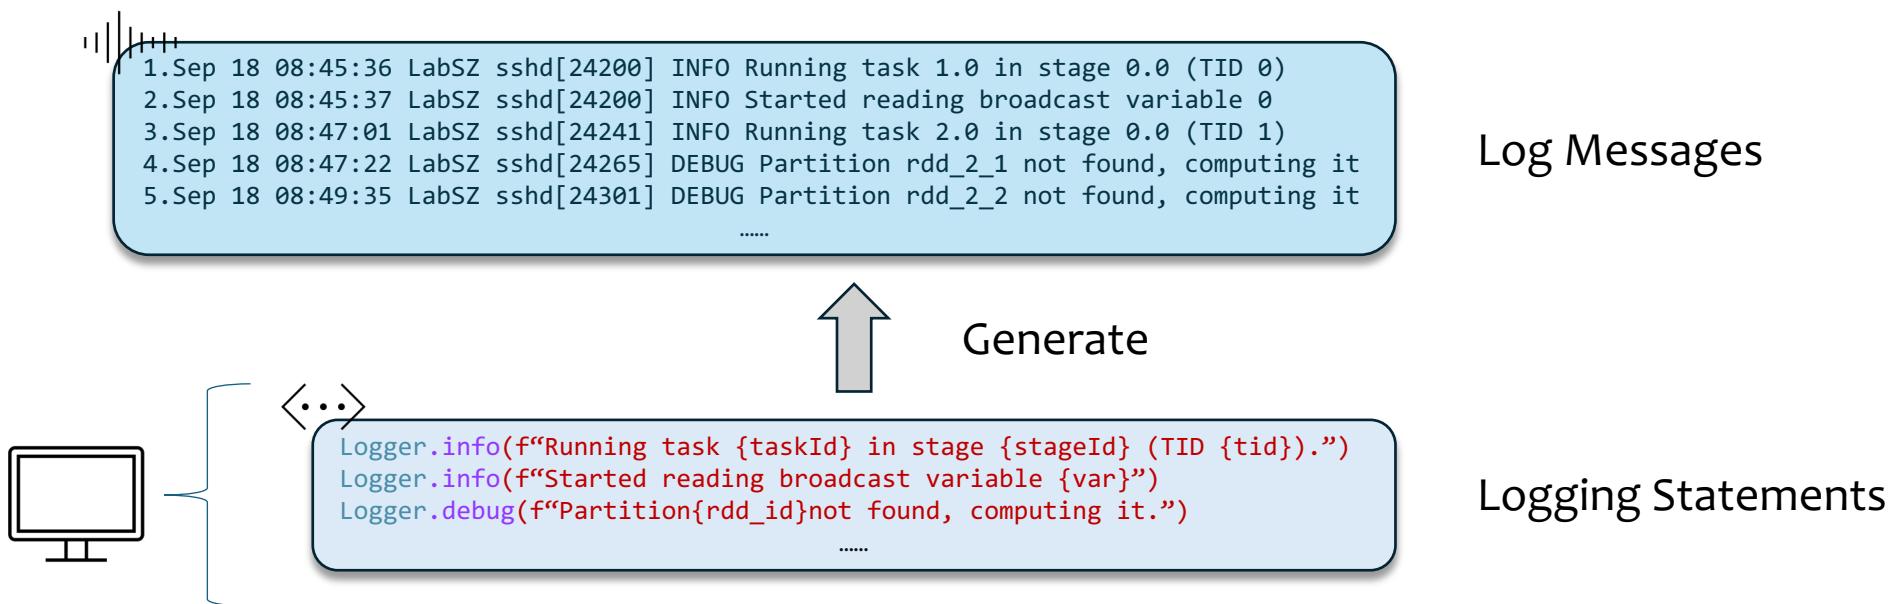
³The Chinese University of Hong Kong (Shenzhen)





System Log

Logs are generated by logging statements and record runtime information of a system.





System Log

Automated Log Analysis is important for:



Anomaly Detection



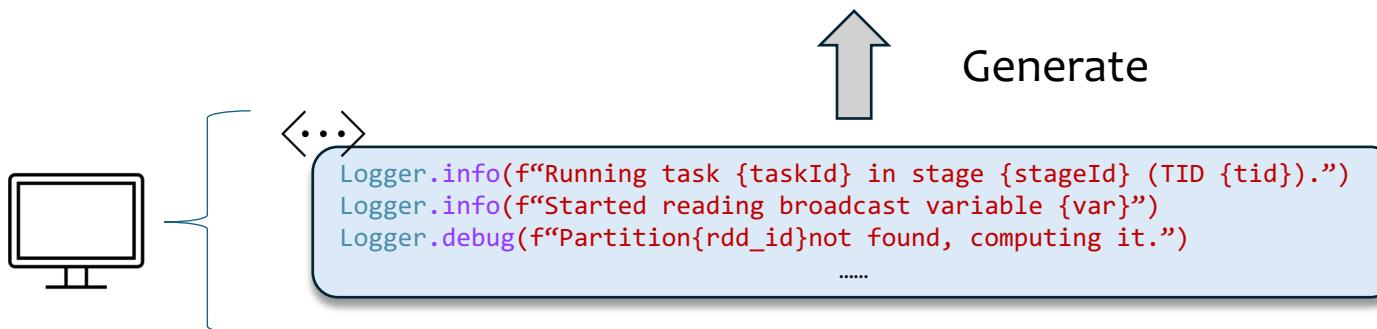
Fault Localization



Root Cause Analysis

```
1.Sep 18 08:45:36 LabSZ sshd[24200] INFO Running task 1.0 in stage 0.0 (TID 0)
2.Sep 18 08:45:37 LabSZ sshd[24200] INFO Started reading broadcast variable 0
3.Sep 18 08:47:01 LabSZ sshd[24241] INFO Running task 2.0 in stage 0.0 (TID 1)
4.Sep 18 08:47:22 LabSZ sshd[24265] DEBUG Partition rdd_2_1 not found, computing it
5.Sep 18 08:49:35 LabSZ sshd[24301] DEBUG Partition rdd_2_2 not found, computing it
.....
```

Log Messages



Logging Statements



Log Parsing

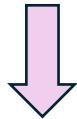
Logs are unstructured, which are not suitable for understanding, analyzing and storing.



Log Parsing

Converting raw logs into structured formats is a critical prerequisite step!

```
Sep 18 08:47:22 DEBUG Partition rdd_2_1 not found in 127.0.0.1  
Sep 18 08:47:35 DEBUG Partition rdd_2_2 not found in 127.0.0.1  
Sep 18 08:47:49 DEBUG Partition rdd_2_1 not found in 127.0.0.1  
Sep 18 08:48:17 DEBUG Partition rdd_2_3 not found in 127.0.0.1
```

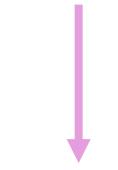


Timestamp	Event	Level	Log Templates	Parameters
Sep 18 08:47:22	e1	DEBUG	Partition <*> not found in <*>	rdd_2_1 127.0.0.1
Sep 18 08:47:35	e1	DEBUG	Partition <*> not found in <*>	rdd_2_2 127.0.0.1
Sep 18 08:47:49	e1	DEBUG	Partition <*> not found in <*>	rdd_2_1 127.0.0.1
Sep 18 08:48:17	e1	DEBUG	Partition <*> not found in <*>	rdd_2_3 127.0.0.1

Substantial Raw Log



How to understand and process these complex logs?



Log Analysis



Four same events happened on **three different partitions** (rdd_2_1, rdd_2_2, rdd_2_3) and **the same address** (127.0.0.1)

Structured Parsed Log



Log Parsing

Goal: to distinguish between constant part and variable part.



Timestamp	Component	Level	Log Templates	Parameters
Sep 18 08:45:36	LabSZ sshd[24200]	INFO	Running task <*> in stage <*> (TID <*>)	1.0 0.0 0
Sep 18 08:45:37	LabSZ sshd[24200]	INFO	Started reading broadcast variable <*>	0
Sep 18 08:47:01	LabSZ sshd[24241]	INFO	Running task <*> in stage <*> (TID <*>)	2.0 0.0 1
Sep 18 08:47:22	LabSZ sshd[24265]	DEBUG	Partition <*> not found, computing it	rdd_2_1
Sep 18 08:49:35	LabSZ sshd[24301]	DEBUG	Partition <*> not found, computing it	rdd_2_2
.....				

Parsed Results



Log Parsing

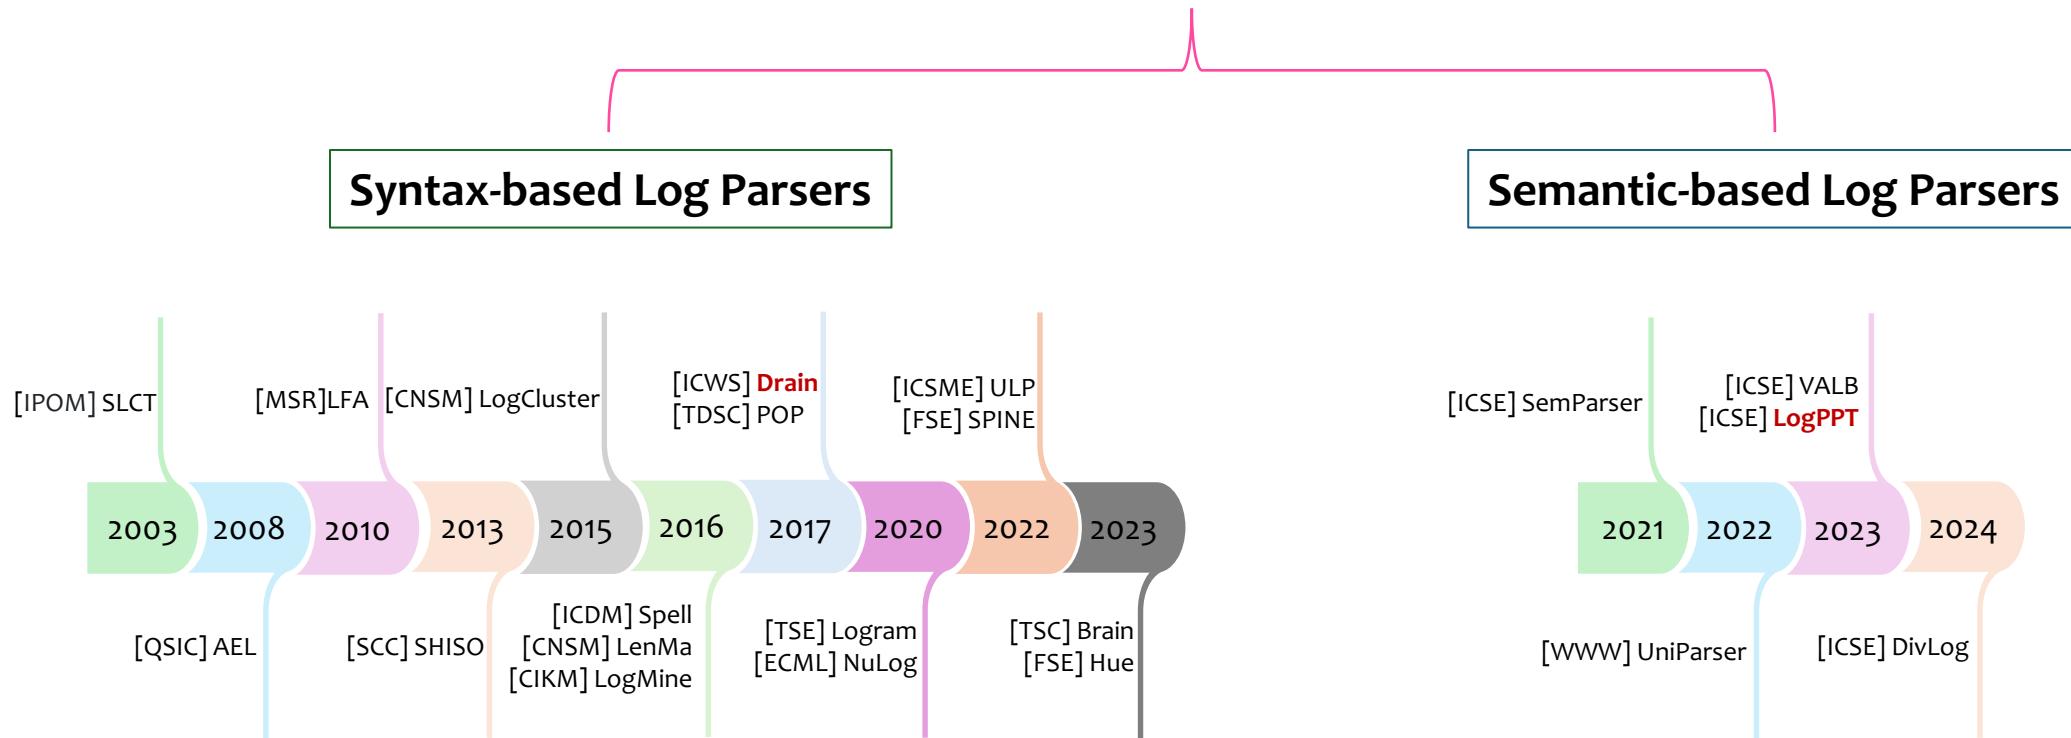
```
1. Sep 18 08:45:36 LabSZ sshd[24200] INFO Running task 1.0 in stage 0.0 (TID 0)
2. Sep 18 08:45:37 LabSZ sshd[24200] INFO Started reading broadcast variable 0
3. Sep 18 08:47:01 LabSZ sshd[24241] INFO Running task 2.0 in stage 0.0 (TID 1)
4. Sep 18 08:47:22 LabSZ sshd[24265] DEBUG Partition rdd_2_1 not found, computing it
5. Sep 18 08:49:35 LabSZ sshd[24301] DEBUG Partition rdd_2_2 not found, computing it
.....
```

Raw Logs



Related Work

A wide range of **data-driven** methods have been proposed to parse logs.



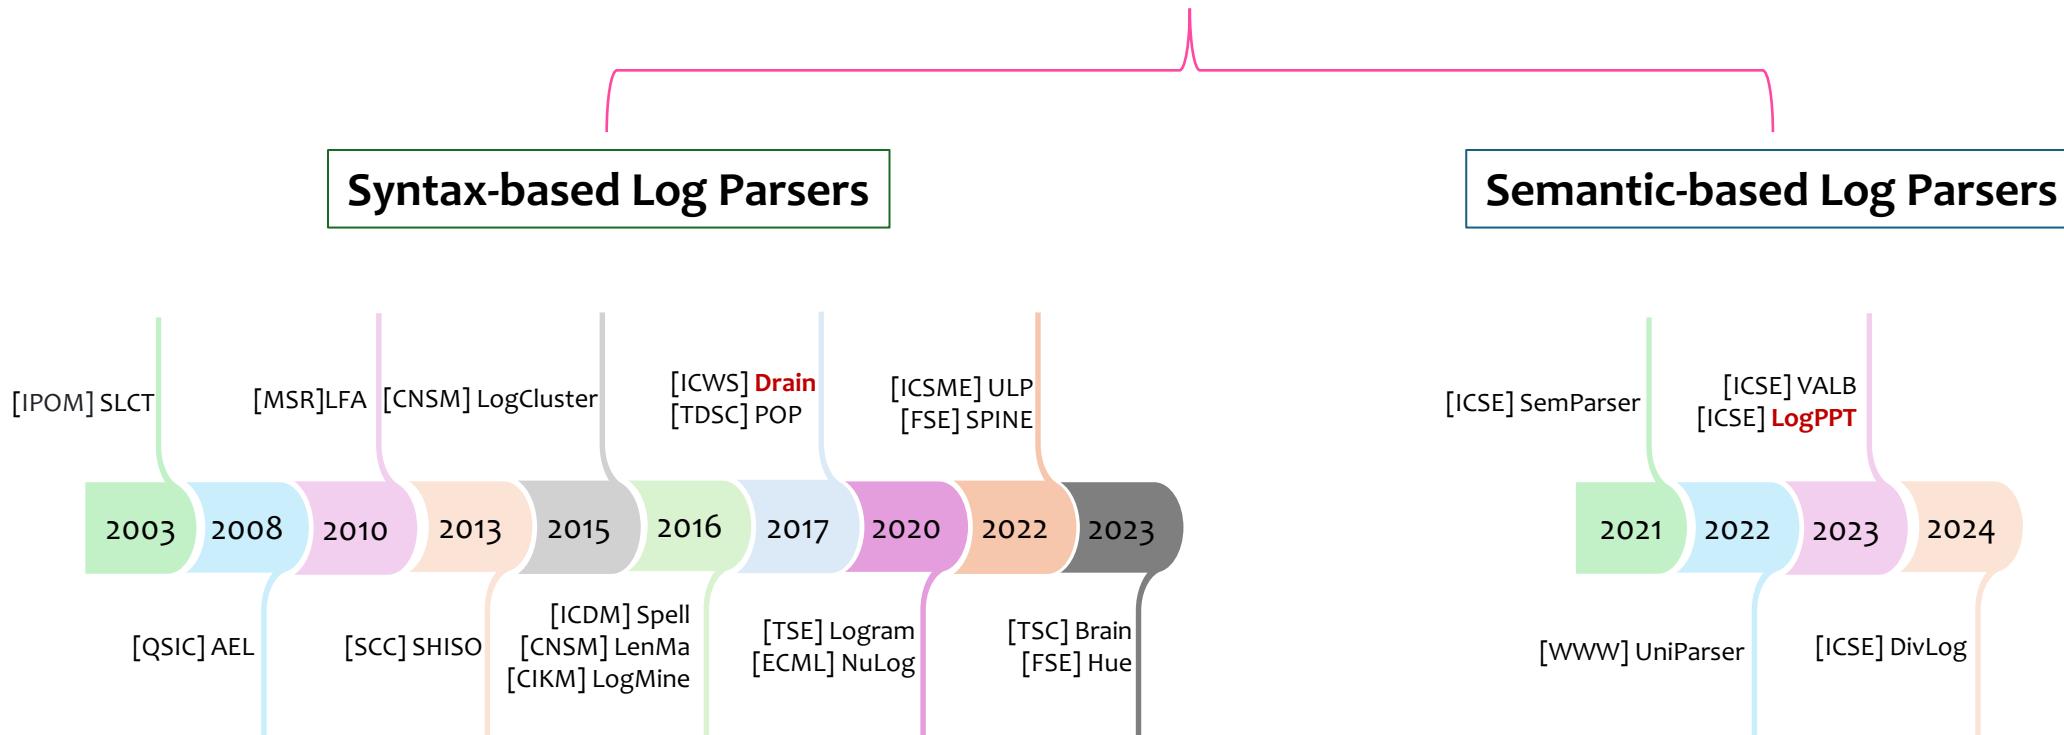


Limitation of Existing Parsers

Logs are continuously evolving!



The performance of existing log parsers in practice remains unsatisfactory.



"Guidelines for assessing the accuracy of log message template identification techniques", ICSE'22

"Investigating and improving log parsing in practice", FSE'22

"A Large-scale Evaluation for Log Parsing Techniques: How Far are We?", ISSTA'24



Limitation of Existing Parsers

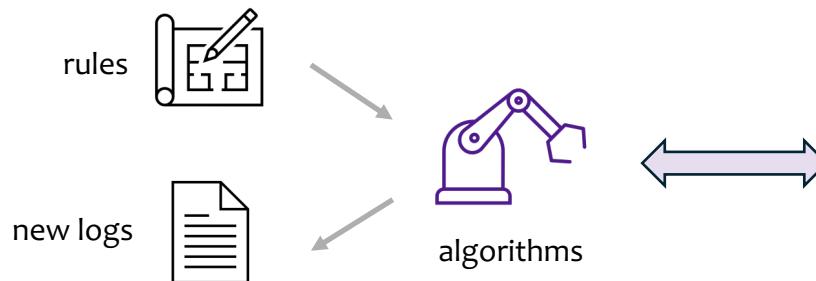
Logs are continuously evolving!



The performance of existing log parsers in practice remains unsatisfactory.

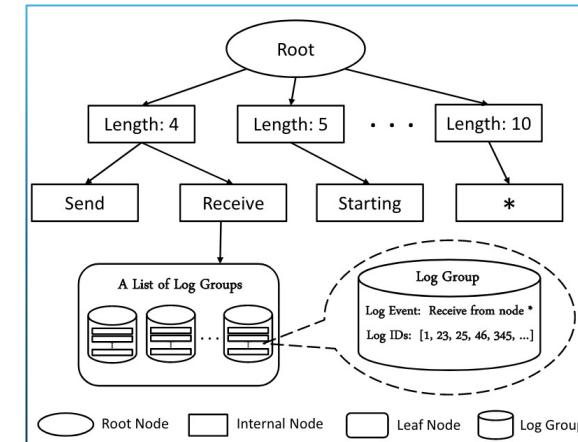
Syntax-based Log Parsers

Utilizing specific features or heuristics,
e.g., log length and word frequency.



heavily rely on manual-crafted rules!

Semantic-based Log Parsers



Example: Drain



Limitation of Existing Parsers

Logs are continuously evolving!



The performance of existing log parsers in practice remains unsatisfactory.

Syntax-based Log Parsers

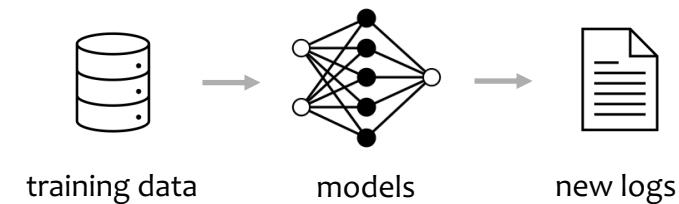
Utilizing specific features or heuristics,
e.g., log length and word frequency.



heavily rely on manual-crafted rules!

Semantic-based Log Parsers

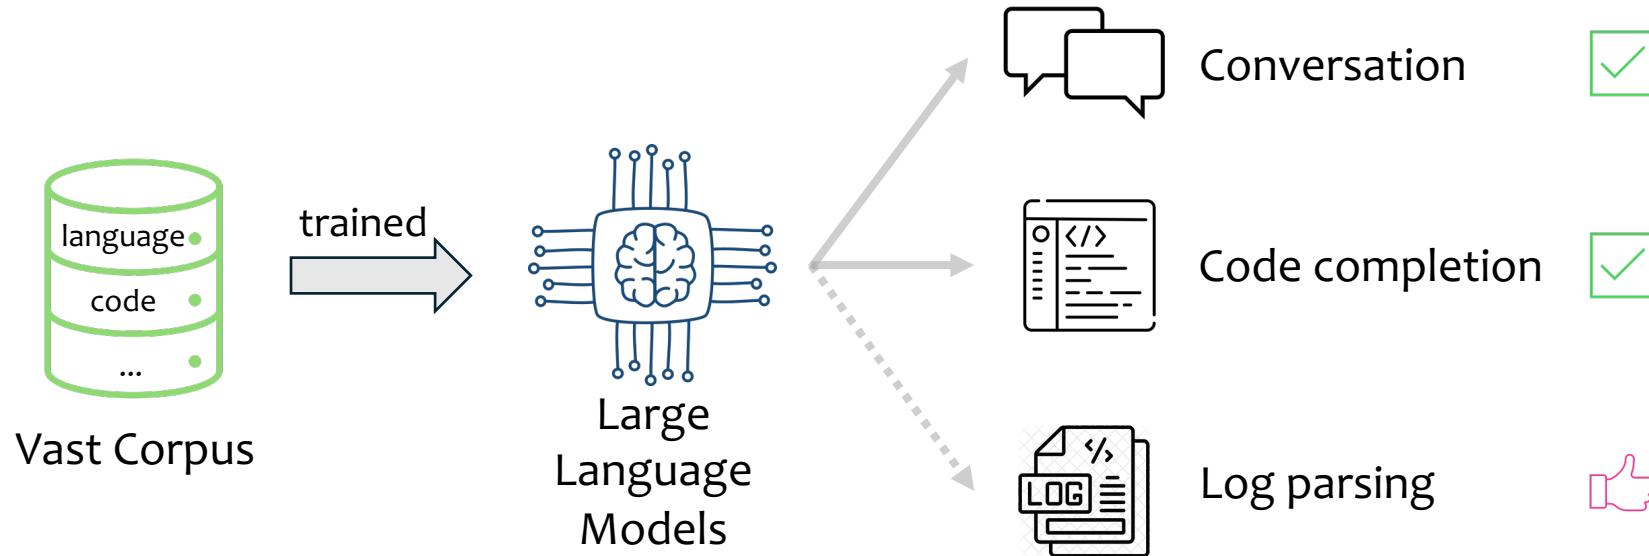
Utilizing deep learning models to learn
semantics and system-specific patterns.



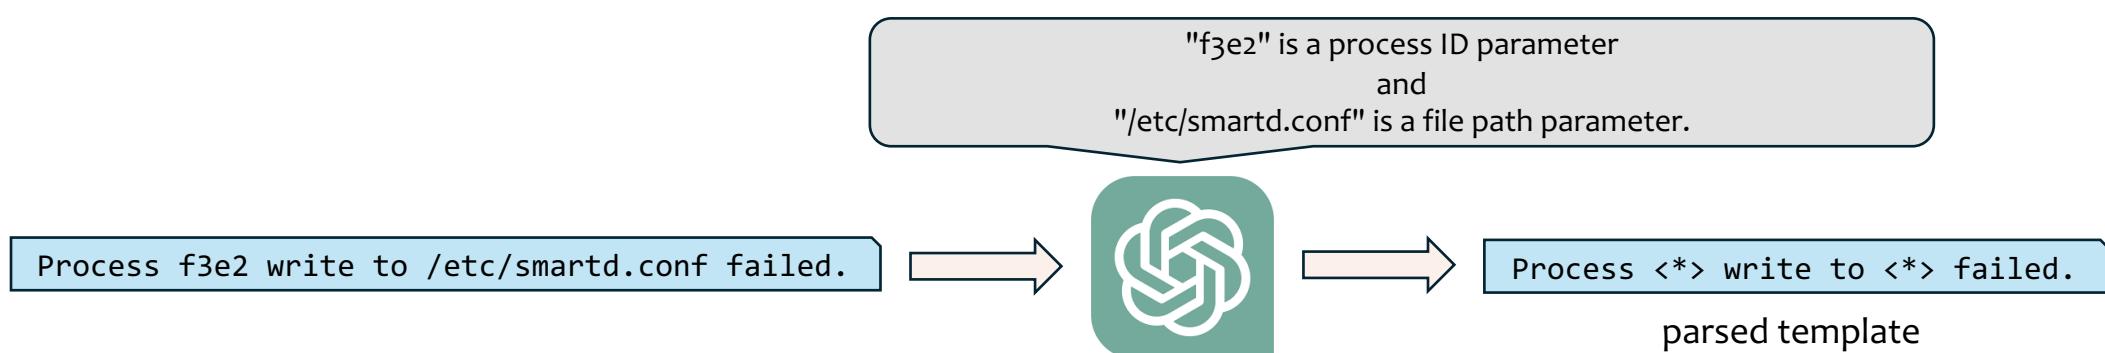
restricted by limited training data!



Opportunity: LLMs



Leveraging the power of LLMs to accurately parse logs is promising!





Challenges

Problem: How could we design a practical LLM-based log parsing framework?

Challenge 1: **Specialization**

- LLMs are not trained for the task of log parsing, resulting in compromised accuracy.

Challenge 2: **Consistency**

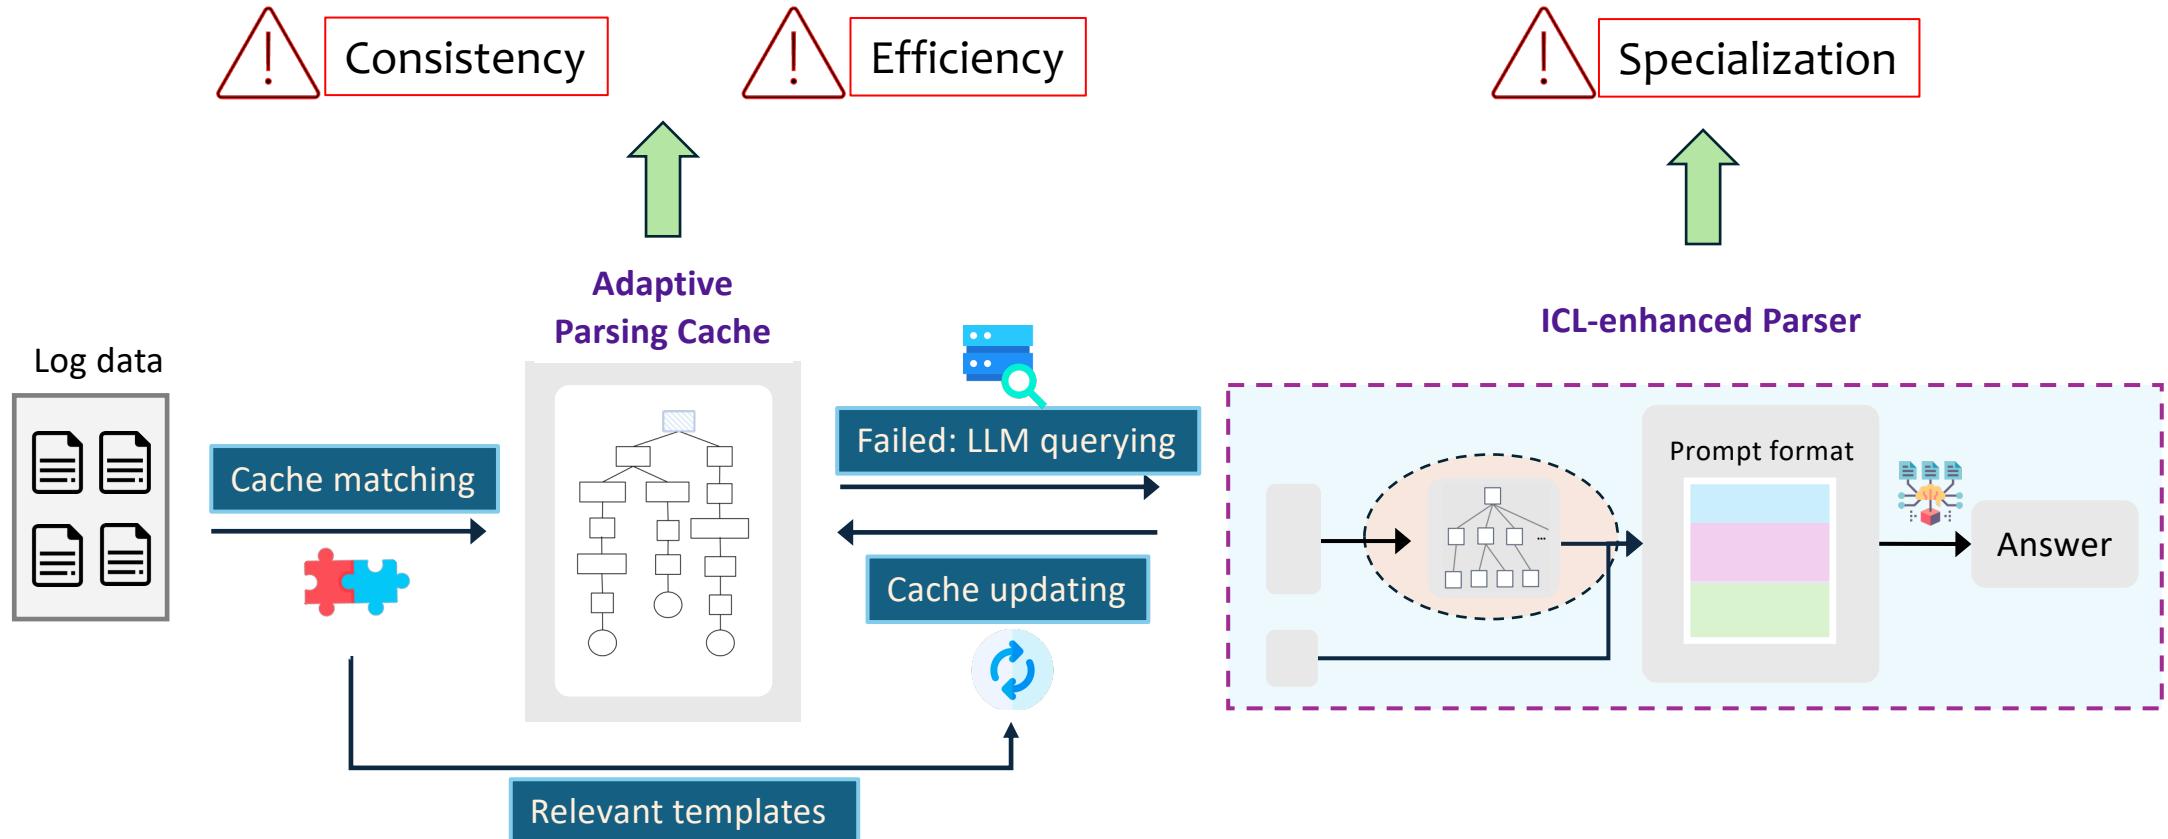
- The unstable responses of LLMs could lead to inconsistent log templates.

Challenge 3: **Efficiency**

- The latency and cost of querying LLMs fails to meet the requirements in large-scale scenarios.



Our Approach



Overall framework of LILAC



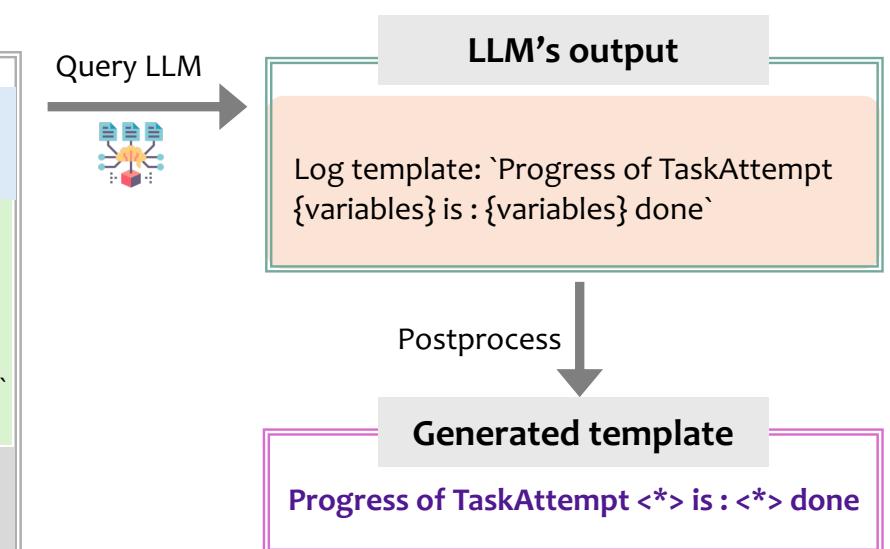
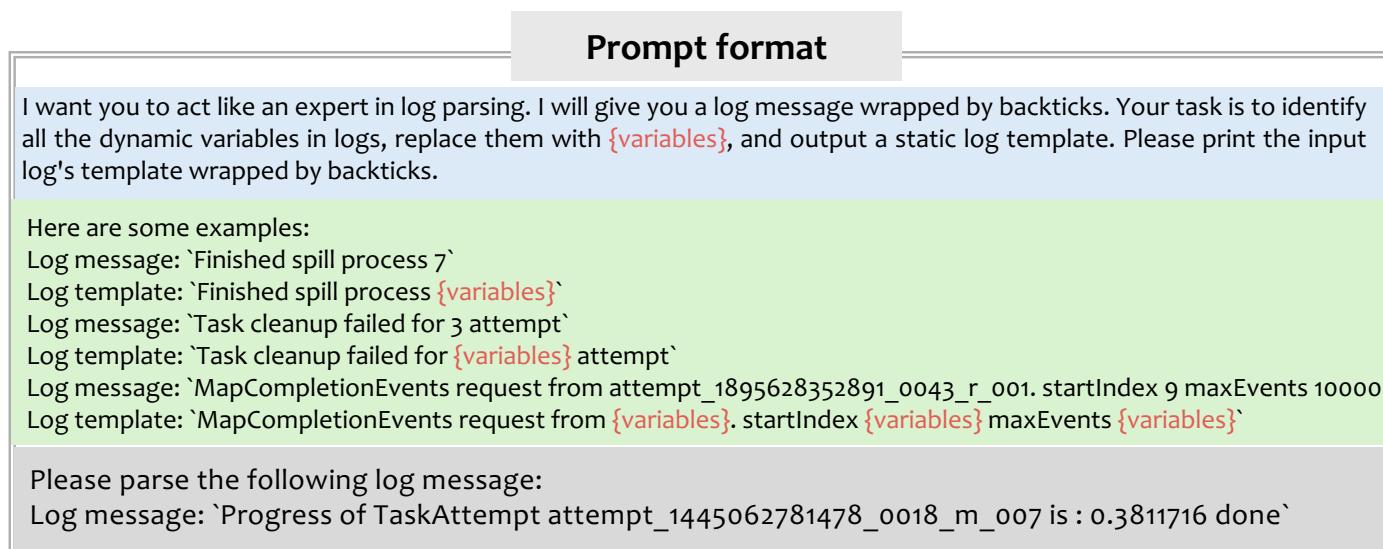
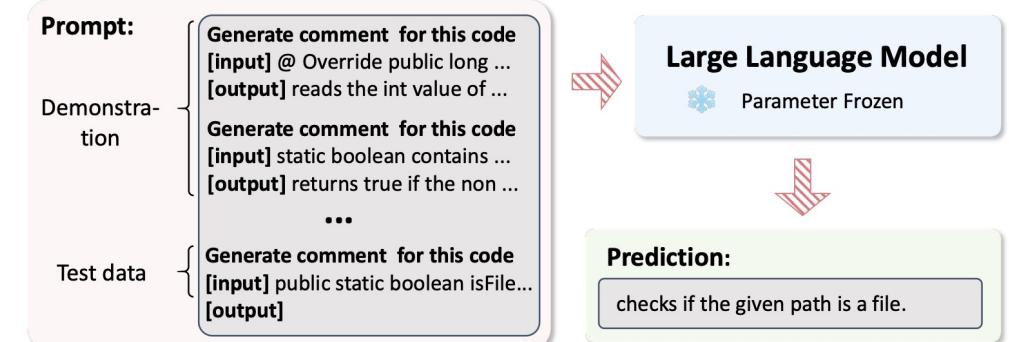
ICL-enhanced Parser



Specialization

The **in-context learning (ICL)** paradigm has been proven to be effective in various downstream tasks.

Understanding log parsing task without model tuning.





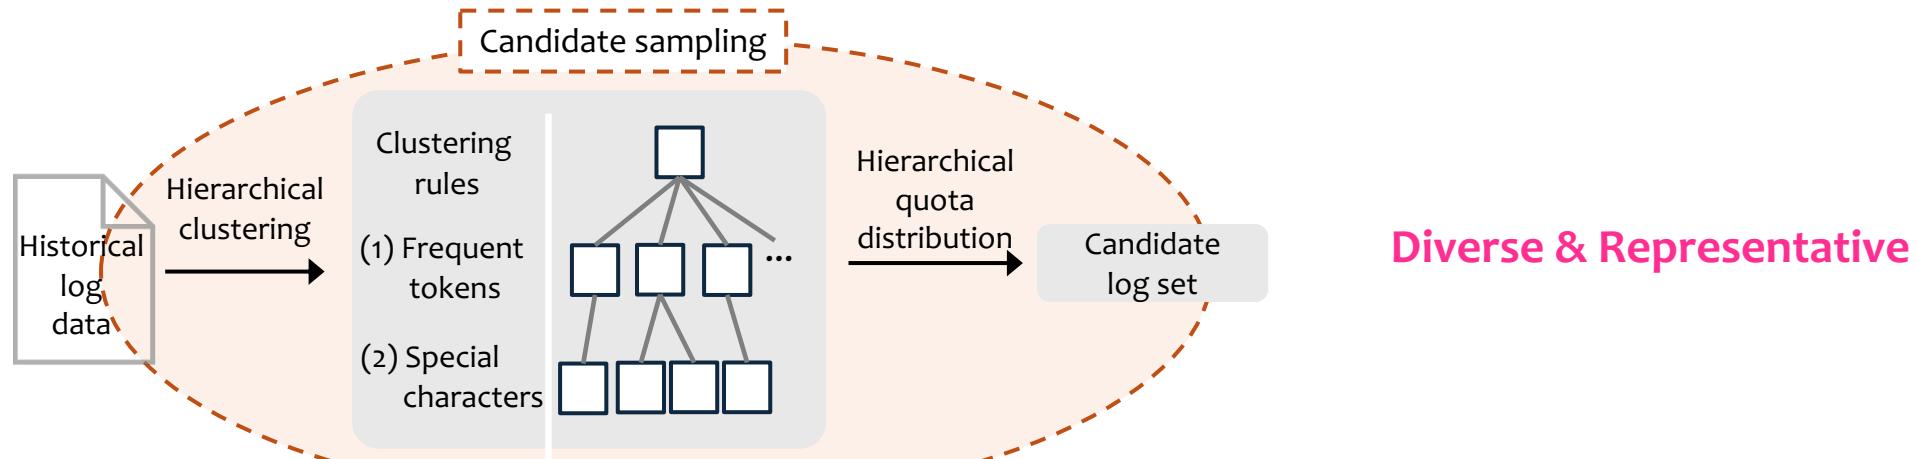
ICL-enhanced Parser

The quality of demonstrations can affect the ICL performance.

Problem: how to select **high-quality** demonstration examples from **substantial** log data?

Our Design

- Hierarchical candidate sampling from historical logs
- Demonstration selection from candidate log set





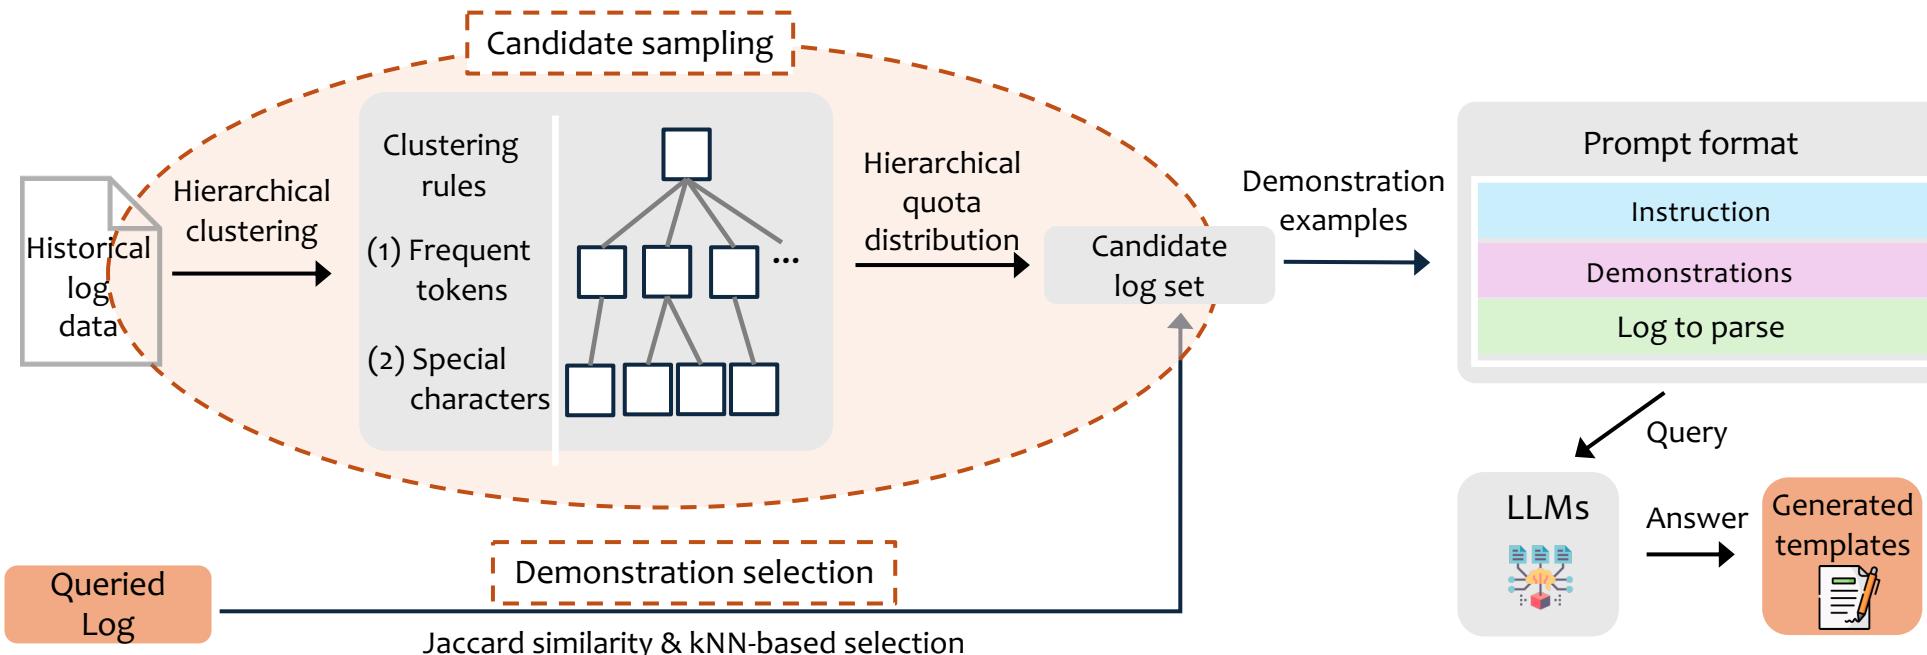
ICL-enhanced Parser

The quality of demonstrations can affect the ICL performance.

Problem: how to select **high-quality** demonstration examples from **substantial** log data?

Our Design

- *Hierarchical candidate sampling from historical logs*
- Demonstration selection from candidate log set





Adaptive Parsing Cache



Consistency



Efficiency

Our Design

Tree-based adaptive parsing cache to **store** and **refine** the generated log templates from the ICL-enhanced Parser

Motivation: the number of *log templates* is significantly smaller than the number of *logs* in real-world systems.



Adaptive Parsing Cache



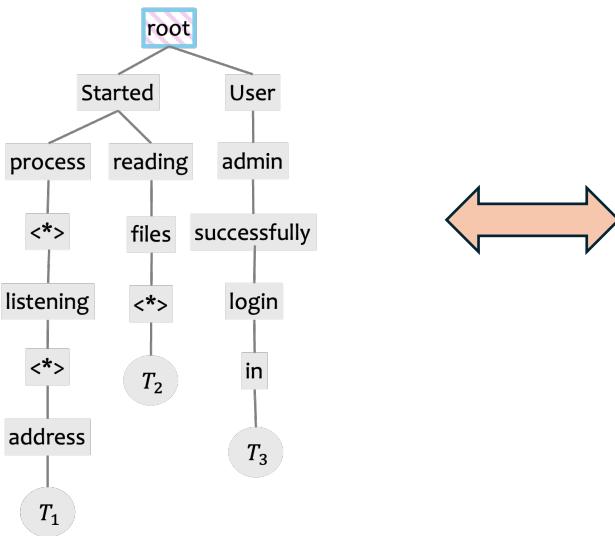
Consistency



Efficiency

Our Design

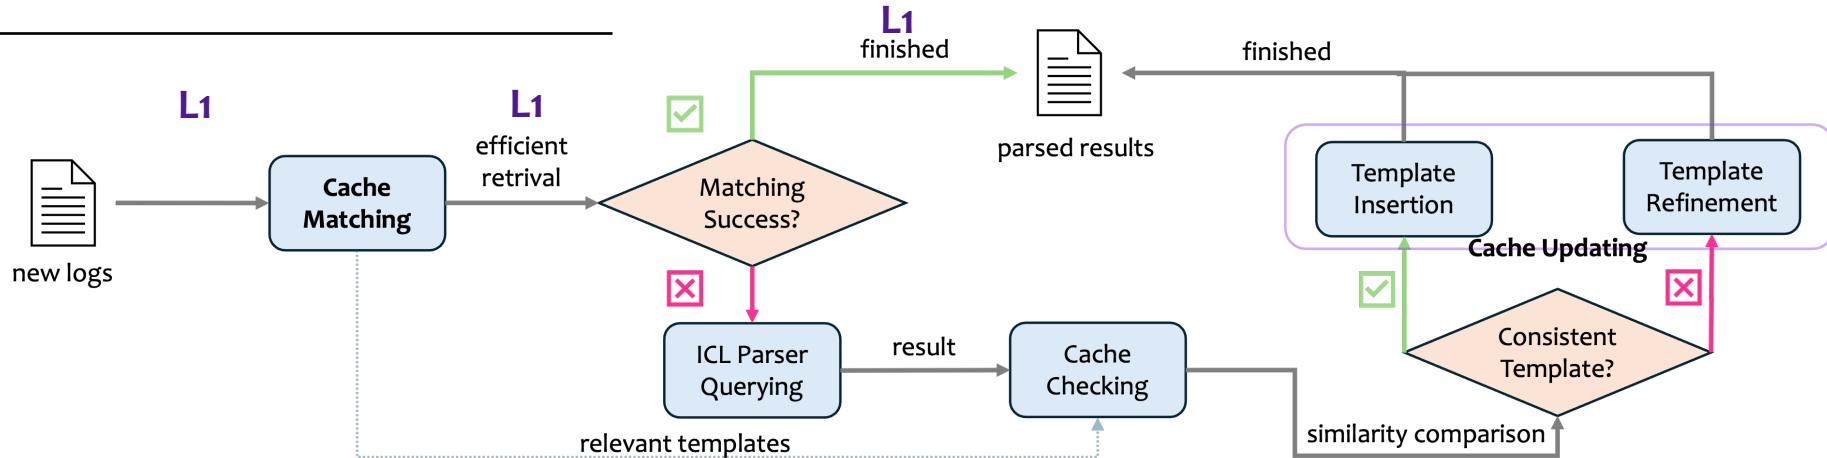
Tree-based adaptive parsing cache to **store** and **refine** the generated log templates from the ICL-enhanced Parser



Three Cached Templates:

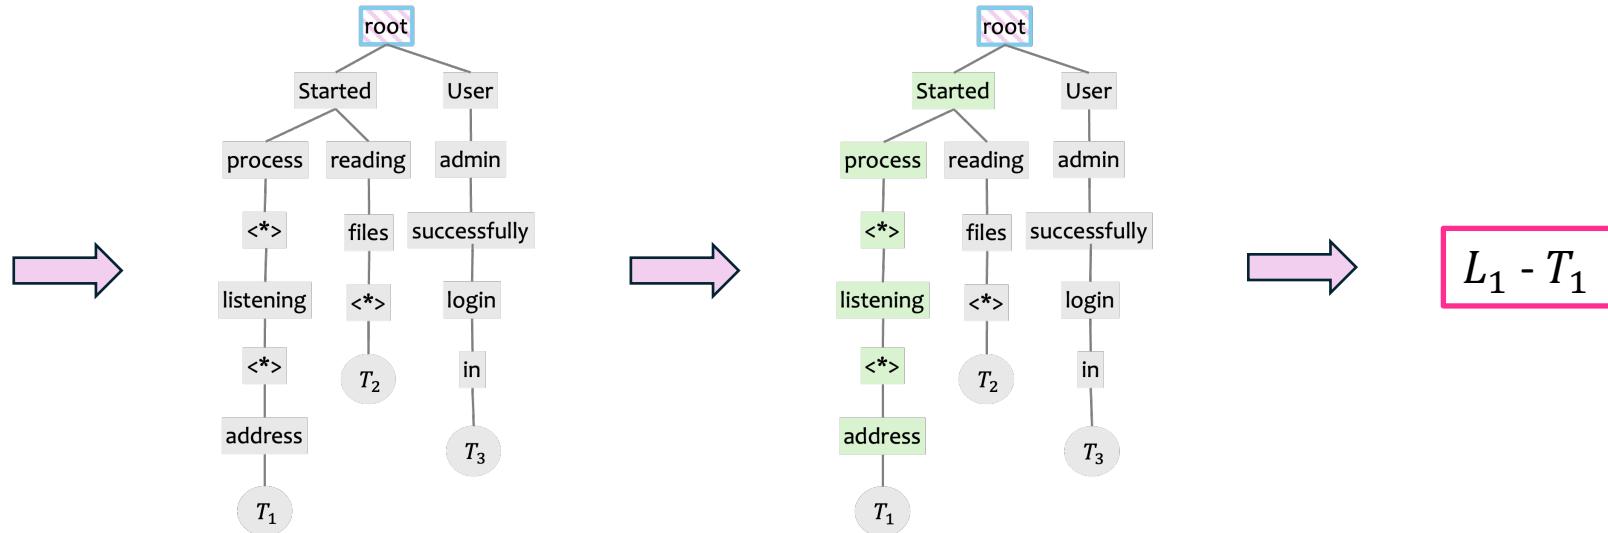
- T_1 : Started process $<*>$ listening $<*>$ port
- T_2 : Started reading $<*>$ files
- T_3 : User admin successfully log in

Adaptive Parsing Cache

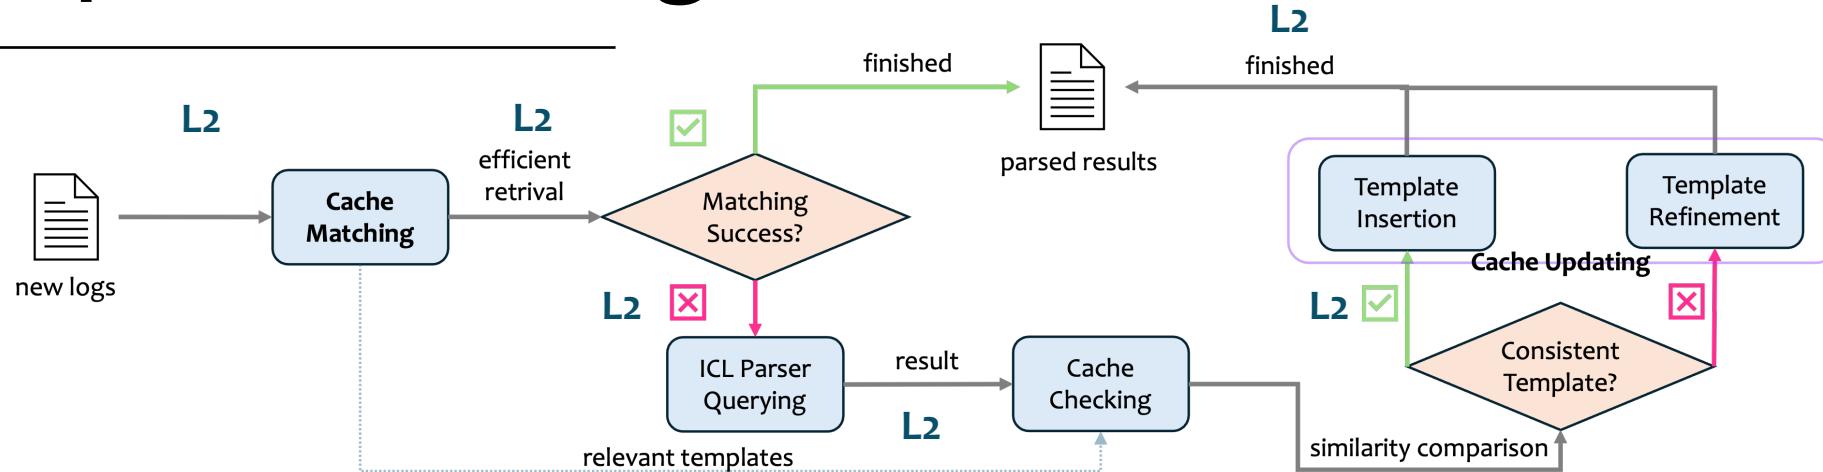


Avoid duplicate LLM queries!

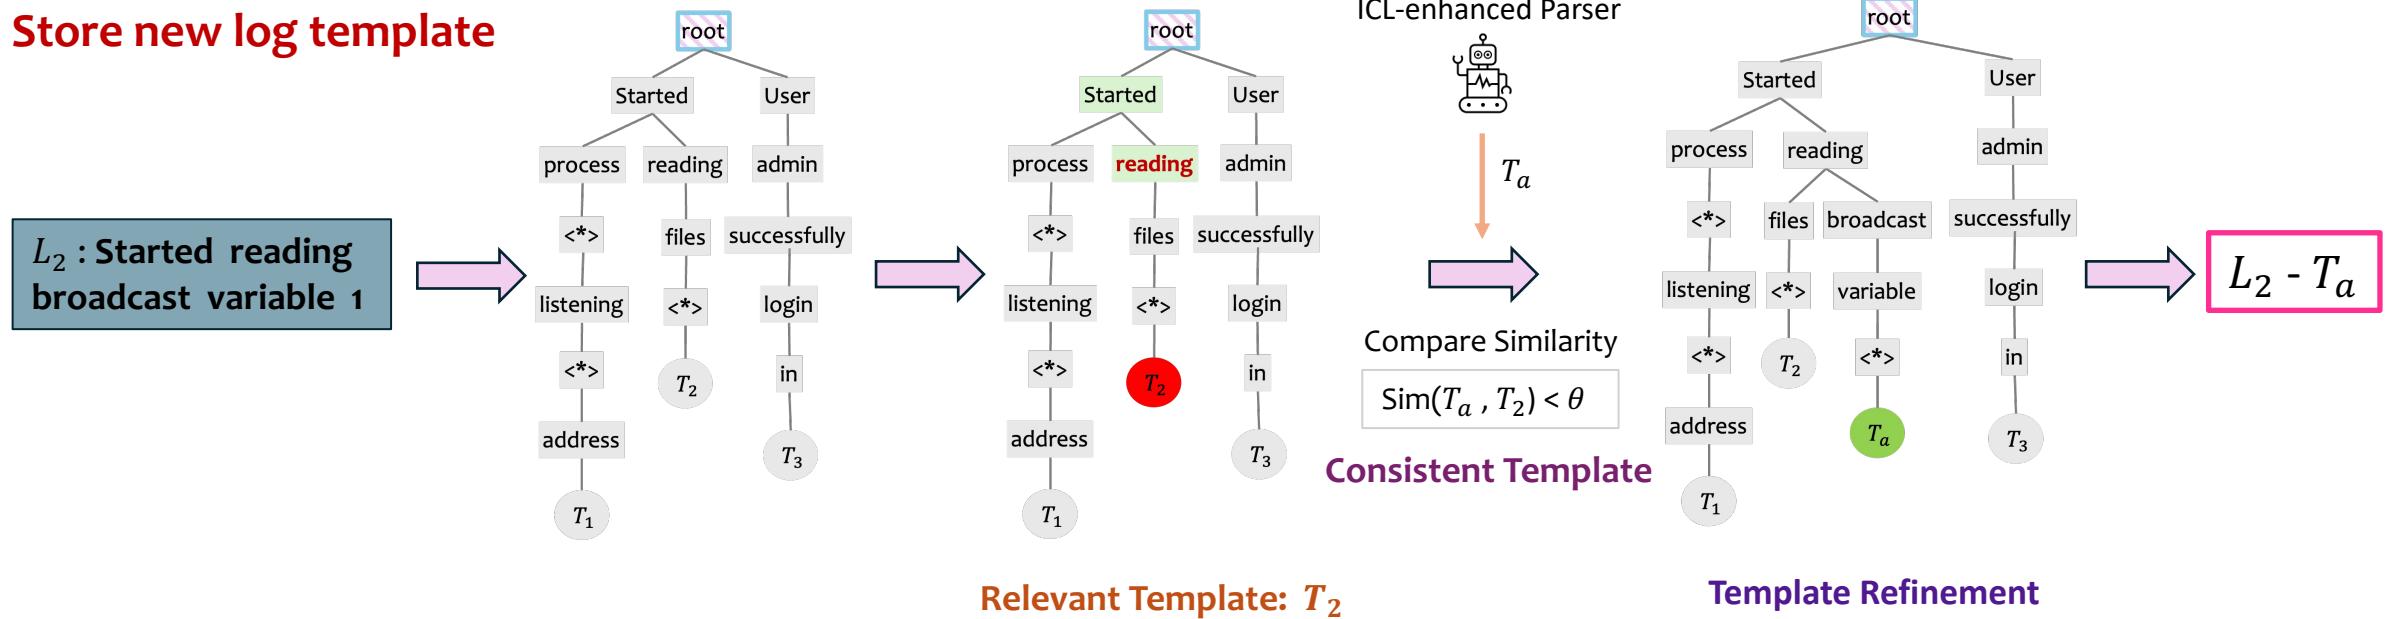
L_1 : Started process 1616
listening 127.0.0.1 address



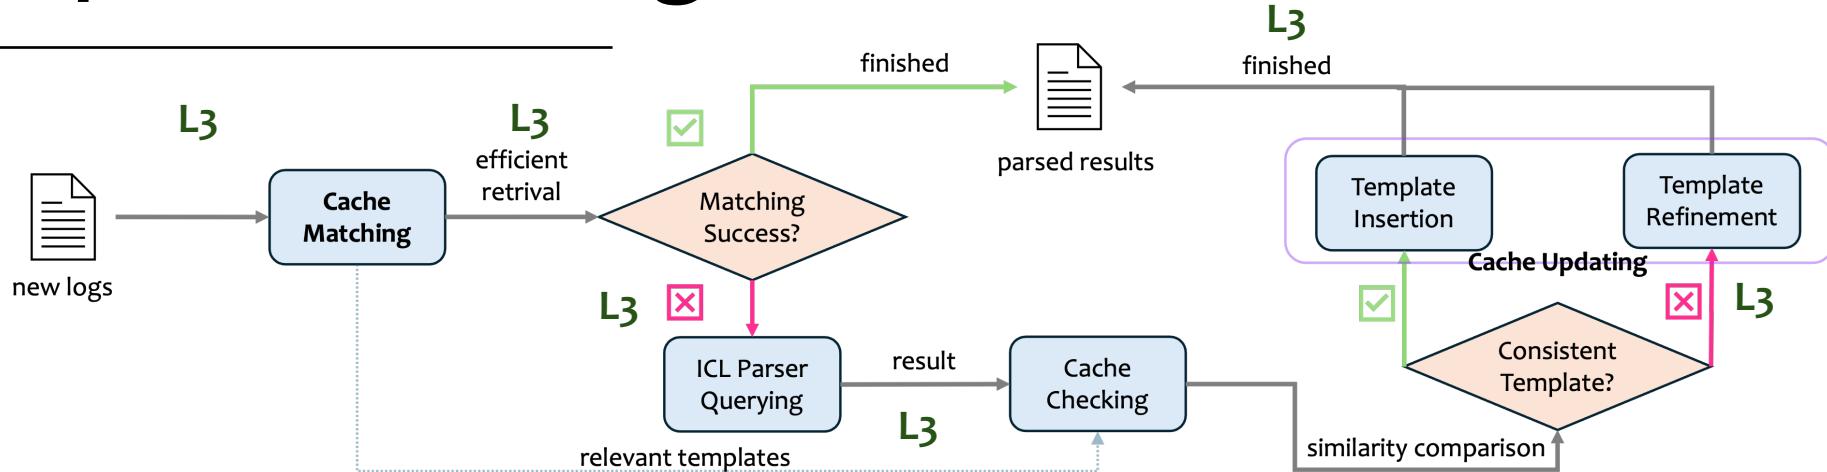
Adaptive Parsing Cache



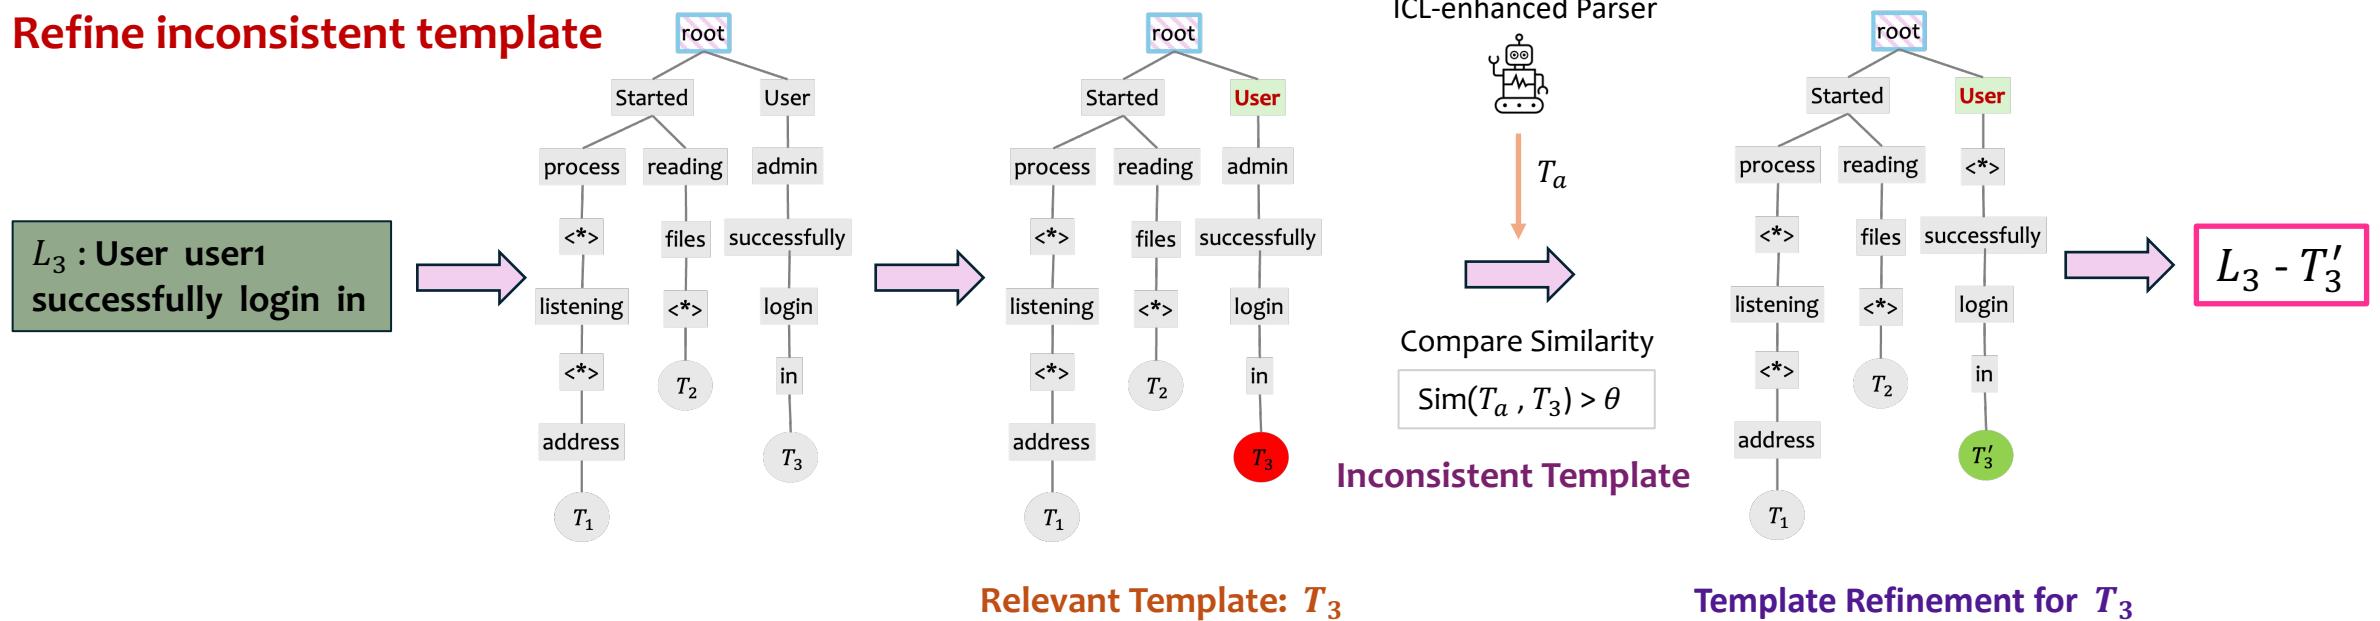
Store new log template



Adaptive Parsing Cache

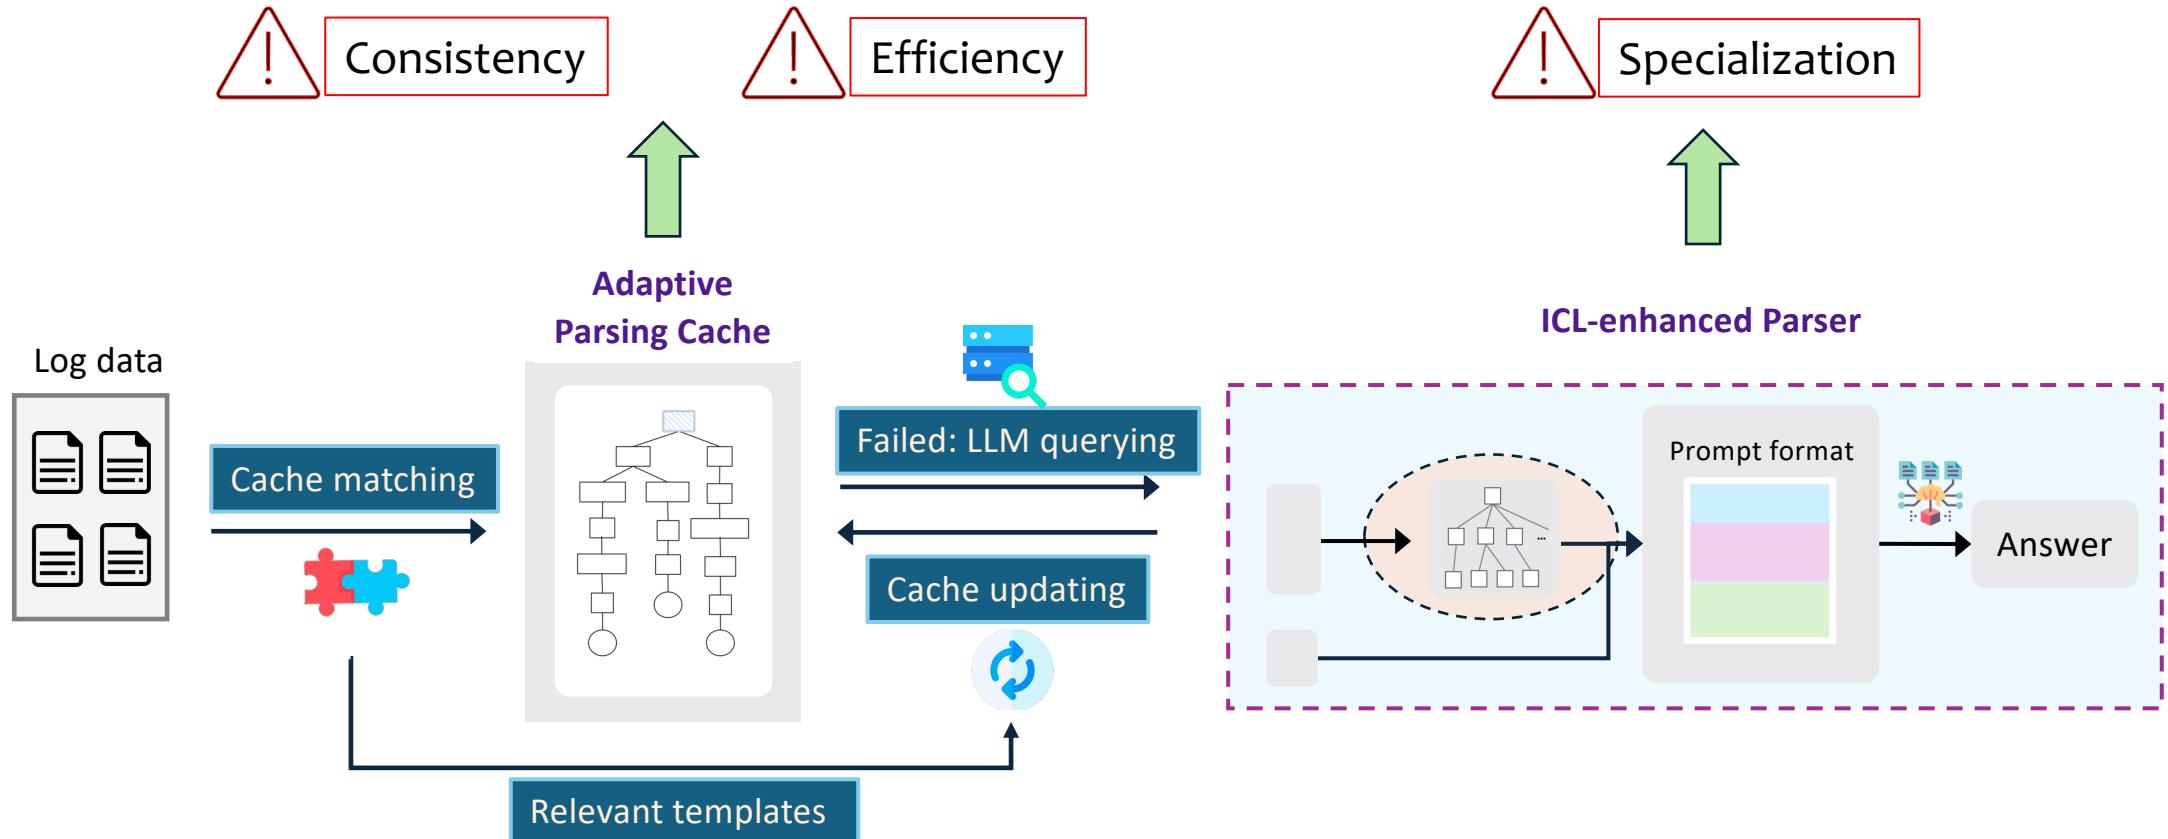


Refine inconsistent template





Our Approach



The collaboration between the two components resulted in effective and efficient log parsing.



Evaluation

Research Questions

- RQ1: What is the **effectiveness** of LILAC?
- RQ2: What is the **contribution** of each design?
- RQ3: What are the **capabilities** with different LLMs?
- RQ4: What is the **efficiency** of LILAC?

Datasets



Loghub-2.0

System	Dataset	# Log Templates	# Log Messages
Distributed systems	Hadoop	236	179,993
	HDFS	46	11,167,740
	OpenStack	48	207,632
	Spark	236	16,075,117
	Zookeeper	89	74,273
Super-computer systems	BGL	320	4,631,261
	HPC	74	429,987
	Thunderbird	1,241	16,601,745
Operating systems	Linux	338	23,921
	Mac	626	100,314
Server application	Apache	29	51,977
	OpenSSH	38	638,946
Standalone software	HealthApp	156	212,394
	Proxifier	11	21,320
Average		249.1	3,601,187

Settings

- LLM: gpt-3.5-turbo-0613
- Historical data rate: 20%
- Candidate samples: 32
- Demonstration examples: 3

State-of-the-art Baselines

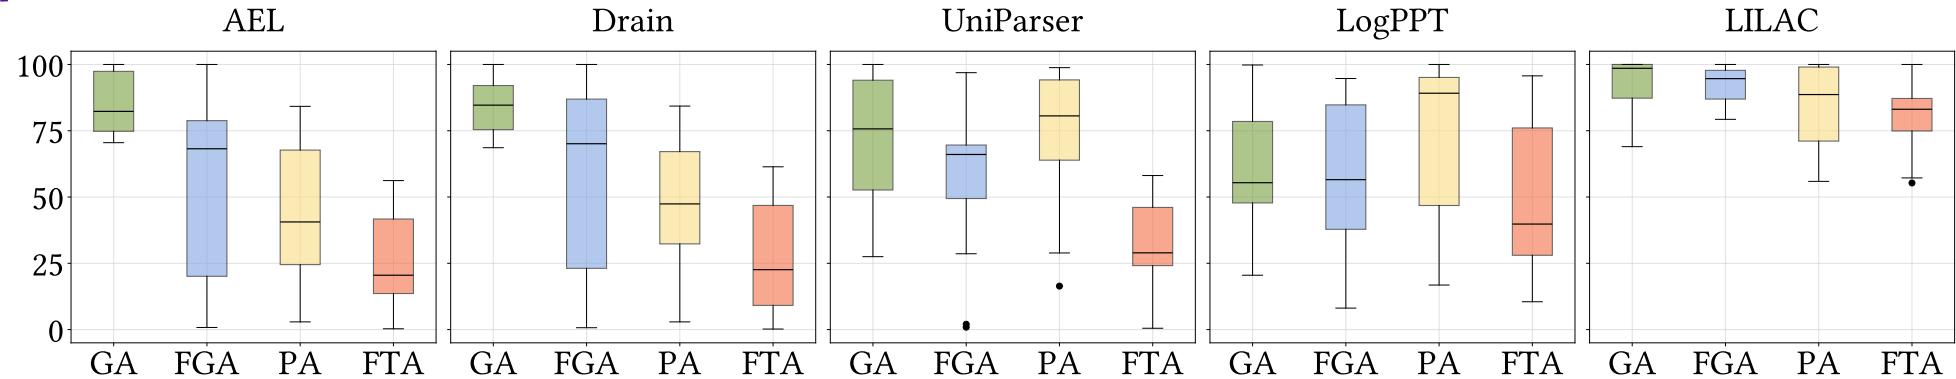
- AEL (*syntax-based*)
- Drain (*syntax-based*)
- UniParser (*semantic-based*)
- LogPPT (*semantic-based*)

Metrics

- Grouping Accuracy (GA)
- F1 score of Grouping Accuracy (FGA)
- Parsing Accuracy (PA)
- F1 score of Parsing Accuracy (FTA)

Evaluation

- **RQ1: Effectiveness**



- **RQ2: Ablation**

	GA	FGA	PA	FTA
LILAC	92.7	92.4	84.2	81.0
w/o ICL	83.5 (↓ 9.9%)	76.5 (↓ 17.2%)	62.6 (↓ 25.6%)	58.4 (↓ 27.9%)
w/ random selection	84.2 (↓ 9.2%)	80.6 (↓ 14.6%)	74.4 (↓ 11.6%)	66.7 (↓ 17.7%)
w/ random sampling	87.6 (↓ 5.5%)	80.9 (↓ 12.4%)	77.5 (↓ 8.0%)	68.3 (↓ 15.7%)
w/ LogPPT sampling	91.3 (↓ 1.5%)	84.8 (↓ 8.2%)	79.7 (↓ 5.3%)	74.9 (↓ 7.5%)

	GA	FGA	PA	FTA
LogPPT (original)	61.2	57.4	73.6	47.8
w/ parsing cache	90.8 (↑ 48.4%)	89.0 (↑ 55.1%)	78.0 (↑ 6.0%)	67.4 (↑ 41.0%)

- **LILAC outperforms all other state-of-the-art log parsers.**
- All designs of ICL-enhanced parsers and adaptive parsing cache contribute to the overall performance.

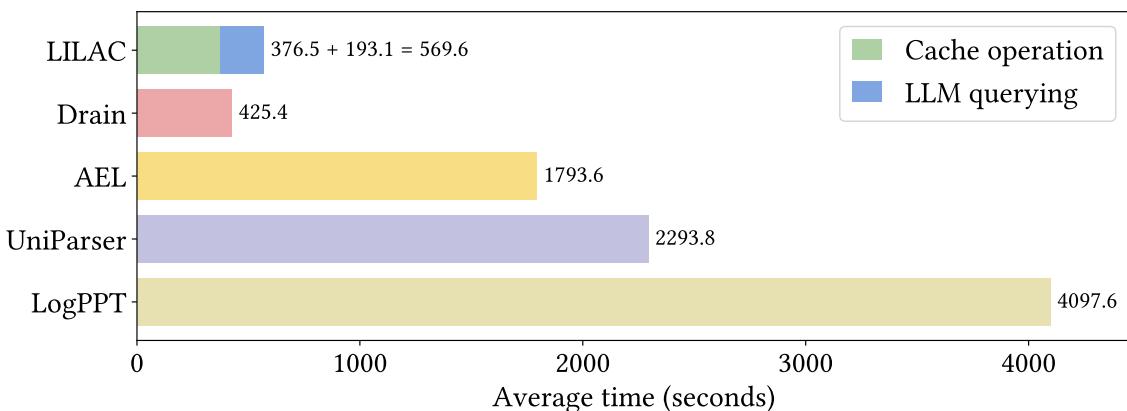


Evaluation

- RQ3: Different LLMs

	GA	FGA	PA	FTA
ChatGPT (175B)	92.7	92.4	84.2	81.0
Davinci (175B)	91.9 (↓ 0.9%)	92.9 (↑ 0.5%)	87.1 (↑ 3.4%)	81.5 (↑ 0.6%)
Curie (13B)	90.1 (↓ 2.8%)	87.6 (↓ 5.2%)	77.8 (↓ 7.6%)	71.2 (↓ 12.1%)

- RQ4: Efficiency

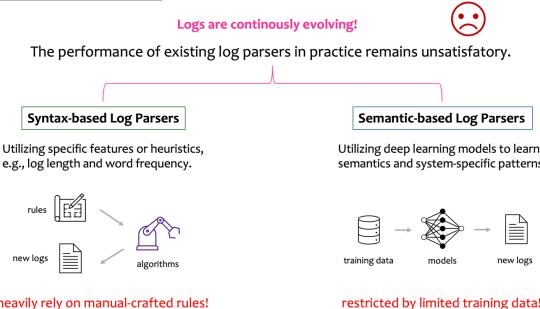


- LILAC consistently achieves high performance with different LLMs.
- The performance of LILAC can be influenced by the capabilities of LLMs.

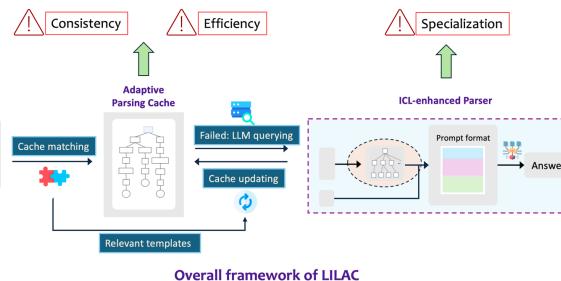
- LILAC successfully reduces millions of LLM queries to an average of 279.7.
- LILAC surpasses semantic-based methods and is comparable to the fastest syntax-based method in efficiency.

Conclusion

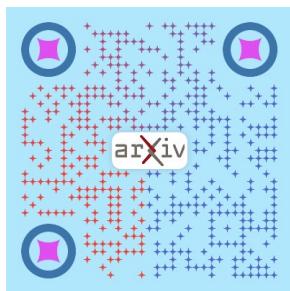
Limitation of existing parsers



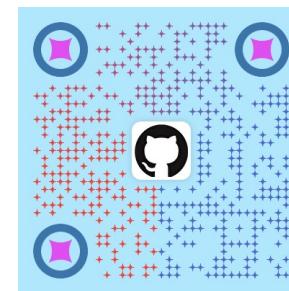
Our Approach



Motivation



Pre-print paper



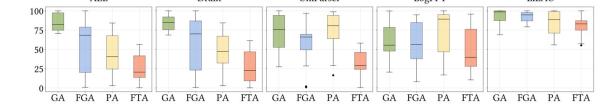
Artifact



LogPAI website

Evaluation

RQ1: Effectiveness



RQ2: Ablation

	GA	FGA	PA	FTA
LILAC	92.9	92.9	84.8	81.8
w/o ICL	88.4 (± 9.9%)	76.3 (± 17.2%)	62.4 (± 25.4%)	58.1 (± 27.9%)
w/ random selection	84.2 (± 9.2%)	80.4 (± 14.0%)	74.4 (± 11.6%)	66.7 (± 17.7%)
w/ random sampling	87.6 (± 5.5%)	80.9 (± 12.4%)	77.5 (± 8.0%)	68.3 (± 15.7%)
w/ LogPPT sampling	91.3 (± 1.5%)	84.1 (± 8.2%)	79.7 (± 5.3%)	74.9 (± 7.5%)
LogPPT (original)	61.2	57.4	73.6	47.8
w/ parsing cache	90.8 (± 48.4%)	89.0 (± 55.1%)	78.0 (± 6.0%)	67.4 (± 41.0%)

- LILAC outperforms all other state-of-the-art log parsers.
- All designs of ICL-enhanced parsers and adaptive parsing cache contribute to the overall performance.

Experiment

Thank you for listening

Q & A

Contact: zhjiang@link.cuhk.edu.hk