



Demystifying and Extracting Fault-indicating Information from Logs for Failure Diagnosis

Junjie Huang¹, Zhihan Jiang¹, Jinyang Liu¹, Yintong Huo¹, Jiazhen Gu¹,
Zhuangbin Chen², Cong Feng³, Hui Dong³, Zengyin Yang³, Michael R. Lyu¹

¹The Chinese University of Hong Kong

²Sun Yat-sen University ³Huawei Cloud

Read the paper!



香港中文大學
The Chinese University of Hong Kong



中山大學
SUN YAT-SEN UNIVERSITY




HUAWEI CLOUD



➤ Background

- Online service systems suffer from unplanned interruptions and failures
- IT companies must take timely and effective measures to respond to failures





➤ Background


- Log is an important source for failure diagnosis
 - Logs record a vast range of software runtime events
 - At least 31% failure diagnosis practices rely on logs in a commercial bank service system [1]
- Examining logs and gaining insights about faults is time-consuming and labor-intensive
 - Software systems produce a large volume of logs due to the scale and complexity

Logging Statements

```
# Logging statements from Spark (spark/storage/BlockManager.scala)
logError(s"Failed to report $blockId to master; giving up.")
logDebug(s"Putting block ${blockId} with replication took $usedTimeMs")
logInfo(s"Writing block $blockId to disk")
```

Log Message

17/08/22 15:50:46 ERROR BlockManager Failed to report rdd_5_1 to master; giving up.
17/08/22 15:50:55 DEBUG BlockManager Putting block rdd_0_1 with replication took 0
17/08/22 15:51:02 INFO BlockManager Writing block rdd_1_3 to disk
17/08/22 15:51:24 DEBUG BlockManager Putting block rdd_2_2 with replication took 0
17/08/22 15:52:36 ERROR BlockManager Failed to report rdd_3_3 to master; giving up.



Background



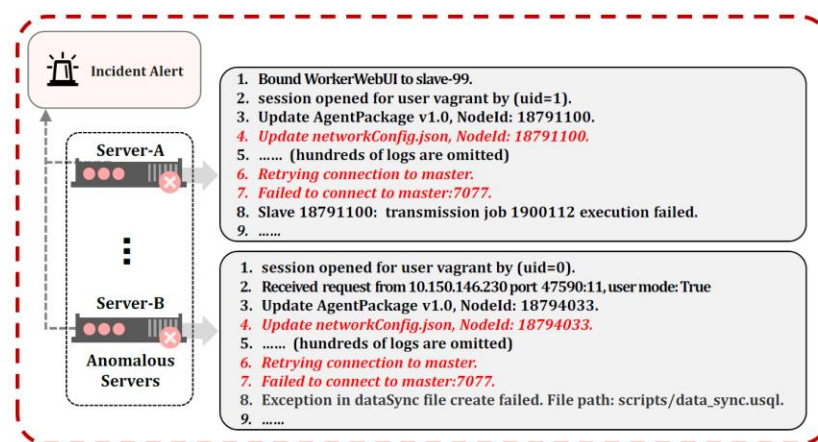
- Current log analysis approaches cannot find direct information for fault diagnosis

Log Analysis Approach	Identified information	Problem
① Log anomaly detection	Anomalous sessions	The sessions can be large (e.g., >100 logs)
② Log clustering	Incident-related logs	The logs can be long
③ Log parsing	Templates and variables	Events and variables may not relate to fault components (e.g., devices and VMs)

```

1 2008-11-09 20:55:54 PacketResponder 0 for block blk_321 terminating
2 2008-11-09 20:55:54 Received block blk_321 of size 67108864 from /10.251.195.70
3 2008-11-09 20:55:54 PacketResponder 2 for block blk_321 terminating
4 2008-11-09 20:55:54 Received block blk_321 of size 67108864 from /10.251.126.5
5 2008-11-09 21:56:50 10.251.126.5:50010:Got exception while serving blk_321 to /10.251.127.243:
6 2008-11-10 03:58:04 Verification succeeded for blk_321
7 2008-11-10 10:36:37 Deleting block blk_321 file /mnt/hadoop/dfs/data/current/subdir1/blk_321
8 2008-11-10 10:36:50 Deleting block blk_321 file /mnt/hadoop/dfs/data/current/subdir51/blk_321
  
```

① Anomalous log sessions



② Incident-related logs (in red)

```

17/08/22 15:50:46 ERROR BlockManager Failed to report rdd_5_1 to master; giving up.
17/08/22 15:50:55 DEBUG BlockManager Putting block rdd_0_1 with replication took 0
17/08/22 15:51:02 INFO BlockManager Writing block rdd_1_3 to disk
17/08/22 15:51:24 DEBUG BlockManager Putting block rdd_2_2 with replication took 0
17/08/22 15:52:36 ERROR BlockManager Failed to report rdd_3_3 to master; giving up.
  
```

Structured Logs

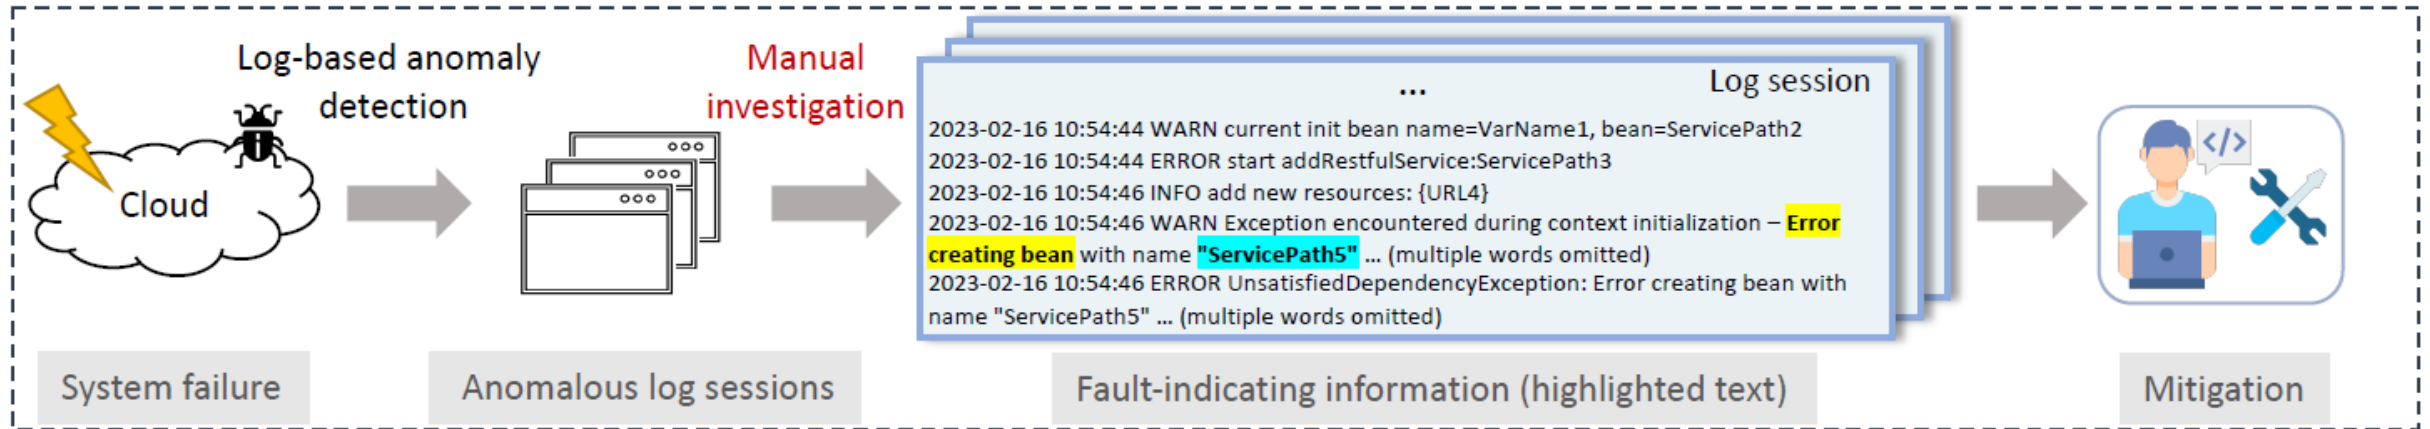
Datetime	Level	Component	Log Template	Variables
17/08/22 15:50:46	ERROR	BlockManager	Failed to report <*> to master; giving up.	rdd_5_1
17/08/22 15:50:55	DEBUG	BlockManager	Putting block <*> with replication took <*>	rdd_0_1, 0
17/08/22 15:51:02	INFO	BlockManager	Writing block <*> to disk	rdd_1_3
17/08/22 15:51:24	DEBUG	BlockManager	Putting block <*> with replication took <*>	rdd_2_2, 0
17/08/22 15:52:36	ERROR	BlockManager	Failed to report <*> to master; giving up.	rdd_3_3

③ Parsed log templates and variables



➤ Our works

- ***Fault-indicating information***: the fine-grained log segments that
 - indicate fault-related information
 - provide actionable insights to guide fault diagnosis



➤ A preliminary study on fault-indicating information



- Manually investigate how logs are leveraged by engineers through history reports
 - Fault-indicating description (FID) describes the symptoms of a fault
 - Fault-indicating parameter (FIP) denotes the precise position of the fault that needs action

CATEGORIES OF FAULT-INDICATING INFORMATION IN LOGS.

Category	Subtype	Example	Number
Description	Error Message	... url detection error! agent <u>taskId:f292c7e596d5435d9b9e9b9f47e1f872</u> , retCode is empty	34/96
	Missing Component	... execute template error, reason is Host name must not be empty	23/96
	Abnormal Behavior	... reader request line for <u>192.168.132.245:8080(https)</u> failed, read line timed-out	26/96
	Wrong Status	... httpCode is 404. requestEntity's type is GET. requestEntity's url is <u>/users/orders/task</u> . please check!	13/96
Parameter	Address	... httpCode is 404. requestEntity's type is GET. requestEntity's url is <u>/users/orders/task</u> . please check!	18/71
	Component ID	... query ... failed. <u>historyid=51890bae-57c6-47a3-b37d-62df9d2f3c87</u>	28/71
	Parameter Name	... cannot get topicInfo for consumer by topic: <u>alarm_and_event_data</u>	25/71

¹ Descriptions is marked in **BLOD** and parameters is in UNDERLINE.

➤ A preliminary study on fault-indicating information



- Manually investigate how logs are leveraged by engineers through history reports
 - Fault-indicating description (FID) describes the symptoms of a fault
 - Fault-indicating parameter (FIP) denotes the precise position of the fault that needs action
- **Challenges:**
 - **Large volume of logs:** Only 1.7% of log lines contain either FID or FIP
 - **Noisy semantics in logs:** Only 14.1% of words are FID or FIP



LoFI: A method to extract log fault-indicating information

- Two stages in LoFI: Log selection and prompt-based extraction

Input: Anomalous Log Sessions

...

l_1 : WARN current init bean name=VarName1, bean=ServicePath2

l_2 : ERROR start addRestfulService: ServicePath3

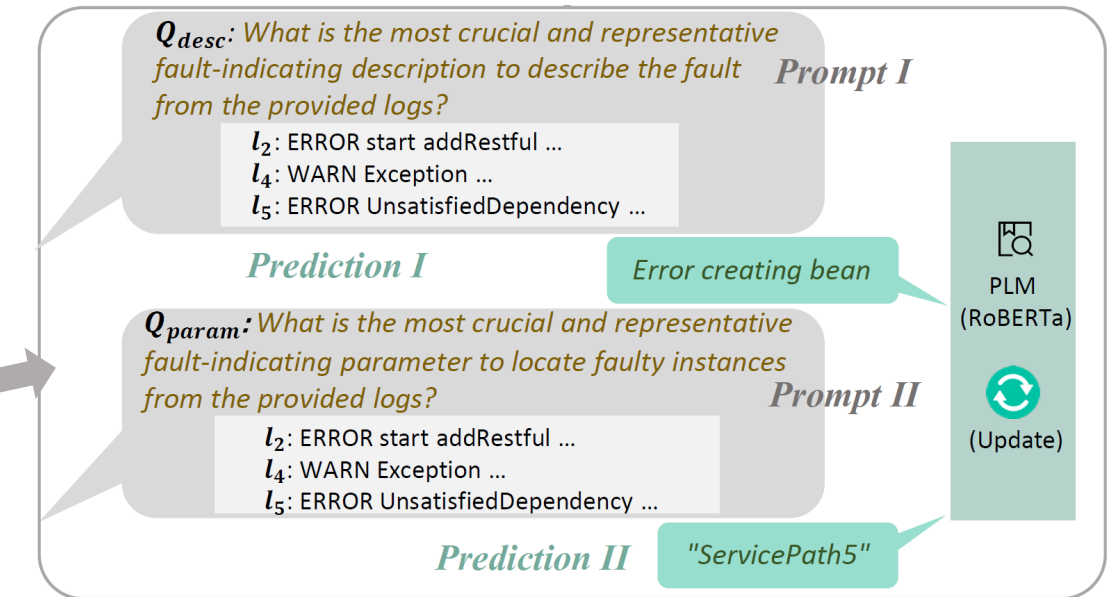
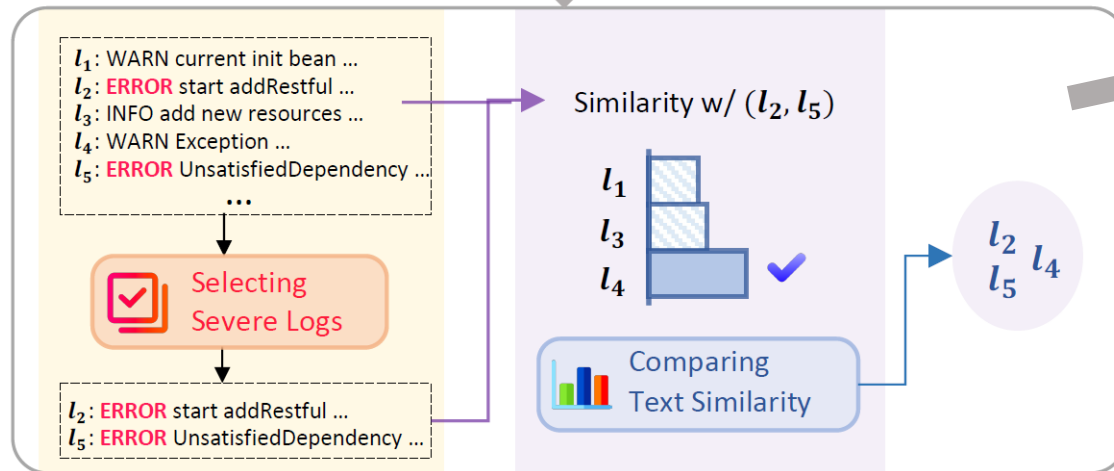
l_3 : INFO add new resources: {Url4}

l_4 : WARN Exception encountered during context initialization Error creating bean with name 'ServicePath5'

l_5 : ERROR UnsatisfiedDependencyException: Error creating bean with name 'ServicePath5'

...

1. Log Selection

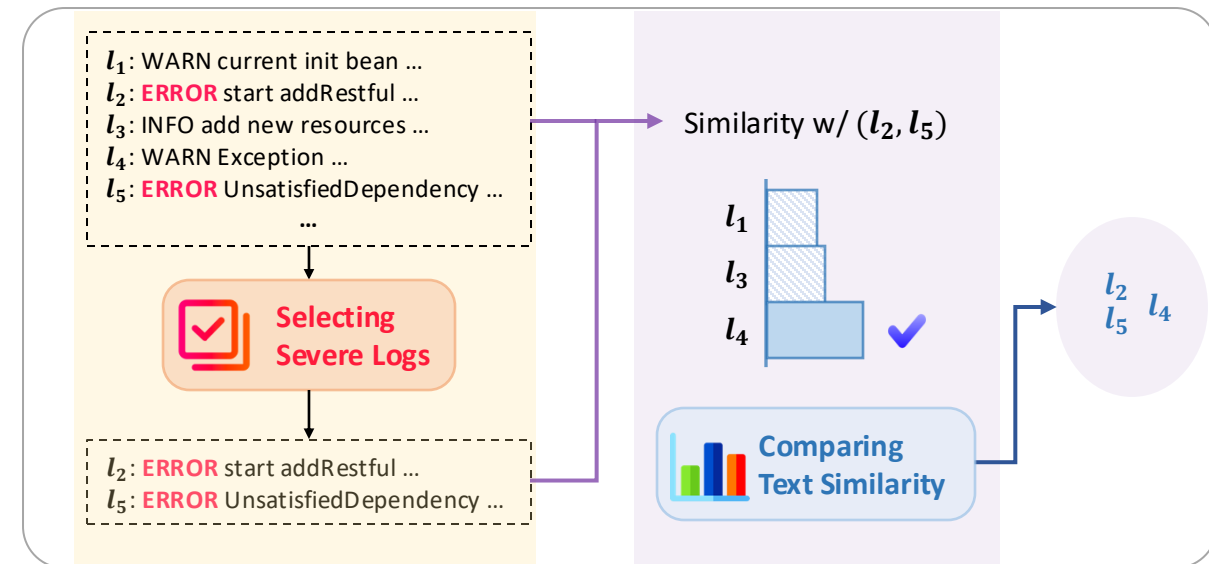


Output: Fault-indicating Information

Fault-indicating description: Error creating bean
Fault-indicating parameter: "ServicePath5"

➤ Stage-1: log selection

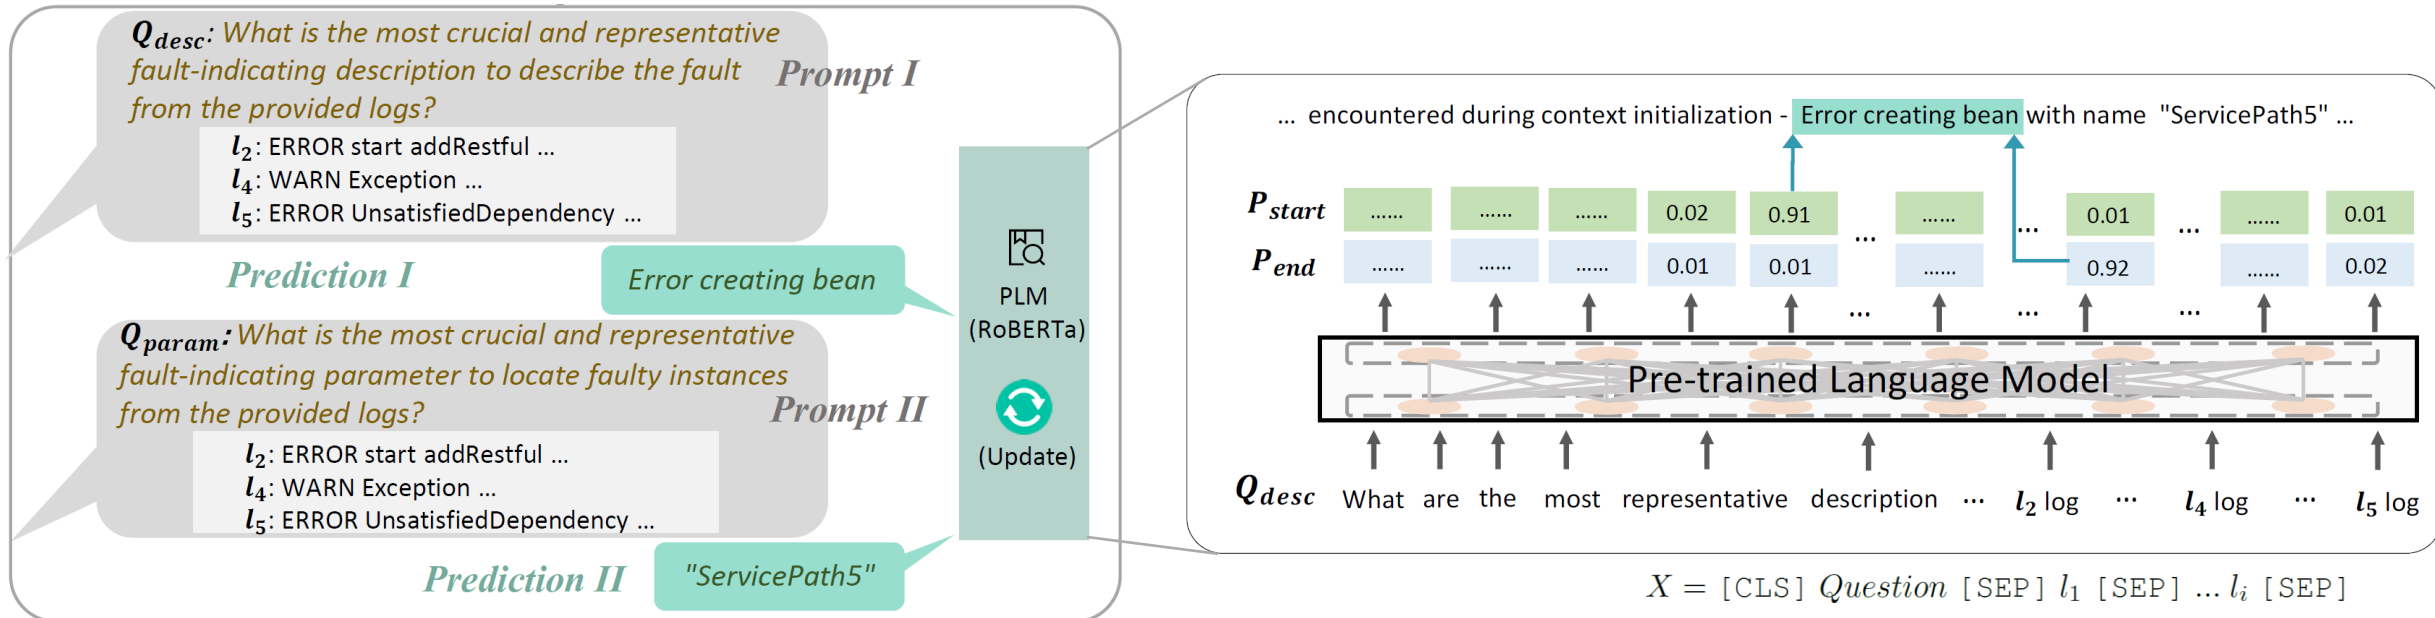
- Problem: abnormal log sessions contain a large volume of logs
 - Irrelevant information can hinder extraction
 - Language models struggle to understand long context
- Method:
 - 1st step: Select logs with severe logging levels
 - i.e., FATAL > ERROR > WARN > INFO > DEBUG > TRACE > Others
 - 2nd step: Identify semantically similar logs
 - RoBERTa-base encoder yields log embeddings
 - Rank the textual similarity of embeddings





Stage-2: prompt-based extraction

- Extraction using pretrained language models (PLM) with prompt-based tuning
 - We use two prompt questions for FID and FIP
 - PLM predicts start and end position of fault-indicating information
 - Fine-tune the PLM with logs and ground truth fault-indicating information
 - We use UniXCoder as the PLM



- Research questions
 - RQ1: How effective is LoFI in the *offline* setting?
 - RQ2: How log selection affects the results of LoFI?
 - RQ3: How prompt-based tuning affects the results of LoFI?
 - RQ4: How LoFI helps SREs to diagnose in *online* setting?
- Datasets
 - **FIBench**: 71 examples collected from Apache Spark
 - **Industry**: 88 examples collected from CloudA

DATASET STATISTICS

Dataset	Total Logs	Logs per Session	Faults	FID	FIP
FIBench	1,225,287	39.9	71	71	37
Industry	2,721,013	64.3	88	88	68



➤ RQ1: Effectiveness of LoFI

- LoFI performs fault-indicating information extraction well!
 - LoFI achieves high accuracy in recognizing FID and FIP
 - LoFI outperforms all baselines by a large margin
- FID prediction is more accurate than FIP prediction

Method	FIBench-FID			FIBench-FIP			Industry-FID			Industry-FIP		
	Precision	Recall	F1	Precision	Recall	F1	Precision	Recall	F1	Precision	Recall	F1
TF-IDF*	3.4	2.6	2.8	2.8	1.4	1.8	5.3	4.2	4.4	2.5	2.4	2.4
TextRank*	12.8	12.8	12.3	0.0	0.0	0.0	17.9	18.0	17.1	5.2	3.6	4.2
LogSummary*	4.5	5.4	4.5	3.6	1.1	1.7	16.2	13.5	14.2	6.9	4.9	5.5
ChatGPT-Zeroshot*	<u>59.6</u>	30.1	38.2	9.7	1.3	2.2	<u>47.2</u>	29.9	33.2	32.1	33.3	32.2
ChatGPT-ICL	<u>53.3</u>	<u>51.6</u>	<u>49.6</u>	<u>46.5</u>	<u>44.4</u>	<u>44.9</u>	<u>45.1</u>	<u>33.3</u>	<u>35.9</u>	<u>41.3</u>	<u>38.3</u>	<u>37.0</u>
LoFI (ours)	87.4	87.6	87.4	80.6	80.6	80.6	73.8	72.0	72.2	70.0	60.9	62.8

¹ We use * to denote unsupervised methods, others are supervised ones.



➤ RQ2: How log selection (LS) affects LoFI?

- The two-stage LS in full LoFI is better than other LS methods
- Diagnosing only with logs in error and warning levels is not enough
 - Compression ratio (CR) is high, indicating noisy logs are kept


RESULTS OF DIFFERENT LOG SELECTION (LS) METHODS (%)

	FIBench				Industry			
	CR	Acc.	F1-FID	F1-FIP	CR	Acc.	F1-FID	F1-FIP
Full LoFI	39.9	100.0	87.4	80.6	62.7	98.2	72.2	62.8
- LS=HighestCtx	38.1	87.2	77.6	78.2	59.1	91.1	67.6	58.9
- LS=Highest	34.5	76.9	69.9	72.2	48.2	91.1	67.0	60.8
- LS=ErrorWarn	7.9	71.8	65.4	59.3	86.4	100.0	61.0	29.5
- LS=Error	4.0	23.1	16.7	13.9	44.6	85.7	64.0	56.5
- w/o LS	-	-	76.9	75.6	-	-	50.2	43.7

* CR denotes compression ratio, Acc. denotes the selection accuracy.

- LS=Highest (highest level logs)
- LS=ErrorWarn (error and warning level logs)
- LS=HighestCtx (highest level and context logs)
- LS=Error (error level logs)

➤ RQ3: How prompt-based tuning affects the results?



- Prompt variants:
 - Using prompt question is helpful for extracting fault-indicating information
 - Prompt question with explicit instructions can improve performance
 - Fine-tuning can improve the performance even with a small amount of training data
- Pretraining Language Models:
 - Using UniXCoder performs better than using other models

RESULTS OF PROMPT-BASED TUNING VARIANTS (%)

	FIBench		Industry	
	F1-FID	F1-FIP	F1-FID	F1-FIP
Full LoFI	87.4	80.6	72.2	62.8
- Prompt=LessInfo	81.3	78.3	68.6	57.9
- w/o Prompt	58.6	29.0	64.4	25.8
- w/o Tuning	34.8	1.8	11.3	5.6
- PLM=CodeBERT	81.5	76.4	69.6	49.5
- PLM=RoBERTa	83.8	78.3	64.7	41.4
- PLM=BERT	77.5	78.7	65.4	14.6

- Prompt=LessInfo (prompt with less information)
 - + “Fault-indicating descriptions in the following logs:” + logs
 - + “Fault-indicating parameters in the following logs:” + logs
- w/o Prompt (without prompt question)
 - w/o Tuning (without fine-tuning on the training set)



➤ RQ4: How LoFI assists in online diagnosis?

- User study after one-month deployment:
 - 10 engineers rate the 50 fault examples during the deployment period
- Q1: Does FID accurately represent anomalous events?
 - Average rating: 4.34 (5-point Likert scale)
- Q2: Does FIP accurately identify anomalous components?
 - Average rating: 4.02 (5-point Likert scale)
- Q3: Would extracting fault-indicating information aid in fault diagnosis?
 - All participants agree!

Conclusions

- Fault-indicating information are useful
 - Fault-indicating descriptions describe the symptoms of a fault
 - Fault-indicating parameters denote the precise position of the fault that needs action
- LoFI can successfully extracts fault-indicating information
 - Log selection and prompt-based tuning are effective in improving the performance



Q & A